
MONTRÉAL.AI ACADEMY: ARTIFICIAL INTELLIGENCE 101

FIRST WORLD-CLASS OVERVIEW OF AI FOR ALL

VIP AI 101 CHEATSHEET

A PREPRINT

Vincent Boucher*
MONTRÉAL.AI
Montreal, Quebec, Canada
info@montreal.ai

March 3, 2020

ABSTRACT

For the purpose of entrusting all sentient beings with powerful AI tools to learn, deploy and scale AI in order to enhance their prosperity, to settle planetary-scale problems and to inspire those who, with AI, will shape the 21st Century, **MONTRÉAL.AI** introduces this **VIP AI 101 CheatSheet** for All.

**MONTRÉAL.AI is preparing a global network of education centers.*

*****ALL OF EDUCATION, FOR ALL.** MONTRÉAL.AI is developing a teacher (**Saraswati AI**) and an agent learning to orchestrate synergies amongst academic disciplines (**Polymatheia AI**).*

Curated Open-Source Codes and Science: <http://www.academy.montreal.ai/>.

Keywords AI-First · Artificial Intelligence · Deep Learning · Reinforcement Learning · Symbolic AI

1 AI-First

TODAY'S ARTIFICIAL INTELLIGENCE IS POWERFUL AND ACCESSIBLE TO ALL. AI is capable of transforming industries and opens up a world of new possibilities. **What's important is what you do with AI and how you embrace it.** To pioneer AI-First innovations advantages: start by exploring how to apply AI in ways never thought of.

The Emerging Rules of the AI-First Era: Search and Learning.

*"Search and learning are general purpose methods that continue to scale with increased computation, even as the available computation becomes very great." — Richard Sutton in *The Bitter Lesson**

The Best Way Forward For AI².

"... so far as I'm concerned, system 1 certainly knows language, understands language... system 2... it does involve certain manipulation of symbols... Gary Marcus ... Gary proposes something that seems very natural... a hybrid architecture... I'm influenced by him... if you look introspectively at the way the mind works... you'd get to that distinction between implicit and explicit... explicit looks like symbols." — Nobel Laureate Danny Kahneman at AAAI-20 Fireside Chat with Daniel Kahneman <https://vimeo.com/390814190>

In *The Next Decade in AI*³, Gary Marcus proposes a hybrid, knowledge-driven, reasoning-based approach, centered around cognitive models, that could provide the substrate for a richer, more robust AI than is currently possible.

^{*}Founding Chairman at MONTRÉAL.AI <http://www.montreal.ai> and QUÉBEC.AI <http://www.quebec.ai>.

²<https://montrealartificialintelligence.com/aidebate/>

³<https://arxiv.org/abs/2002.06177v3>

2 Getting Started

Tinker with neural networks in the browser with *TensorFlow Playground* <http://playground.tensorflow.org/>.

- **Deep Learning Drizzle** <https://deep-learning-drizzle.github.io>.
- **Papers With Code** (*Learn Python 3 in Y minutes*⁴) <https://paperswithcode.com/state-of-the-art>.
- **Google Dataset Search** (Blog⁵) <https://datasetsearch.research.google.com>.

"Dataset Search has indexed almost 25 million of these datasets, giving you a single place to search for datasets and find links to where the data is." — Natasha Noy

The Measure of Intelligence (*Abstraction and Reasoning Corpus*⁶) <https://arxiv.org/abs/1911.01547>.

❖ Growing Neural Cellular Automata, Mordvintsev et al. <https://distill.pub/2020/growing-ca/>.

2.1 In the Cloud

Colab⁷. Practice Immediately⁸. Labs⁹: Introduction to Deep Learning (MIT 6.S191)

- Free GPU compute via Colab <https://colab.research.google.com/notebooks/welcome.ipynb>.
- Colab can open notebooks directly from GitHub by simply replacing "*http://github.com*" with "*http://colab.research.google.com/github/*" in the notebook URL.
- Colab Pro <https://colab.research.google.com/signup>.

2.2 On a Local Machine

JupyterLab is an interactive development environment for working with notebooks, code and data¹⁰.

- Install Anaconda <https://www.anaconda.com/download/> and launch ‘Anaconda Navigator’
- Update Jupyterlab and launch the application. Under Notebook, click on ‘Python 3’

"If we truly reach AI, it will let us know." — Garry Kasparov

3 Deep Learning

After the **Historical AI Debate**¹¹: “*Yoshua Bengio and Gary Marcus on the Best Way Forward for AI*” <https://montrealartificialintelligence.com/aidebate/>, there have been clarifications on the term “**deep learning**”¹².

“Deep learning is inspired by neural networks of the brain to build learning machines which discover rich and useful internal representations, computed as a composition of learned features and functions.” — Yoshua Bengio

“DL is constructing networks of parameterized functional modules and training them from examples using gradient-based optimization.” — Yann LeCun

“... replace symbols by vectors and logic by continuous (or differentiable) functions.” — Yann LeCun

Deep learning allows computational models that are composed of multiple processing layers to learn REPRESENTATIONS of (raw) data with multiple levels of abstraction[2]. At a high-level, neural networks are either encoders, decoders, or a combination of both¹³. Introductory course <http://introtodeeplearning.com>. See also Table 1.

⁴<https://learnxinyminutes.com/docs/python3/>

⁵<https://blog.google/products/search/discovering-millions-datasets-web/>

⁶<https://github.com/fchollet/ARC>

⁷<https://medium.com/tensorflow/colab-an-easy-way-to-learn-and-use-tensorflow-d74d1686e309>

⁸<https://colab.research.google.com/github/madewithml/practicalAI/>

⁹https://colab.research.google.com/github/aamini/introtodeeplearning_labs

¹⁰<https://blog.jupyter.org/jupyterlab-is-ready-for-users-5a6f039b8906>

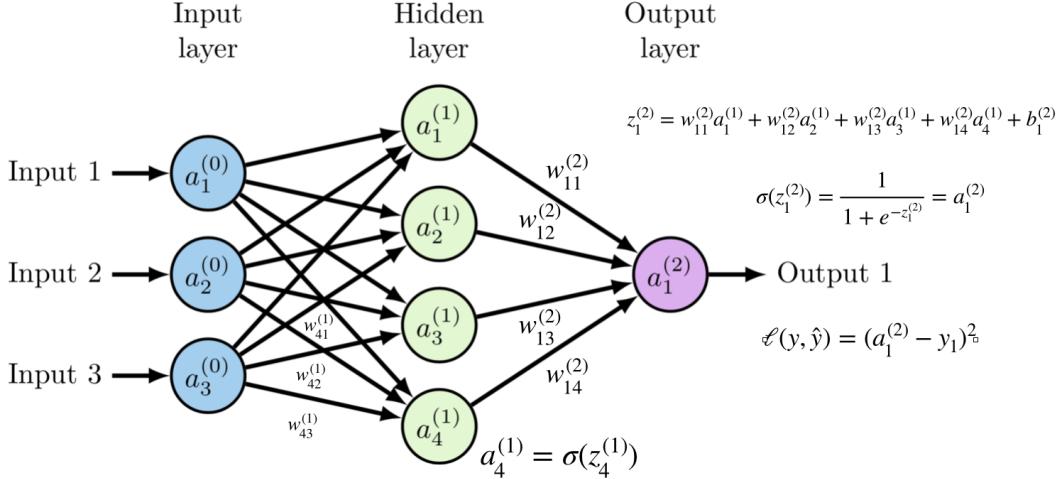
¹¹<https://www.zdnet.com/article/devils-in-the-details-in-bengio-marcus-ai-debate/>

¹²<https://www.zdnet.com/article/whats-in-a-name-the-deep-learning-debate/>

¹³<https://github.com/lexfridman/mit-deep-learning>

Table 1: Types of Learning, by Alex Graves at NeurIPS 2018

Name	With Teacher	Without Teacher
Active	<i>Reinforcement Learning / Active Learning</i>	<i>Intrinsic Motivation / Exploration</i>
Passive	<i>Supervised Learning</i>	<i>Unsupervised Learning</i>

Figure 1: **Multilayer perceptron (MLP).**

Deep learning (*distributed representations + composition*) is a general-purpose learning procedure.

"When you first study a field, it seems like you have to memorize a zillion things. You don't. What you need is to identify the 3-5 core principles that govern the field. The million things you thought you had to memorize are various combinations of the core principles." — J. Reed

"1. Multiply things together
 2. Add them up
 3. Replaces negatives with zeros
 4. Return to step 1, a hundred times."
 — Jeremy Howard

- ❖ Linear Algebra. Prof. Gilbert Strang¹⁴.
- ❖ Dive into Deep Learning <http://d2l.ai>.
- ❖ Minicourse in Deep Learning with PyTorch¹⁵.
- ❖ Introduction to Artificial Intelligence, Gilles Louppe¹⁶.
- ❖ Deep Learning. The full deck of (600+) slides, Gilles Louppe¹⁷.
- ❖ These Lyrics Do Not Exist <https://theselyricsdonotexist.com>.
- ❖ Backward Feature Correction: How Deep Learning Performs Deep Learning¹⁸.
- ❖ A Selective Overview of Deep Learning <https://arxiv.org/abs/1904.05526>.
- ❖ The Missing Semester of Your CS Education <https://missing.csail.mit.edu>.
- ❖ fastai: A Layered API for Deep Learning <https://arxiv.org/abs/2002.04688>.
- ❖ Anatomy of Matplotlib <https://github.com/matplotlib/AnatomyOfMatplotlib>.
- ❖ Data project checklist <https://www.fast.ai/2020/01/07/data-questionnaire/>.
- ❖ Using Nucleus and TensorFlow for DNA Sequencing Error Correction, Colab Notebook¹⁹.

¹⁴<https://ocw.mit.edu/courses/mathematics/18-06-linear-algebra-spring-2010/video-lectures/>

¹⁵<https://github.com/Atcold/pytorch-Deep-Learning-Minicourse>

¹⁶<https://glouppe.github.io/info8006-introduction-to-ai/pdf/lec-all.pdf>

¹⁷<https://github.com/glouppe/info8010-deep-learning/raw/v2-info8010-2019/pdf/lec-all.pdf>

¹⁸<https://arxiv.org/abs/2001.04413>

¹⁹https://colab.research.google.com/github/google/nucleus/blob/master/nucleus/examples/dna_sequencing_error_correction.ipynb

- ❖ PoseNet Sketchbook <https://googlecreativelab.github.io/posenet-sketchbook/>.
- ❖ Removing people from complex backgrounds in real time using TensorFlow.js in the web browser²⁰.
- ❖ A Recipe for Training Neural Networks <https://karpathy.github.io/2019/04/25/recipe/>.
- ❖ TensorFlow Datasets: load a variety of public datasets into TensorFlow programs (Blog²¹ | Colab²²).
- ❖ The Markov-Chain Monte Carlo Interactive Gallery <https://chi-feng.github.io/mcmc-demo/>.
- ❖ NeurIPS 2019 Implementations <https://paperswithcode.com/conference/neurips-2019-12>.
- ❖ Algebra, Topology, Differential Calculus, and Optimization Theory For Computer Science and Machine Learning²³.
- ❖ How to Choose Your First AI Project <https://hbr.org/2019/02/how-to-choose-your-first-ai-project>.
- ❖ Blog | MIT 6.S191 <https://medium.com/tensorflow/mit-introduction-to-deep-learning-4a6f8dde1f0c>.

3.1 Universal Approximation Theorem

The universal approximation theorem states that a feed-forward network with a single hidden layer containing a finite number of neurons can solve any given problem to arbitrarily close accuracy as long as you add enough parameters.

Neural Networks + Gradient Descent + GPU²⁴:

- Infinitely flexible function: *Neural Network* (multiple hidden layers: Deep Learning)²⁵.
- All-purpose parameter fitting: *Backpropagation*²⁶²⁷. Backpropagation is the key algorithm that makes training deep models computationally tractable and highly efficient²⁸. The backpropagation procedure is nothing more than a practical application of the chain rule for derivatives.

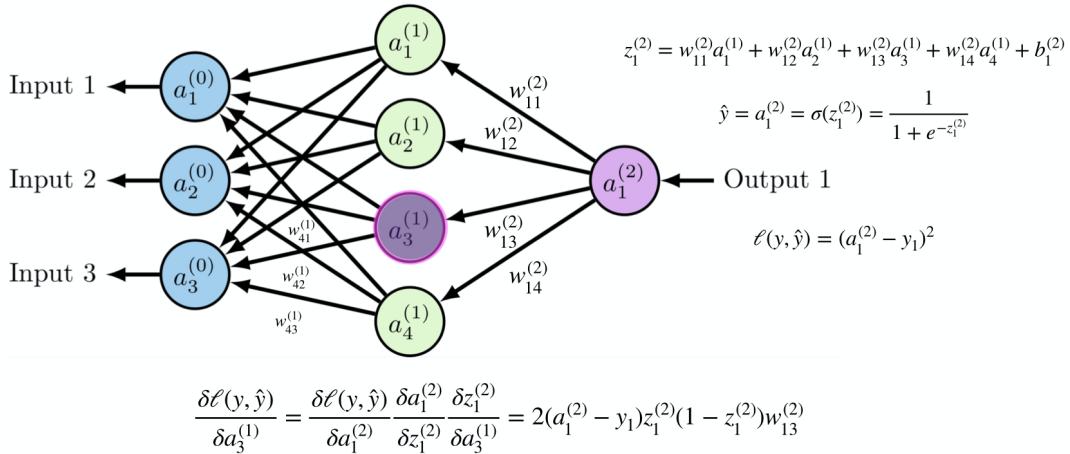


Figure 2: All-purpose parameter fitting: Backpropagation.

- Fast and scalable: *GPU*.

"You have relatively simple processing elements that are very loosely models of neurons. They have connections coming in, each connection has a weight on it, and that weight can be changed through learning." — Geoffrey Hinton

"Deep learning has fully solved the curse of dimensionality. It vanished like an RNN gradient!" — Ilya Sutskever

When a choice must be made, just feed the (raw) data to a deep neural network (Universal function approximators).

²⁰<https://github.com/jasonmayes/Real-Time-Person-Removal>

²¹<https://medium.com/tensorflow/introducing-tensorflow-datasets-c7f01f7e19f3>

²²<https://colab.research.google.com/github/tensorflow/datasets/blob/master/docs/overview.ipynb>

²³<https://drive.google.com/file/d/1sJvLQwxMyu89t2z4Zf9tD707efnbIUyB/view>

²⁴http://wiki.fast.ai/index.php/Lesson_1_Notes

²⁵<http://neuralnetworksanddeeplearning.com/chap4.html>

²⁶https://github.com/DebPanigrahi/Machine-Learning/blob/master/back_prop.ipynb

²⁷<https://www.jeremyjordan.me/neural-networks-training/>

²⁸<https://colah.github.io/posts/2015-08-Backprop/>

3.2 Convolution Neural Networks (Useful for Images | Space)

Richer innate priors : innateness that enables learning.

A significant percentage of Deep Learning breakthroughs comes from reusable constructs and parameters sharing. The deep convolutional network is a construct that reuses weights in multiple locations (parameters sharing in space)²⁹.

"Virtually all modern observers would concede that genes and experience work together; it is "nature and nurture", not "nature versus nurture". No nativist, for instance, would doubt that we are also born with specific biological machinery that allows us to learn. Chomsky's Language Acquisition Device should be viewed precisely as an innate learning mechanism, and nativists such as Pinker, Peter Marler (Marler, 2004) and myself (Marcus, 2004) have frequently argued for a view in which a significant part of a creature's innate armamentarium consists not of specific knowledge but of learning mechanisms, a form of innateness that enables learning." — Gary Marcus, Innateness, AlphaZero, and Artificial Intelligence³⁰

The deep convolutional network, inspired by Hubel and Wiesel's seminal work on early visual cortex, uses hierarchical layers of tiled convolutional filters to mimic the effects of receptive fields, thereby exploiting the local spatial correlations present in images[1]. See Figure 4. Demo <https://ml4a.github.io/demos/convolution/>.

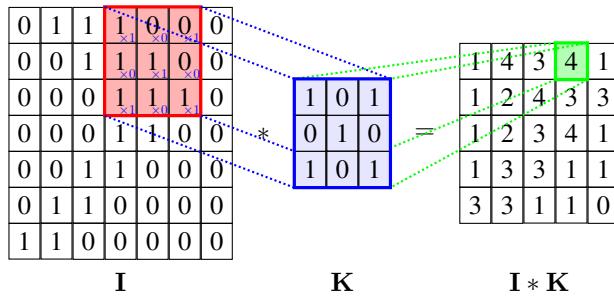


Figure 3: **2D Convolution**. Source: Cambridge Coding Academy

A ConvNet is made up of Layers. Every Layer has a simple API: It transforms an input 3D volume to an output 3D volume with some differentiable function that may or may not have parameters³¹. Reading³².

In images, local combinations of edges form motifs, motifs assemble into parts, and parts form objects³³³⁴.

Representation learning : the language of neural networks. The visual vocabulary of a convolutional neural network seems to emerge from low level features such as edges and orientations, and builds up textures, patterns and composites, ... and builds up even further into complete objects. This relates to Wittgenstein's "language-game" in **Philosophical Investigations**³⁵, where a functional language emerge from simple tasks before defining a vocabulary³⁶.

"DL is essentially a new style of programming – "differentiable programming" – and the field is trying to work out the reusable constructs in this style. We have some: convolution, pooling, LSTM, GAN, VAE, memory units, routing units, etc." — Thomas G. Dietterich

- ❖ CS231N : Convolutional Neural Networks for Visual Recognition³⁷.
- ❖ Deep Plastic Surgery: Robust and Controllable Image Editing with Human-Drawn Sketches. Yang et al.³⁸.
- ❖ TensorSpace (<https://tensorspace.org>) offers interactive 3D visualizations of *LeNet*, *AlexNet* and *Inceptionv3*.

²⁹<https://twitter.com/iamtrask/status/949439556499230720>

³⁰<https://arxiv.org/abs/1801.05667>

³¹<http://cs231n.github.io/convolutional-networks/>

³²<https://ml4a.github.io/ml4a/convnets/>

³³<http://yosinski.com/deepvis>

³⁴<https://distill.pub/2017/feature-visualization/>

³⁵https://en.wikipedia.org/wiki/Philosophical_Investigations

³⁶https://media.neurips.cc/Conferences/NIPS2018/Slides/Deep_Unsupervised_Learning.pdf

³⁷https://www.youtube.com/playlist?list=PLzUTmXVwsnXod6WNdg57Yc3zFx_f-RYsq

³⁸<https://arxiv.org/abs/2001.02890>

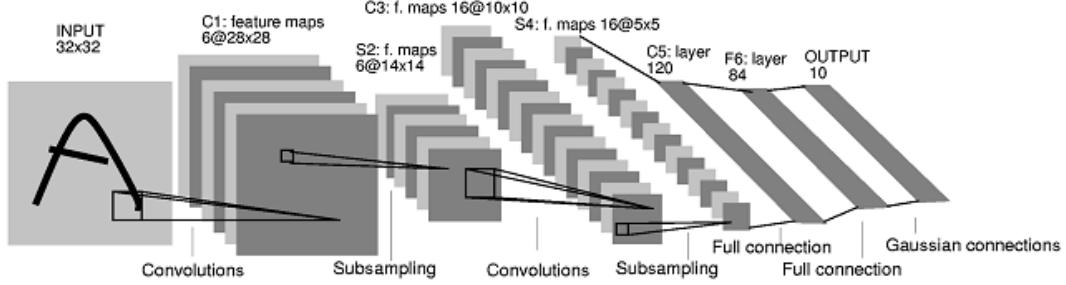


Figure 4: Architecture of LeNet-5, a Convolutional Neural Network. LeCun et al., 1998

3.3 Recurrent Neural Networks (Useful for Sequences | Time)

Recurrent neural networks are networks with loops in them, allowing information to persist³⁹. RNNs process an input sequence one element at a time, maintaining in their hidden units a ‘state vector’ that implicitly contains information about the history of all the past elements of the sequence[2]. For sequential inputs. See Figure 5.

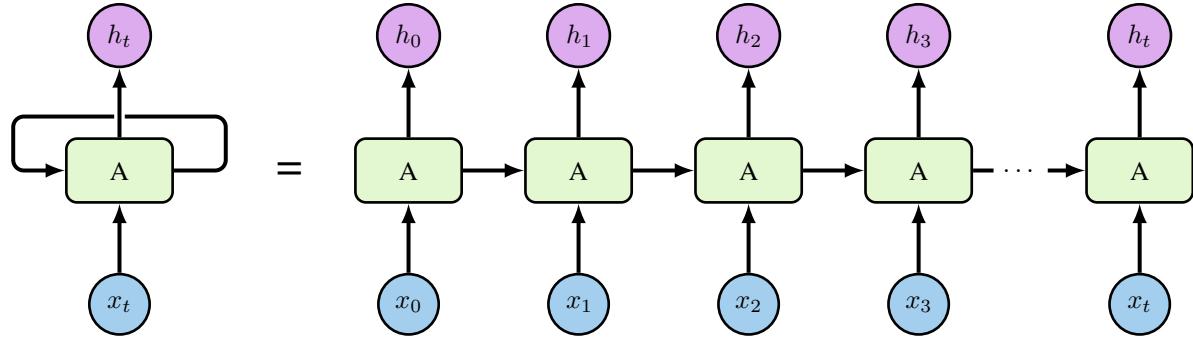


Figure 5: RNN Layers Reuse Weights for Multiple Timesteps.

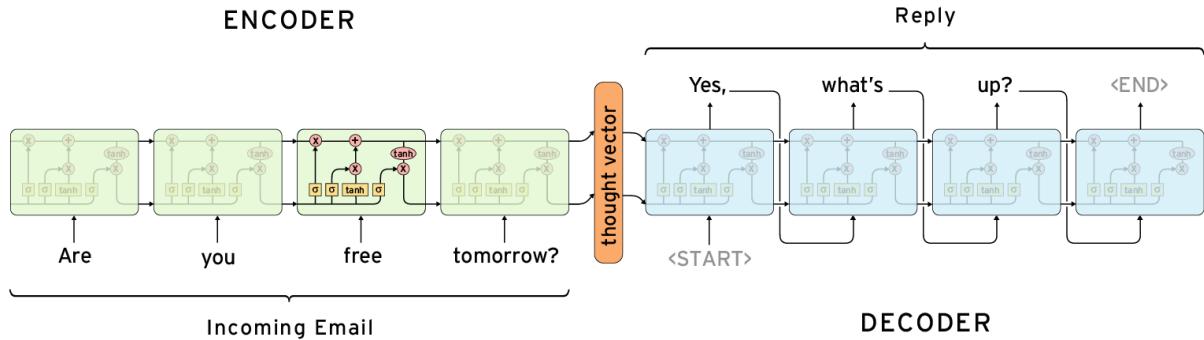


Figure 6: Google Smart Reply System is built on a pair of recurrent neural networks. Diagram by Chris Olah

"I feel like a significant percentage of Deep Learning breakthroughs ask the question “how can I reuse weights in multiple places?” – Recurrent (LSTM) layers reuse for multiple timesteps – Convolutional layers reuse in multiple locations. – Capsules reuse across orientation." — Andrew Trask

- ❖ CS224N : Natural Language Processing with Deep Learning⁴⁰.
- ❖ Long Short-Term-Memory (LSTM), Sepp Hochreiter and Jürgen Schmidhuber⁴¹.

³⁹<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

⁴⁰https://www.youtube.com/playlist?list=PLU40WL80194IJzQtileLTqGZuXtG1LMP_

⁴¹<https://www.bioinf.jku.at/publications/older/2604.pdf>

- ❖ The Unreasonable Effectiveness of Recurrent Neural Networks, blog (2015) by Andrej Karpathy⁴².
- ❖ Understanding LSTM Networks <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>.
- ❖ Can Neural Networks Remember? Slides by Vishal Gupta: http://vishalgupta.me/deck/char_lstms/.

3.4 Transformers

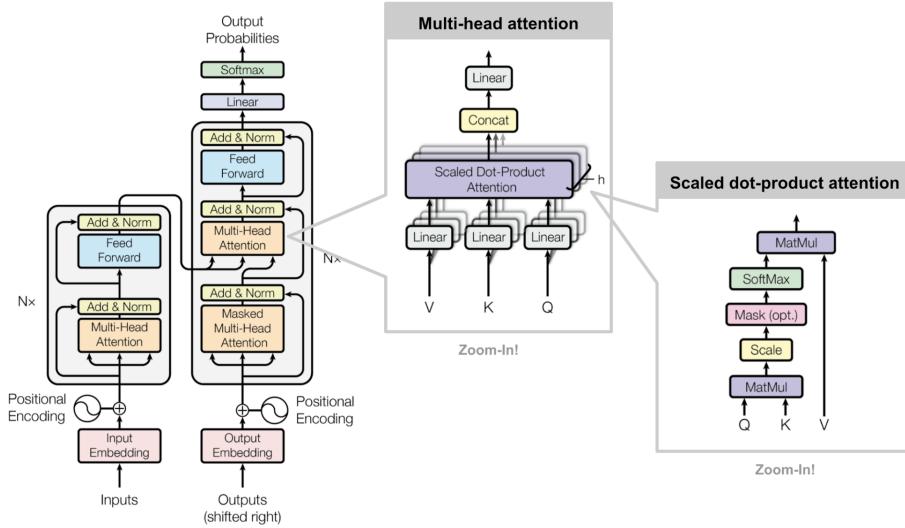


Figure 7: Attention Is All You Need. Vaswani et al., 2017 : <https://arxiv.org/abs/1706.03762>.

Transformers are generic, simple and exciting machine learning architectures designed to process a connected set of units (tokens in a sequence, pixels in an image, etc.) where the only interaction between units is through self-attention. Transformers' performance limit seems purely in the hardware (how big a model can be fitted in GPU memory)⁴³.

The fundamental operation of transformers is **self-attention** (a sequence-to-sequence operation, Figure 8): *an attention mechanism relating different positions of a single sequence in order to compute a representation of the same sequence*⁴⁴.

Let's call the input vectors (of dimension k) :

$$\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_t \quad (1)$$

Let's call the corresponding output vectors (of dimension k) :

$$\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_t \quad (2)$$

The **self attention** operation takes a weighted average over all the input vectors :

$$\mathbf{y}_i = \sum_j w_{ij} \mathbf{x}_j \quad (3)$$

The weight w_{ij} is derived from a function over \mathbf{x}_i and \mathbf{x}_j . The simplest option is the dot product (with softmax) :

$$w_{ij} = \frac{e^{\mathbf{x}_i^T \mathbf{x}_j}}{\sum_j e^{\mathbf{x}_i^T \mathbf{x}_j}} \quad (4)$$

Transformers are Graph Neural Networks⁴⁵.

⁴²<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

⁴³<http://www.peterbloem.nl/blog/transformers>

⁴⁴<https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>

⁴⁵<https://graphdeeplearning.github.io/post/transformers-are-gnns/>

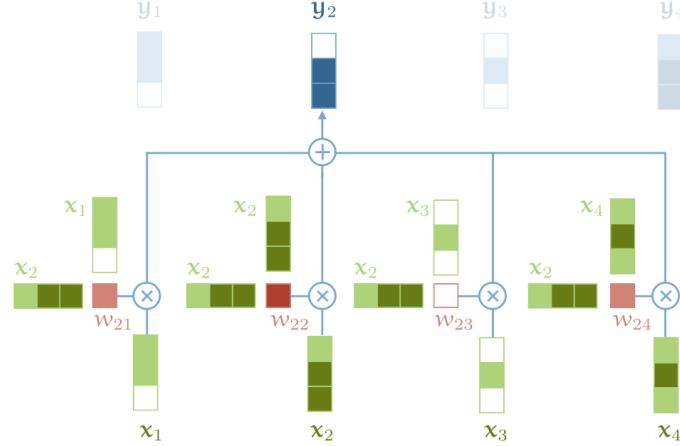


Figure 8: Self-attention. By Peter Bloem : <http://www.peterbloem.nl/blog/transformers>.

- ❖ Making Transformer networks simpler and more efficient⁴⁶.
- ❖ AttentioNN: All about attention in neural networks described as colab notebooks⁴⁷.
- ❖ Attention Is All You Need, Vaswani et al. <https://arxiv.org/abs/1706.03762>.
- ❖ How to train a new language model from scratch using Transformers and Tokenizers⁴⁸.
- ❖ The Illustrated Transformer <http://jalammar.github.io/illustrated-transformer/>.
- ❖ The annotated transformer (code) <http://nlp.seas.harvard.edu/2018/04/03/attention.html>.
- ❖ Attention and Augmented Recurrent Neural Networks <https://distill.pub/2016/augmented-rnns/>.
- ❖ Transformer model for language understanding. Tutorial showing how to write Transformer in TensorFlow 2.0⁴⁹.
- ❖ Transformer in TensorFlow 2.0 (code) <https://www.tensorflow.org/beta/tutorials/text/transformer>.
- ❖ **Write With Transformer.** By Hugging Face: <https://transformer.huggingface.co>.

3.4.1 Natural Language Processing (NLP) | BERT: A New Era in NLP

BERT (Bidirectional Encoder Representations from Transformers)[6] is a *deeply bidirectional, unsupervised language representation*, pre-trained using only a plain text corpus (in this case, Wikipedia)⁵⁰.

- Reading: Unsupervised pre-training of an LSTM followed by supervised fine-tuning[7].
- TensorFlow code and pre-trained models for BERT <https://github.com/google-research/bert>.
- Better Language Models and Their Implications⁵¹.

"I think transfer learning is the key to general intelligence. And I think the key to doing transfer learning will be the acquisition of conceptual knowledge that is abstracted away from perceptual details of where you learned it from." — Demis Hassabis

- ❖ Towards a Conversational Agent that Can Chat About... Anything⁵².
- ❖ How to Build OpenAI's GPT-2: "The AI That's Too Dangerous to Release"⁵³.
- ❖ A Primer in BERTology: What we know about how BERT works, Rogers et al., 2020⁵⁴.
- ❖ Play with BERT with your own data using TensorFlow Hub https://colab.research.google.com/github/google-research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb.

⁴⁶<https://ai.facebook.com/blog/making-transformer-networks-simpler-and-more-efficient/>

⁴⁷<https://github.com/zaidalyafeai/AttentioNN>

⁴⁸<https://huggingface.co/blog/how-to-train>

⁴⁹<https://www.tensorflow.org/tutorials/text/transformer>

⁵⁰<https://ai.googleblog.com/2018/11/open-sourcing-bert-state-of-art-pre.html>

⁵¹<https://blog.openai.com/better-language-models/>

⁵²<https://ai.googleblog.com/2020/01/towards-conversational-agent-that-can.html>

⁵³<https://blog.floydhub.com/gpt2/>

⁵⁴<https://arxiv.org/abs/2002.12327>

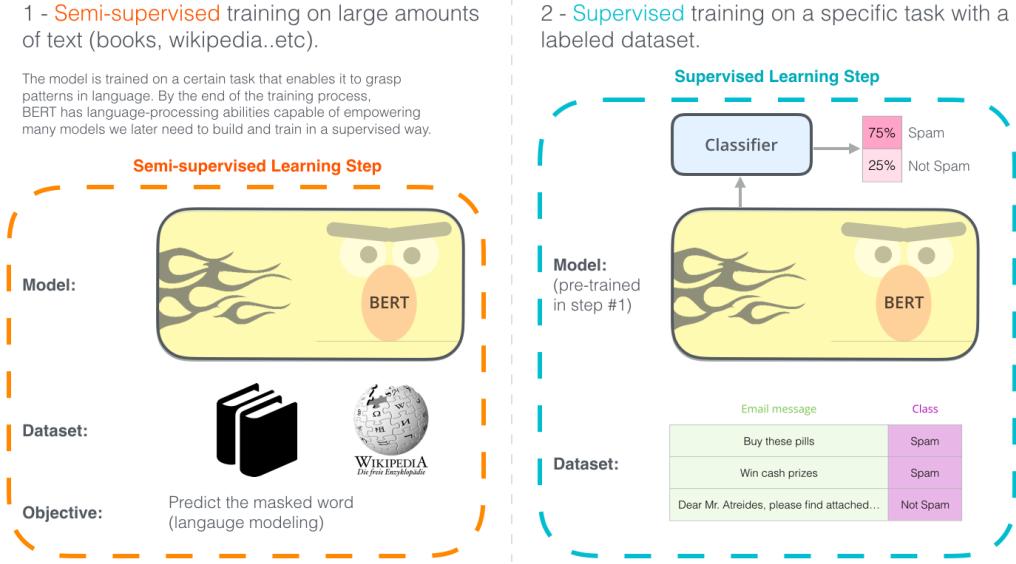


Figure 9: The two steps of how BERT is developed. Source <https://jalammar.github.io/illustrated-bert/>.

3.5 Unsupervised Learning

True intelligence will require independent learning strategies.

"Give a robot a label and you feed it for a second; teach a robot to label and you feed it for a lifetime." — Pierre Sermanet

Unsupervised learning is a paradigm for creating AI that learns without a particular task in mind: learning for the sake of learning⁵⁵. It captures some characteristics of the joint distribution of the observed random variables (learn the underlying structure). The variety of tasks include density estimation, dimensionality reduction, and clustering.[4]⁵⁶.

"The unsupervised revolution is taking off!" — Alfredo Canziani

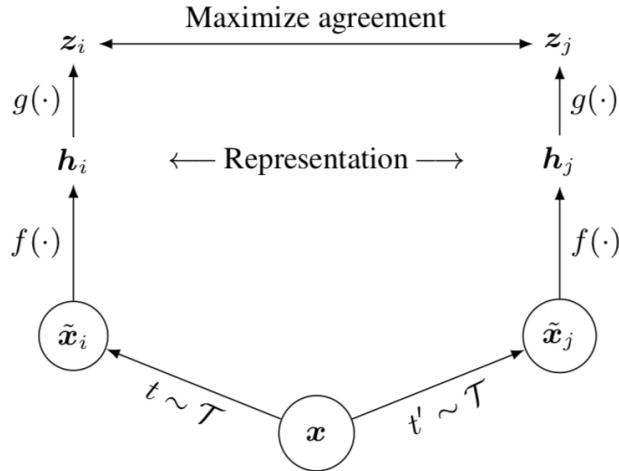


Figure 10: A Simple Framework for Contrastive Learning of Visual Representations, Chen et al., 2020

⁵⁵<https://deepmind.com/blog/unsupervised-learning/>

⁵⁶https://media.neurips.cc/Conferences/NIPS2018/Slides/Deep_Unsupervised_Learning.pdf

Self-supervised learning is derived from unsupervised learning where the data provides the supervision. E.g. Word2vec⁵⁷, a technique for learning vector representations of words, or word **embeddings**. An embedding is a mapping from discrete objects, such as words, to vectors of real numbers⁵⁸.

"The next revolution of AI won't be supervised." — Yann LeCun

- ❖ Self-Supervised Image Classification, Papers With Code⁵⁹.
- ❖ Self-supervised learning and computer vision, Jeremy Howard⁶⁰.
- ❖ Momentum Contrast for Unsupervised Visual Representation Learning, He et al.⁶¹
- ❖ Data-Efficient Image Recognition with Contrastive Predictive Coding, Hénaff et al.⁶²
- ❖ A Simple Framework for Contrastive Learning of Visual Representations, Chen et al.⁶³
- ❖ FixMatch: Simplifying Semi-Supervised Learning with Consistency and Confidence, Sohn et al.⁶⁴
- ❖ Self-Supervised Learning of Pretext-Invariant Representations, Ishan Misra, Laurens van der Maaten⁶⁵.

3.5.1 Generative Adversarial Networks

Simultaneously train two models: a generative model G that captures the data distribution, and a discriminative model D that estimates the probability that a sample came from the training data rather than G. The training procedure for G is to maximize the probability of D making a mistake. This framework corresponds to a minimax two-player game[3].

$$\min_{\theta_g} \max_{\theta_d} [\mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D_{\theta_d}(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D_{\theta_d}(G_{\theta_g}(z)))]] \quad (5)$$

"What I cannot create, I do not understand." — Richard Feynman

Goodfellow et al. used an interesting analogy where the generative model can be thought of as analogous to a team of counterfeiters, trying to produce fake currency and use it without detection, while the discriminative model is analogous to the police, trying to detect the counterfeit currency. Competition in this game drives both teams to improve their methods until the counterfeits are indistinguishable from the genuine articles. See Figure 11.

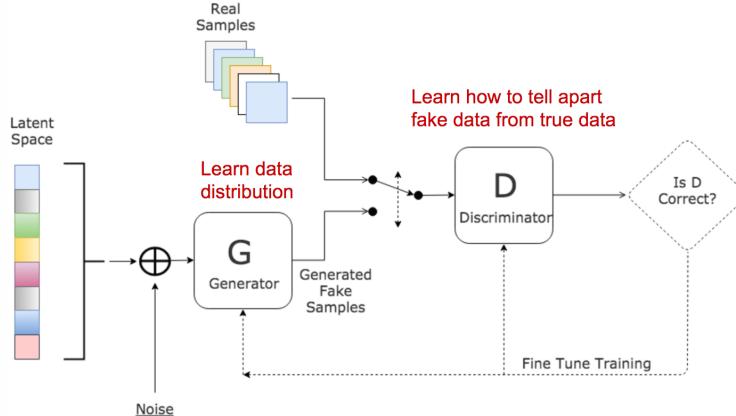


Figure 11: GAN: Neural Networks Architecture Pioneered by Ian Goodfellow at University of Montreal (2014).

StyleGAN: A Style-Based Generator Architecture for Generative Adversarial Networks

- Paper <http://stylegan.xyz/paper> | Code <https://github.com/NVlabs/stylegan>.

⁵⁷<https://jalamar.github.io/illustrated-word2vec/>

⁵⁸<http://projector.tensorflow.org>

⁵⁹<https://paperswithcode.com/task/self-supervised-image-classification>

⁶⁰https://www.fast.ai/2020/01/13/self_supervised/

⁶¹<https://arxiv.org/abs/1911.05722>

⁶²<https://arxiv.org/abs/1905.09272>

⁶³<https://arxiv.org/abs/2002.05709>

⁶⁴<https://arxiv.org/abs/2001.07685>

⁶⁵<https://arxiv.org/abs/1912.01991>

- **StyleGAN for art.** Colab <https://colab.research.google.com/github/ak9250/stylegan-art>.
 - This Person Does Not Exist <https://thispersondoesnotexist.com>.
 - Which Person Is Real? <http://www.whichfaceisreal.com>.
 - This Resume Does Not Exist <https://thisresumedoestnotexist.com>.
 - This Waifu Does Not Exist <https://www.thiswaifudoesnotexist.net>.
 - Encoder for Official TensorFlow Implementation <https://github.com/Puzer/stylegan-encoder>.
 - How to recognize fake AI-generated images. By Kyle McDonald⁶⁶.
- ❖ 100,000 Faces Imagined by a GAN <https://generated.photos>.
- ❖ Introducing **TF-GAN**: A lightweight GAN library for TensorFlow 2.0⁶⁷.
- ❖ Generative Adversarial Networks (GANs) in 50 lines of code (PyTorch)⁶⁸.
- ❖ Few-Shot Adversarial Learning of Realistic Neural Talking Head Models⁶⁹.
- ❖ Wasserstein GAN <http://www.depthfirstlearning.com/2019/WassersteinGAN>.
- ❖ GANpaint Paint with GAN units <http://gandissect.res.ibm.com/ganpaint.html>.
- ❖ A Review on Generative Adversarial Networks: Algorithms, Theory, and Applications. Gui et al.⁷⁰.
- ❖ CariGANs: Unpaired Photo-to-Caricature Translation. Cao et al.: <https://cari-gan.github.io>.
- ❖ Infinite-resolution (CPPNs, GANs and TensorFlow.js) <https://thispicturdoesnotexist.com>.
- ❖ PyTorch pretrained BigGAN <https://github.com/huggingface/pytorch-pretrained-BigGAN>.
- ❖ GANSynth: Generate high-fidelity audio with GANs! Colab <http://goo.gl/magenta/gansynth-demo>.
- ❖ SC-FEGAN: Face Editing Generative Adversarial Network <https://github.com/JoYoungjoo/SC-FEGAN>.
- ❖ Demo of BigGAN in an official Colaboratory notebook (backed by a GPU) https://colab.research.google.com/github/tensorflow/hub/blob/master/examples/colab/biggan_generation_with_tf_hub.ipynb

3.5.2 Variational AutoEncoder

Variational Auto-Encoders⁷¹ (VAEs) are powerful models for learning low-dimensional representations See Figure 12. Disentangled representations are defined as ones where a change in a single unit of the representation corresponds to a change in single factor of variation of the data while being invariant to others (Bengio et al. (2013)).

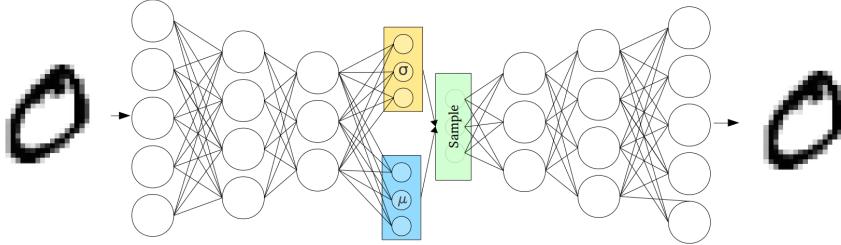


Figure 12: Variational Autoencoders (VAEs): Powerful Generative Models.

- ❖ Colab⁷²: "Debiasing Facial Detection Systems." **AIEthics**
- ❖ Reading: Disentangled VAE's (DeepMind 2016) <https://arxiv.org/abs/1606.05579>.
- ❖ Slides: A Few Unusual Autoencoders <https://colinraffel.com/talks/vector2018few.pdf>.
- ❖ **MusicVAE**: Learning latent spaces for musical scores <https://magenta.tensorflow.org/music-vae>.
- ❖ Generative models in **Tensorflow 2** <https://github.com/timsainb/tensorflow2-generative-models/>.
- ❖ **SpaceSheet**: Interactive Latent Space Exploration with a Spreadsheet <https://vusd.github.io/spacesheet/>.

⁶⁶<https://medium.com/@kcimc/how-to-recognize-fake-ai-generated-images-4d1f6f9a2842>

⁶⁷<https://medium.com/tensorflow/introducing-tf-gan-a-lightweight-gan-library-for-tensorflow-2-0-36d767e1abae>

⁶⁸<https://medium.com/@devnag/generative-adversarial-networks-gans-in-50-lines-of-code-pytorch-e81b79659e3f>

⁶⁹<https://arxiv.org/abs/1905.08233>

⁷⁰<https://arxiv.org/abs/2001.06937>

⁷¹<https://arxiv.org/abs/1906.02691v2>

⁷²https://colab.research.google.com/github/aamini/introtodeeplearning_labs/blob/master/lab2/Part2_debiasing_solution.ipynb

3.5.3 Capsule

Stacked Capsule Autoencoders. The inductive biases in this unsupervised version of capsule networks give rise to object-centric latent representations, which are learned in a self-supervised way—simply by reconstructing input images. Clustering learned representations is enough to achieve unsupervised state-of-the-art classification performance on MNIST (98.5%). Reference: blog by Adam Kosiorek.⁷³ Code⁷⁴.

Capsules learn *equivariant object representations* (applying any transformation to the input of the function has the same effect as applying that transformation to the output of the function).

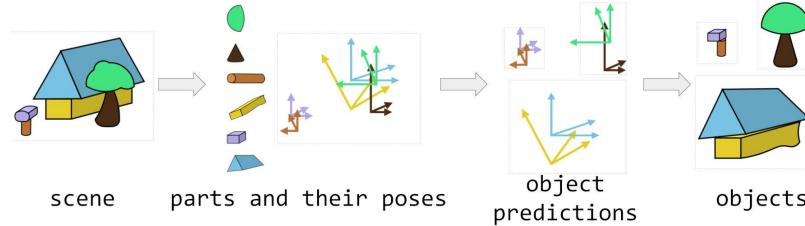


Figure 13: Stacked Capsule Autoencoders. Image source: Blog by Adam Kosiorek.

4 Autonomous Agents

We are on the dawn of *The Age of Artificial Intelligence*.

"In a moment of technological disruption, leadership matters." — Andrew Ng

An **autonomous agent** is any device that perceives its environment and takes actions that maximize its chance of success at some goal. At the bleeding edge of AI, autonomous agents can learn from experience, simulate worlds and orchestrate meta-solutions. Here's an informal definition⁷⁵ of the *universal intelligence* of agent π^{76} :

$$\Upsilon(\pi) := \sum_{\mu \in E} 2^{-K(\mu)} V_\mu^\pi \quad (6)$$

"Intelligence measures an agent's ability to achieve goals in a wide range of environments." — Shane Legg

4.1 Deep Reinforcement Learning

Reinforcement learning (RL) studies how an agent can learn how to achieve goals in a complex, uncertain environment (Figure 14) [5]. Recent superhuman results in many difficult environments combine deep learning with RL (*Deep Reinforcement Learning*). See Figure 15 for a taxonomy of RL algorithms.

- ❖ An Opinionated Guide to ML Research⁷⁷.
- ❖ CS 188 : Introduction to Artificial Intelligence⁷⁸.
- ❖ Introduction to Reinforcement Learning by DeepMind⁷⁹.
- ❖ "My Top 10 Deep RL Papers of 2019" by Robert Tjarko Lange⁸⁰.
- ❖ Deep tic-tac-toe <https://zackakil.github.io/deep-tic-tac-toe/>.
- ❖ CS 287: Advanced Robotics⁸¹. <https://people.eecs.berkeley.edu/~pabbeel/cs287-fa19/>.

⁷³http://akosiorek.github.io/ml/2019/06/23/stacked_capsule_autoencoders.html

⁷⁴https://github.com/google-research/google-research/tree/master/stacked_capsule_autoencoders

⁷⁵<https://arxiv.org/abs/0712.3329>

⁷⁶Where μ is an environment, K is the Kolmogorov complexity function, E is the space of all computable reward summable environmental measures with respect to the reference machine U and the value function V_μ^π is the agent's "ability to achieve".

⁷⁷<http://joschu.net/blog/opinionated-guide-ml-research.html>

⁷⁸<https://inst.eecs.berkeley.edu/~cs188/fa18/>

⁷⁹<https://www.youtube.com/watch?v=2pWv7G0vuf0&list=PLqYmG7hTraZDM-0YHWgPebj2MfCFzF0bQ>

⁸⁰<https://roberttlange.github.io/posts/2019/12/blog-post-9/>

⁸¹<https://people.eecs.berkeley.edu/~pabbeel/cs287-fa19/exam/cs287-fa19-exam-study-handout.pdf>

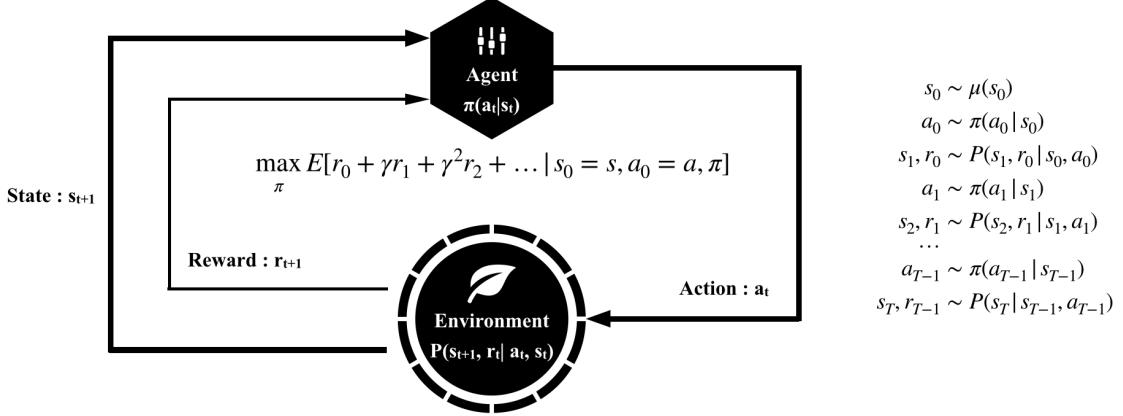


Figure 14: An Agent Interacts with an Environment.

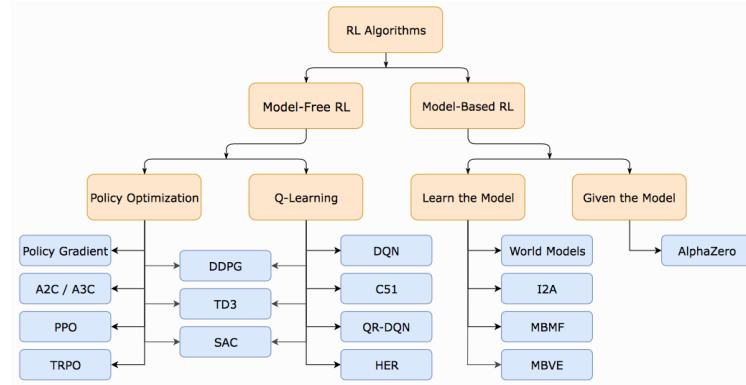


Figure 15: A Taxonomy of RL Algorithms. Source: Spinning Up in Deep RL by Achiam et al. | OpenAI

4.1.1 Model-Free RL | Value-Based

The goal in RL is to train the agent to maximize the discounted sum of all future rewards R_t , called the return:

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \quad (7)$$

The Q-function captures the expected total future reward an agent in state s can receive by executing a certain action a :

$$Q(s, a) = E[R_t] \quad (8)$$

The optimal policy should choose the action a that maximizes $Q(s, a)$:

$$\pi^*(s) = \operatorname{argmax}_a Q(s, a) \quad (9)$$

- **Q-Learning:** Playing Atari with Deep Reinforcement Learning (DQN). Mnih et al, 2013[10]. See Figure 17.

"There's no limit to intelligence." — David Silver

- ❖ Q-Learning in enormous action spaces via amortized approximate maximization, de Wiele et al.⁸².
- ❖ TF-Agents (DQN Tutorial) | Colab <https://colab.research.google.com/github/tensorflow/agents>.

⁸²<https://arxiv.org/abs/2001.08116>

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	AA
2	Project	Maintainer	Framework	Execution	DQN	Rainbow	REINFORCE	A2C	PPO	DDPG	SAC	TDS	REINFORCE	A2C	PPO		n-step return	prioritized experience replay	distributional value function approximation	hyperbolic discounting	dist observations support	?	X	Y	Z	AA	
3	OpenAI baseline	OpenAI	Tensorflow	✓	✗	✗	✗	✗	✓	✓	✓	✗	✗	✗	✓	✓	✗	✓	✗	✗	✗	✗	last commit June	8 Stars	8.1k	commit activity 1/month	code size 1.22 MB
4	stable_baselines	Ashley Ruffin,	Tensorflow	✓	✗	✗	✗	✗	✗	✓	✓	✓	✗	✗	✗	✓	✓	✓	✓	✗	✗	✗	last commit last monday	0 Stars	1.1k	commit activity 6/month	code size 893 kB
5	CatalystRL	Sergey Kolesnikov	PyTorch	✗?	✓	✓	✗	✗	✗	✓	✗	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	last commit today	0 Stars	718	commit activity 19/month	code size 411 kB
6	Ray.rllib	Ray Team	Tensorflow	✗?	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	?	last commit today	0 Stars	8k	commit activity 100/month	code size 4.92 MB
7	TF agents	Google	Tensorflow	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit yesterday	0 Stars	749	commit activity 34/month	code size 2.18 MB
8	Horizon	Facebook	PyTorch	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit today	0 Stars	2k	commit activity 28/month	code size 1.04 MB
9	Coach	Intel	Tensorflow	✗?	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit today	0 Stars	1.4k	commit activity 7/month	code size 1.99 MB
10	Garage	community	Tensorflow	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit yesterday	0 Stars	40k	commit activity 27/month	code size 1.54 MB
11	SLM-Lab	Wah Loon Kong, Laura Grasser	PyTorch	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit last monday	0 Stars	848	commit activity 147/month	code size 315 kB
12	Dopamine	Google	Tensorflow	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit July	0 Stars	5.1k	commit activity 3/month	code size 2.34 MB
13	OpenAI Gymnasium	OpenAI	Tensorflow	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit today	0 Stars	3.2k	commit activity 21/kb	code size 403 kB
14	tf1	DeepMind	Tensorflow	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit march	0 Stars	654	commit activity 0/month	code size 12.1 kB
15	scalable_agent	DeepMind	Tensorflow	?	✗	✗	✗	✗	✗	?	?	?	?	?	?	?	?	?	?	?	?	?	last commit march	0 Stars	1.9k	commit activity 0/month	code size 364 kB
16	ELF	Facebook	PyTorch	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	last commit July	0 Stars	4k	commit activity 0/month	code size 1.91 kB
17	keras-d	Matthias Piapert	Tensorflow	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit May	0 Stars	1.4k	commit activity 0/month	code size 95.9 kB
18	Iosifitkov	Ilya Kostrikov	PyTorch	✓	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit last saturday	0 Stars	777	commit activity 2/month	code size 30.4 kB
19	Rainbow	Kai Arulkumaran	PyTorch	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit June	0 Stars	23k	commit activity 0/month	code size 668 kB
20	Vel	Jerry (?)	PyTorch	✓	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit march	0 Stars	2.4k	commit activity 0/month	code size 870 kB
21	tensorforce		Tensorflow	✓	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit April 2016	0 Stars	1.6k	commit activity 0/month	code size 1.07 MB
22	RL-Adventure		PyTorch	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit march	0 Stars	416	commit activity 0/month	code size 4.15 kB
23	DeepRL-Tutorials		PyTorch	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit April	0 Stars	342	commit activity 0/month	code size 467 kB	
24	surreal		PyTorch	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	✗	?	last commit June	0 Stars	11k	commit activity 12/month	code size 2.14 MB	
25	lagent		Tensorflow	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	last commit June 2017	0 Stars	71	commit activity 0/month	code size 1.39 MB
26	dmcontrol		Tensorflow	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	last commit July	0 Stars	251	commit activity 2/month	code size 73.9 kB
27	atari2d		Tensorflow	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	last commit July	0 Stars	251	commit activity 2/month	code size 73.9 kB
28	gymar	WhiRL	PyTorch	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	?	last commit July	0 Stars	251	commit activity 2/month	code size 73.9 kB

Figure 16: Open-Source RL Algorithms https://docs.google.com/spreadsheets/d/1EeFPd-XIQ3mq_9snT1AZSsFY7Hbnmd7P5bbT8LPuMn0/

DQN Training Algorithm

Algorithm 1: deep Q-learning with experience replay.

```

Initialize replay memory  $D$  to capacity  $N$ 
Initialize action-value function  $Q$  with random weights  $\theta$ 
Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$ 
For episode = 1,  $M$  do
    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequence  $\phi_1 = \phi(s_1)$ 
    For  $t = 1, T$  do
        With probability  $\epsilon$  select a random action  $a_t$ 
        otherwise select  $a_t = \text{argmax}_a Q(\phi(s_t), a; \theta)$ 
        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$ 
        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$ 
        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $D$ 
        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $D$ 
        Set  $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$ 
        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$ 
        Every  $C$  steps reset  $\hat{Q} = Q$ 
    End For
End For

```



Figure 17: DQN Training Algorithm. Volodymyr Mnih, Deep RL Bootcamp

4.1.2 Model-Free RL | Policy-Based

An RL agent learns the stochastic policy function that maps state to action and act by sampling policy.

Run a policy for a while (code: <https://gist.github.com/karpathy/a4166c7fe253700972fcbe77e4ea32c5>):

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T) \quad (10)$$

Increase probability of actions that lead to high rewards and decrease probability of actions that lead to low rewards:

$$\nabla_\theta E_\tau[R(\tau)] = E_\tau \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi(a_t | s_t, \theta) R(\tau) \right] \quad (11)$$



Figure 18: Policy Gradient Directly Optimizes the Policy.

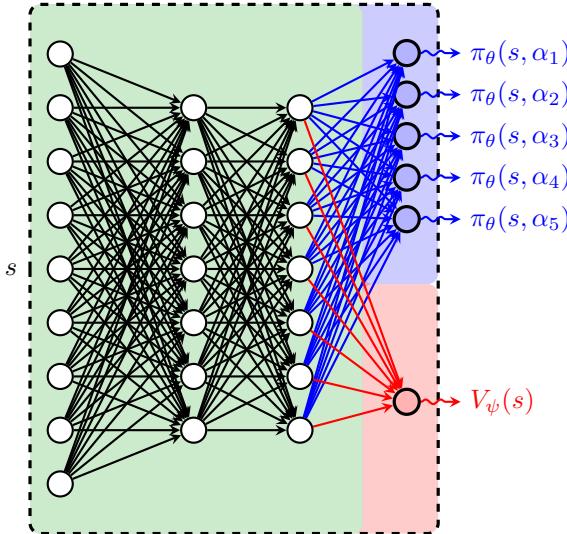


Figure 19: Asynchronous Advantage Actor-Critic (A3C). Source: Petar Velickovic

- **Policy Optimization:** Asynchronous Methods for Deep Reinforcement Learning (A3C). Mnih et al, 2016[8].
- **Policy Optimization:** Proximal Policy Optimization Algorithms (PPO). Schulman et al, 2017[9].

❖ Deep Reinforcement Learning for Playing 2.5D Fighting Games. Li et al.⁸³.

4.1.3 Model-Based RL

In Model-Based RL, the agent generates predictions about the next state and reward before choosing each action.

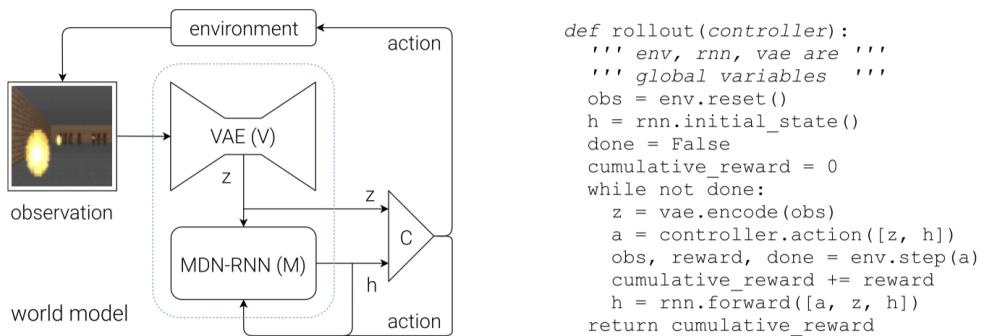


Figure 20: World Model's Agent consists of: Vision (V), Memory (M), and Controller (C). | Ha et al, 2018[11]

⁸³<https://arxiv.org/abs/1805.02070>

- **Learn the Model:** *Recurrent World Models Facilitate Policy Evolution* (World Models⁸⁴). The world model agent can be trained in an unsupervised manner to learn a compressed spatial and temporal representation of the environment. Then, a compact policy can be trained. See Figure 20. Ha et al, 2018[11].

- **Learn the Model:** *Learning Latent Dynamics for Planning from Pixels* <https://planetrl.github.io/>.

- **Given the Model:** *Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm* (AlphaZero). Silver et al, 2017[14]. AlphaGo Zero Explained In One Diagram⁸⁵.

❖ Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model. Schrittwieser et al.⁸⁶.

4.1.4 Toward a General AI-Agent Architecture: SuperDyna (*General Dyna-style RL Agent*)

"Intelligence is the computational part of the ability to predict and control a stream of experience." — Rich Sutton

SuperDyna.⁸⁷ The ambition: a general AI agent for Artificial Biological Reinforcement Learning.

1. Interact with the world: sense, update state and take an action
2. Learn from what just happened: see what happened and learn from it
3. Plan: (while there is time remaining in this time step) imagine hypothetical states and actions you might take
4. Discover : curate options and features and measure how well they're doing

Inner loop of a General Dyna-style RL Agent (SuperDyna)

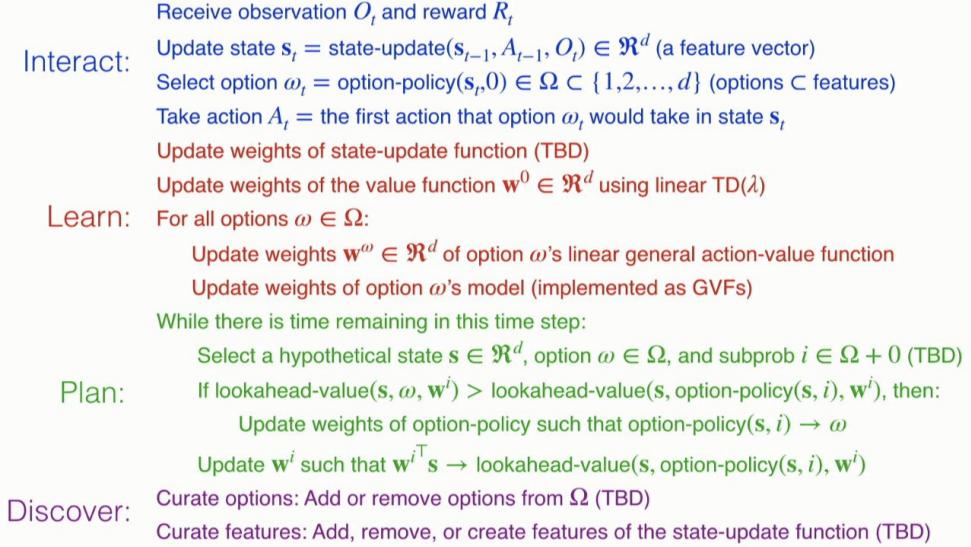


Figure 21: Inner Loop of a General Dyna-Style RL Agent (**SuperDyna**).

The first complete and scalable general AI-agent architecture that has all the most important capabilities and desiderata:

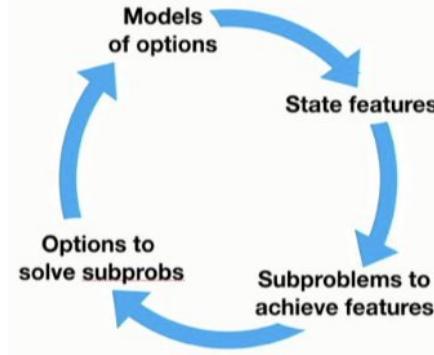
- Acting, learning, planning, model-learning, subproblems, and options.
- Function approximation, partial observability, non-stationarity and stochasticity.
- Discovery of state features, and thereby of subproblems, options and models.
- All feeding back to motivate new, more-abstract features in a virtuous cycle of discovery.

⁸⁴<https://worldmodels.github.io>

⁸⁵https://applied-data.science/static/main/res/alpha_go_zero_cheat_sheet.png

⁸⁶<https://arxiv.org/abs/1911.08265>

⁸⁷<https://insidehpc.com/2020/02/video-toward-a-general-ai-agent-architecture/>

Figure 22: **SuperDyna**: Virtuous cycle of discovery.

Presentation by Richard Sutton (starts at 15 min.)⁸⁸.

"In practice, I work primarily in reinforcement learning as an approach to artificial intelligence. I am exploring ways to represent a broad range of human knowledge in an empirical form—that is, in a form directly in terms of experience—and in ways of reducing the dependence on manual encoding of world state and knowledge." — Richard S. Sutton

4.1.5 Improving Agent Design

Via Reinforcement Learning: Blog⁸⁹. arXiv⁹⁰. ASTool <https://github.com/hardmaru/astool/>.

Via Evolution: Video⁹¹. Evolved Creatures <http://www.karlsims.com/evolved-virtual-creatures.html>.

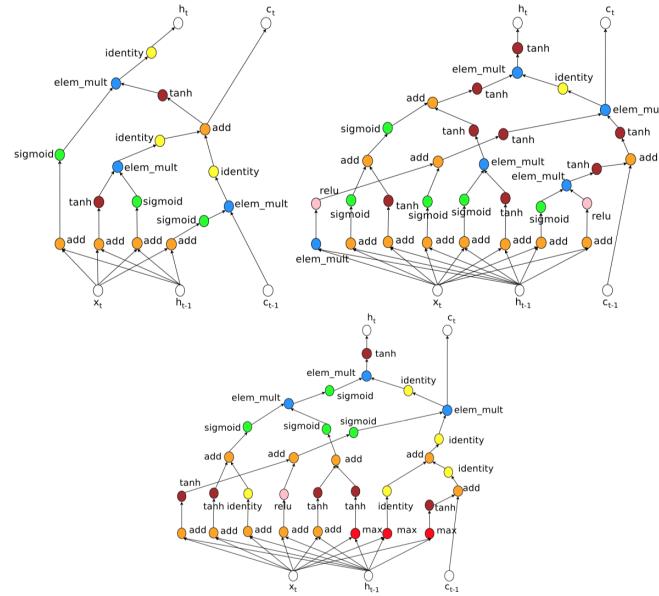


Figure 23: A comparison of the original LSTM cell vs. two new good generated. Top left: LSTM cell. [19]

"The future of high-level APIs for AI is... a problem-specification API. Currently we only search over network weights, thus "problem specification" involves specifying a model architecture. In the future, it will just be: "tell me what data you have and what you are optimizing"." — François Chollet

⁸⁸<https://slideslive.com/38921889/biological-and-artificial-reinforcement-learning-4>

⁸⁹<https://designrl.github.io>

⁹⁰<https://arxiv.org/abs/1810.03779>

⁹¹https://youtu.be/JBgG_VSP7f8

- ❖ Teacher algorithms for curriculum learning of Deep RL in continuously parameterized environments⁹².

4.1.6 OpenAI Baselines

High-quality implementations of reinforcement learning algorithms <https://github.com/openai/baselines>. Colab <https://colab.research.google.com/drive/1KKq9A3dRTq1q6bJmPyF0gg917gQyTjJI>.

4.1.7 Google Dopamine and A Zoo of Agents

Dopamine is a research framework for fast prototyping of reinforcement learning algorithms.⁹³.

A Zoo of Atari-Playing Agents: Code⁹⁴, Blog⁹⁵ and Colaboratory notebook <https://colab.research.google.com/github/uber-research/atari-model-zoo/blob/master/colab/AtariZooColabDemo.ipynb>.

4.1.8 TRFL : TensorFlow Reinforcement Learning

TRFL ("truffle"): a library of reinforcement learning building blocks <https://github.com/deepmind/trfl>.

4.1.9 bsuite : Behaviour Suite for Reinforcement Learning

A collection of experiments that investigate core capabilities of RL agents <http://github.com/deepmind/bsuite>.

4.2 Evolution Strategies (ES)

In her Nobel Prize in Chemistry 2018 Lecture "*Innovation by Evolution: Bringing New Chemistry to Life*" (Nobel Lecture)^{†96}, Prof. Frances H. Arnold said :

"Nature ... invented life that has flourished for billions of years. (...) Equally awe-inspiring is the process by which Nature created these enzyme catalysts and in fact everything else in the biological world. The process is evolution, the grand diversity-generating machine that created all life on earth, starting more than three billion years ago. (...) evolution executes a simple algorithm of diversification and natural selection, an algorithm that works at all levels of complexity from single protein molecules to whole ecosystems." — Prof. Frances H. Arnold

Evolution and neural networks proved a potent combination in nature.

"Evolution is a slow learning algorithm that with the sufficient amount of compute produces a human brain." — Wojciech Zaremba

Natural evolutionary strategy **directly evolves the weights of a DNN** and performs competitively with the best deep reinforcement learning algorithms, including deep Q-networks (DQN) and policy gradient methods (A3C)[21].

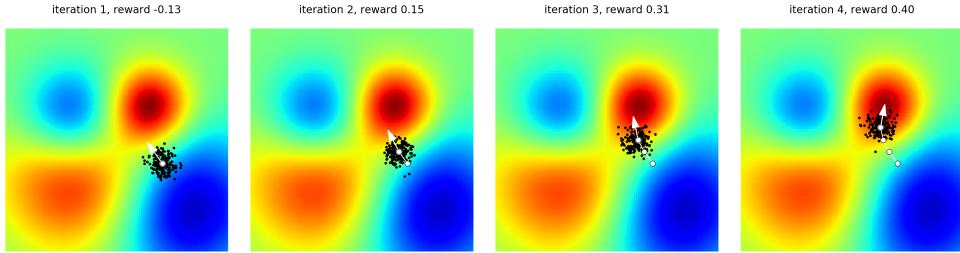


Figure 24: <https://colab.research.google.com/github/karpathy/randomfun/blob/master/es.ipynb>.

Neuroevolution, which harnesses evolutionary algorithms to optimize neural networks, enables capabilities that are typically unavailable to gradient-based approaches, including learning neural network building blocks, architectures and even the algorithms for learning[12].

⁹²<https://arxiv.org/abs/1910.07224>

⁹³<https://github.com/google/dopamine>

⁹⁴<https://github.com/uber-research/atari-model-zoo>

⁹⁵<https://eng.uber.com/atari-zoo-deep-reinforcement-learning/>

⁹⁶<https://onlinelibrary.wiley.com/doi/epdf/10.1002/anie.201907729>

"... evolution — whether biological or computational — is inherently creative, and should routinely be expected to surprise, delight, and even outwit us." — The Surprising Creativity of Digital Evolution, Lehman et al.[22]

The ES algorithm is a “guess and check” process, where we start with some random parameters and then repeatedly:

1. Tweak the guess a bit randomly, and
2. Move our guess slightly towards whatever tweaks worked better.

Neural architecture search has advanced to the point where it can outperform human-designed models[13].

"Caterpillar brains LIQUIFY during metamorphosis, but the butterfly retains the caterpillar's memories!" — M. Levin

"Open-ended" algorithms are algorithms that endlessly create. Brains and bodies evolve together in nature.

"We're machines," says Hinton. *"We're just produced biologically (...)"* — Katrina Onstad, Toronto Life

- ❖ Evolution Strategies⁹⁷.
- ❖ VAE+CPPN+GAN⁹⁸.
- ❖ Demos: ES on CartPole-v1⁹⁹ and ES on LunarLanderContinuous-v2¹⁰⁰.
- ❖ Spiders Can Fly Hundreds of Miles Riding the Earth's Magnetic Fields¹⁰¹.
- ❖ A Visual Guide to ES <http://blog.otoro.net/2017/10/29/visual-evolution-strategies/>.
- ❖ **Xenobots** A scalable pipeline for designing reconfigurable organisms, Kriegman et al.¹⁰². Learn¹⁰³. Evolve¹⁰⁴.

4.3 Self Play

Silver et al.[15] introduced an algorithm based solely on reinforcement learning, without human data, guidance or domain knowledge. Starting tabula rasa (and being its own teacher!), AlphaGo Zero achieved superhuman performance. AlphaGo Zero showed that algorithms matter much more than big data and massive amounts of computation.

"Self-Play is Automated Knowledge Creation." — Carlos E. Perez

Self-play mirrors similar insights from coevolution. Transfer learning is the key to go from self-play to the real world¹⁰⁵.

"Open-ended self play produces: Theory of mind, negotiation, social skills, empathy, real language understanding." — Ilya Sutskever, Meta Learning and Self Play

- ❖ How To Build Your Own MuZero AI Using Python¹⁰⁶.
- ❖ TensorFlow.js Implementation of DeepMind’s AlphaZero Algorithm for Chess. Live Demo¹⁰⁷. | Code¹⁰⁸.
- ❖ An open-source implementation of the AlphaGoZero algorithm <https://github.com/tensorflow/minigo>.
- ❖ ELF OpenGo: An Open Reimplementation of AlphaZero, Tian et al.: <https://arxiv.org/abs/1902.04522>.

4.4 Multi-Agent Populations

"We design a Theory of Mind neural network – a ToMnet – which uses meta-learning to build models of the agents it encounters, from observations of their behaviour alone." — Machine Theory of Mind, Rabinowitz et al.[25]

Cooperative Agents. Learning to Model Other Minds, by OpenAI[24], is an algorithm which accounts for the fact that other agents are learning too, and discovers self-interested yet collaborative strategies. Also: OpenAI Five¹⁰⁹.

⁹⁷<https://lilianweng.github.io/lil-log/2019/09/05/evolution-strategies.html>

⁹⁸https://colab.research.google.com/drive/1_OoZ3z_C5J15gnxDOE9VEMCTs-F18pvM

⁹⁹<https://colab.research.google.com/drive/1bMZWWhdm-mT9NJENWoVewUks7cGV10go>

¹⁰⁰https://colab.research.google.com/drive/1lvyKjFtc_C_8njCKD-MnXEW8LPS2RPr6

¹⁰¹[https://www.cell.com/current-biology/fulltext/S0960-9822\(18\)30693-6](https://www.cell.com/current-biology/fulltext/S0960-9822(18)30693-6)

¹⁰²<https://www.pnas.org/content/early/2020/01/07/1910837117>

¹⁰³<https://cdorgs.github.io>

¹⁰⁴https://github.com/skriegman/reconfigurable_organisms

¹⁰⁵<http://metalearning-symposium.ml>

¹⁰⁶<https://medium.com/applied-data-science/how-to-build-your-own-muzero-in-python-f77d5718061a>

¹⁰⁷<https://frpays.github.io/lc0-js/engine.html>

¹⁰⁸<https://github.com/frpays/lc0-js/>

¹⁰⁹<https://blog.openai.com/openai-five/>

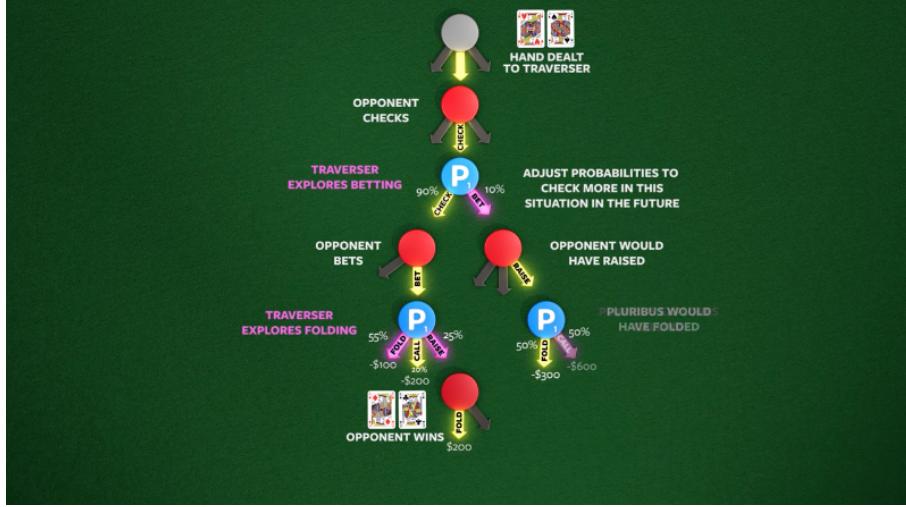


Figure 25: Facebook, Carnegie Mellon build first AI that beats pros in 6-player poker <https://ai.facebook.com/blog/pluribus-first-ai-to-beat-pros-in-6-player-poker>

"Artificial Intelligence is about recognising patterns, Artificial Life is about creating patterns." — Mizuki Oka et al.

Active Learning Without Teacher. In *Intrinsic Social Motivation via Causal Influence in Multi-Agent RL*, Jaques et al. (2018) <https://arxiv.org/abs/1810.08647> propose an intrinsic reward function designed for multi-agent RL (MARL), which awards agents for having a causal influence on other agents' actions. Open-source implementation¹¹⁰.

"Open-ended Learning in Symmetric Zero-sum Games," Balduzzi et al.: <https://arxiv.org/abs/1901.08106>

- ❖ Neural MMO v1.3: A Massively Multiagent Game Environment for Training and Evaluating Neural Networks, Suarez et al.¹¹¹ Project Page <https://jsuarez5341.github.io>, Video¹¹² and Slides¹¹³.
- ❖ Neural MMO: A massively multiagent env. for simulations with many long-lived agents. Code¹¹⁴ and 3D Client¹¹⁵.

4.5 Deep Meta-Learning

Learning to Learn[16].

"The notion of a neural "architecture" is going to disappear thanks to meta learning." — Andrew Trask

- ❖ Stanford CS330: Multi-Task and Meta-Learning, Finn et al., 2019¹¹⁶.
- ❖ Meta Learning Shared Hierarchies[18] (*The Lead Author is in High School!*).
- ❖ Causal Reasoning from Meta-reinforcement Learning <https://arxiv.org/abs/1901.08162>.

4.5.1 MAML: Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks

The goal of *model-agnostic meta-learning for fast adaptation of deep networks* is to train a model on a variety of learning tasks, such that it can solve new learning tasks using only a small number of training samples[20].

$$\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta_i'}) \quad (12)$$

¹¹⁰https://github.com/eugenevinitksy/sequential_social_dilemma_games

¹¹¹<https://arxiv.org/abs/2001.12004>

¹¹²<https://youtube.com/watch?v=DkHopV1RSxw>

¹¹³https://docs.google.com/presentation/d/1tqm_Do9ph-duqqAlx3r9lI5Nbfb9yUfNETXk1Qo4zSw/edit?usp=sharing

¹¹⁴<https://github.com/openai/neural-mmo>

¹¹⁵<https://github.com/jsuarez5341/neural-mmo-client>

¹¹⁶<https://youtube.com/playlist?list=PLoR0Mvodv4rMC6zfYmnD7UG3LVvwaITY5>

A meta-learning algorithm takes in a distribution of tasks, where each task is a learning problem, and it produces a quick learner — a learner that can generalize from a small number of examples[17].

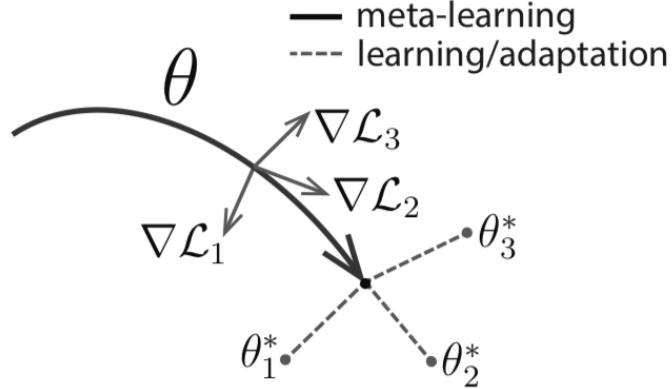


Figure 26: Diagram of Model-Agnostic Meta-Learning (MAML)

- ❖ How to Train MAML (Model-Agnostic Meta-Learning)¹¹⁷.
- ❖ Meta-Learning with Implicit Gradients <https://arxiv.org/abs/1909.04630>.
- ❖ Collaboratory reimplementation of MAML (Model-Agnostic Meta-Learning) in TF 2.0¹¹⁸.
- ❖ Torchmeta: A Meta-Learning library for PyTorch¹¹⁹ <https://github.com/tristandeleu/pytorch-meta>.

4.5.2 The Grand Challenge for AI Research | AI-GAs: AI-Generating Algorithms, an Alternate Paradigm for Producing General Artificial Intelligence

In *AI-GAs: AI-generating algorithms, an alternate paradigm for producing general artificial intelligence*¹²⁰, Jeff Clune describes an exciting path that ultimately may be successful at producing general AI. The idea is to create an AI-generating algorithm (AI-GA), which automatically learns how to produce general AI.

Three Pillars are essential for the approach: (1) **Meta-learning architectures**, (2) **Meta-learning the learning algorithms themselves**, and (3) **Generating effective learning environments**.

- **The First Pillar**, meta-learning architectures, could potentially discover the building blocks : *convolution, recurrent layers, gradient-friendly architectures, spatial transformers, etc.*
- **The Second Pillar**, meta-learning learning algorithms, could potentially learn the building blocks : *intelligent exploration, auxiliary tasks, efficient continual learning, causal reasoning, active learning, etc.*
- **The Third Pillar**, generating effective and fully expressive learning environments, could learn things like : *co-evolution / self-play, curriculum learning, communication / language, multi-agent interaction, etc.*

On Earth,

"(. . .) a remarkably simple algorithm (Darwinian evolution) began producing solutions to relatively simple environments. The ‘solutions’ to those environments were organisms that could survive in them. Those organisms often created new niches (i.e. environments, or opportunities) that could be exploited. Ultimately, that process produced all of the engineering marvels on the planet, such as jaguars, hawks, and the human mind." — Jeff Clune

Turing Complete (universal computer) : an encoding that enables the creation any possible learning algorithm.
Darwin Complete : an environmental encoding that enables the creation of any possible learning environment.

- ❖ **Learning to Continually Learn.** Beaulieu et al. <https://arxiv.org/abs/2002.09571>.
- ❖ **Fully Differentiable Procedural Content Generation through Generative Playing Networks.** Bontrager et al.¹²¹

¹¹⁷<https://medium.com/towards-artificial-intelligence/how-to-train-maml-model-agnostic-meta-learning-90aa093f8e46>

¹¹⁸<https://colab.research.google.com/github/mari-linhares/tensorflow-maml/blob/master/maml.ipynb>

¹¹⁹<https://medium.com/pytorch/torchmeta-a-meta-learning-library-for-pytorch-f76c2b07ca6d>

¹²⁰<https://arxiv.org/abs/1905.10985>

¹²¹<https://arxiv.org/abs/2002.05259>

5 Symbolic AI

- ❖ On neural-symbolic computing: suggested readings on foundations of the field. Luis Lamb¹²².
- ❖ **Neural-Symbolic Learning and Reasoning: A Survey and Interpretation.** Besold et al.¹²³
- ❖ **Neural Module Networks for Reasoning over Text.** Gupta et al.¹²⁴ Code.¹²⁵
- ❖ **The compositionality of neural networks: integrating symbolism and connectionism.** Hupkes et al.¹²⁶
- ❖ **Neuro-symbolic A.I. is the future of artificial intelligence. Here's how it works.** Luke Dormehl¹²⁷.
- ❖ **DDSP: Differentiable Digital Signal Processing.** Engel et al. Blog¹²⁸, Colab¹²⁹, Paper¹³⁰ and Code¹³¹.
- ❖ **Differentiable Reasoning on Large Knowledge Bases and Natural Language.** Minervini et al.¹³² Open-source neuro-symbolic reasoning framework, in TensorFlow <https://github.com/uclnlp/gntp>.

6 Environments

Platforms for training autonomous agents.

"Run a physics sim long enough and you'll get intelligence." — Elon Musk

6.1 OpenAI Gym

"Situation awareness is the perception of the elements in the environment within a volume of time and space, and the comprehension of their meaning, and the projection of their status in the near future." — Endsley (1987)

The OpenAI Gym <https://gym.openai.com/> (Blog¹³³ | GitHub¹³⁴) is a toolkit for developing and comparing reinforcement learning algorithms. What makes the gym so great is a common API around environments.

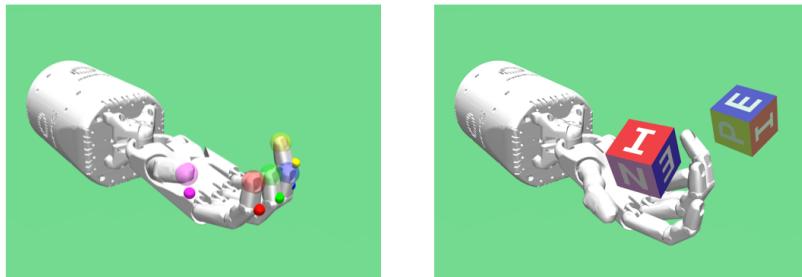


Figure 27: Robotics Environments <https://blog.openai.com/ingredients-for-robotics-research/>

"By framing the approach within the popular OpenAI Gym framework, design firms can create more realistic environments – for instance, incorporate strength of materials, safety factors, malfunctioning of components under stressed conditions, and plug existing algorithms into this framework to optimize also for design aspects such as energy usage, easy-of-manufacturing, or durability." — David Ha¹³⁵

How to create new environments for Gym¹³⁶. **Minimal example with code and agent** (*evolution strategies on foo-v0*):

¹²²<https://twitter.com/luislamb/status/1218575842340634626>

¹²³<https://arxiv.org/abs/1711.03902>

¹²⁴<https://arxiv.org/abs/1912.04971>

¹²⁵<https://nitishgupta.github.io/nmn-drop>

¹²⁶<https://arxiv.org/abs/1908.08351>

¹²⁷<https://www.digitaltrends.com/cool-tech/neuro-symbolic-ai-the-future/>

¹²⁸<http://magenta.tensorflow.org/ddsp>

¹²⁹<http://g.co/magenta/ddsp-demo>

¹³⁰<http://g.co/magenta/ddsp-paper>

¹³¹<http://github.com/magenta/ddsp>

¹³²<https://arxiv.org/abs/1912.10824>

¹³³<https://blog.openai.com/openai-gym-beta/>

¹³⁴<https://github.com/openai/gym>

¹³⁵<https://designrl.github.io>

¹³⁶<https://github.com/openai/gym/blob/master/docs/creating-environments.md>

1. Download `gym-foo` <https://drive.google.com/file/d/1r2A8J9CJjIQNwss246gATeDOLLMzpUT-/view?usp=sharing>
2. `cd gym-foo`
3. `pip install -e .`
4. `python ES-foo.py`

He're another more difficult (*for the agent!*) new environment for Gym (*evolution strategies on foo-v3*):

1. Download `gym-foo-v3`¹³⁷
2. `cd gym-foo-v3`
3. `pip install -e .`
4. `python ES-foo-v3.py`

- ❖ OpenAI Gym Environment for Trading¹³⁸.
- ❖ Fantasy Football AI Environment <https://github.com/njustesen/ffai>.
- ❖ Create custom gym environments from scratch — A stock market example¹³⁹.
- ❖ IKEA Furniture Assembly Environment <https://clvrai.github.io/furniture/>.
- ❖ Minimalistic Gridworld Environment <https://github.com/maximecb/gym-minigrid>.
- ❖ **DoorGym**: A Scalable Door Opening Environment and Baseline Agent, Urakami et al., 2019¹⁴⁰.
- ❖ **gym-gazebo2**, a toolkit for reinforcement learning using ROS 2 and Gazebo, Lopez et al., 2019¹⁴¹.
- ❖ OFFWORLD GYM Open-access physical robotics environment for real-world reinforcement learning¹⁴².
- ❖ Safety Gym: environments to evaluate agents with safety constraints <https://github.com/openai/safety-gym>.

6.2 DeepMind Lab

DeepMind Lab: A customisable 3D platform for agent-based AI research <https://github.com/deepmind/lab>.

- DeepMind Control Suite https://github.com/deepmind/dm_control.
- Convert DeepMind Control Suite to OpenAI Gym Envs <https://github.com/zuoxingdong/dm2gym>.

6.3 Unity ML-Agents

Unity ML Agents allows to create environments where intelligent agents (*Single Agent, Cooperative and Competitive Multi-Agent and Ecosystem*) can be trained using RL, neuroevolution, or other ML methods <https://unity3d.ai>.

- Getting Started with Marathon Environments for Unity ML-Agents¹⁴³ <https://github.com/Unity-Technologies/marathon-envs>.
- Arena: A General Evaluation Platform and Building Toolkit for Multi-Agent Intelligence¹⁴⁴.

6.4 AI Habitat

AI Habitat enables training of embodied AI agents (virtual robots) in a highly photorealistic and efficient 3D simulator, before transferring the learned skills to reality. By Facebook AI Research <https://aihabitat.org/>.

Why the name Habitat? Because that's where AI agents live!

¹³⁷<https://drive.google.com/file/d/1cGncsXJ56UUKC09MaRWJVTnxiQEnLuxS/view?usp=sharing>

¹³⁸<https://github.com/hackthemarket/gym-trading>

¹³⁹<https://towardsdatascience.com/creating-a-custom-openai-gym-environment-for-stock-trading-be532be3910e>

¹⁴⁰<https://arxiv.org/abs/1908.01887>

¹⁴¹<https://arxiv.org/abs/1903.06278>

¹⁴²<https://gym.offworld.ai>

¹⁴³<https://towardsdatascience.com/gettingstartedwithmarathonenvs-v0-5-0a-c1054a0b540c>

¹⁴⁴<https://arxiv.org/abs/1905.08085>

6.5 POET: Paired Open-Ended Trailblazer

Diversity is the premier product of evolution. Endlessly generate increasingly complex and diverse learning environments¹⁴⁵. Open-endedness could generate learning algorithms reaching human-level intelligence[23].

- Implementation of the POET algorithm <https://github.com/uber-research/poet>.

7 Deep-Learning Hardware



Figure 28: Edge TPU - Dev Board <https://coral.ai/products/dev-board/>

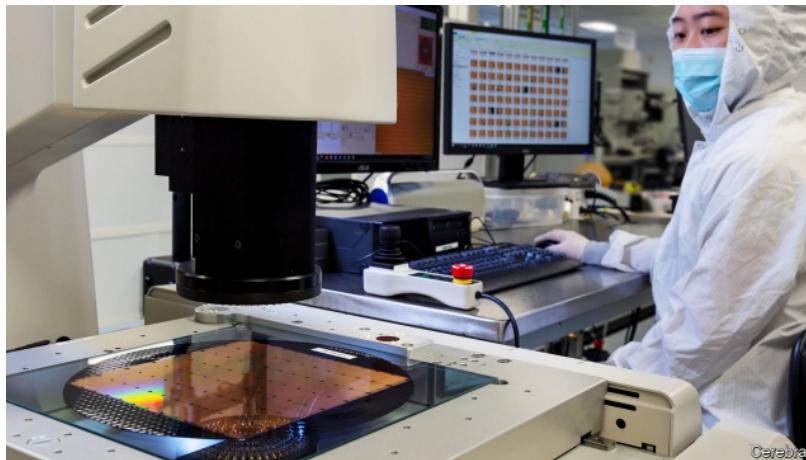


Figure 29: The world's largest chip : Cerebras Wafer Scale Engine <https://www.cerebras.net>

- ❖ Which GPU(s) to Get for Deep Learning, by Tim Dettmers¹⁴⁶.
- ❖ A Full Hardware Guide to Deep Learning, by Tim Dettmers¹⁴⁷.
- ❖ Jetson Nano. A small but mighty AI computer to create intelligent systems¹⁴⁸.
- ❖ Build AI that works offline with Coral Dev Board, Edge TPU, and TensorFlow Lite, by Daniel Situnayake¹⁴⁹.

¹⁴⁵<https://eng.uber.com/poet-open-ended-deep-learning/>

¹⁴⁶<http://timdettmers.com/2019/04/03/which-gpu-for-deep-learning/>

¹⁴⁷<http://timdettmers.com/2018/12/16/deep-learning-hardware-guide/>

¹⁴⁸<https://www.nvidia.com/en-us/autonomous-machines/embedded-systems/jetson-nano/>

¹⁴⁹<https://medium.com/tensorflow/build-ai-that-works-offline-with-coral-dev-board-edge-tpu-and-tensorflow-lite-70>

8 Deep-Learning Software

8.1 TensorFlow

TensorFlow Hub is a library for reusable machine learning modules <https://www.tensorflow.org/hub>.
TensorFlow.js allows machine learning to happen within the web browser <https://www.tensorflow.org/js/>.

- TensorFlow Lite for Microcontrollers¹⁵⁰.
- TensorBoard in Jupyter Notebooks¹⁵¹. Colab¹⁵².
- TensorFlow 2.0 + Keras Crash Course. Colab¹⁵³.
- tf.keras (TensorFlow 2.0) for Researchers: Crash Course. Colab¹⁵⁴.
- TensorFlow Tutorials <https://www.tensorflow.org/tutorials>.
- TensorFlow 2.0: basic ops, gradients, data preprocessing and augmentation, training and saving. Colab¹⁵⁵.

8.2 PyTorch

- PyTorch primer. Colab¹⁵⁶.
- PyTorch internals <http://blog.ezyang.com/2019/05/pytorch-internals/>

9 AI Art | A New Day Has Come in Art Industry



Figure 30: On October 25, 2018, the first AI artwork ever sold at Christie’s auction house fetched USD 432,500.

The code (*art-DCGAN*) for the first artificial intelligence artwork ever sold at Christie’s auction house (Figure 23) is a modified implementation of DCGAN focused on generative art: <https://github.com/robbiebarrat/art-dcgan>.

- **TensorFlow Magenta**. An open source research project exploring the role of ML in the creative process.¹⁵⁷.
- **Magenta Studio**. A suite of free music-making tools using machine learning models!¹⁵⁸.
- **Style Transfer Tutorial** https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/r2/tutorials/generative/style_transfer.ipynb
- **AI x AR Paper Cubes** <https://experiments.withgoogle.com/paper-cubes>.
- **Photo Wake-Up** <https://grail.cs.washington.edu/projects/wakeup/>.

¹⁵⁰<https://www.tensorflow.org/lite/microcontrollers>

¹⁵¹https://www.tensorflow.org/tensorboard/tensorboard_in_notebooks

¹⁵²https://colab.research.google.com/github/tensorflow/blob/master/docs/tensorboard_in_notebooks.ipynb

¹⁵³<https://colab.research.google.com/drive/1UCJt8EYjlzCs1H1d1X0iDGYJsHKwu-NO>

¹⁵⁴<https://colab.research.google.com/drive/14CvUNTaX10FHDfaKaaZzrBsvMfhCOHIR>

¹⁵⁵https://colab.research.google.com/github/zaidalyafeai/Notebooks/blob/master/TF_2_0.ipynb

¹⁵⁶<https://colab.research.google.com/drive/1DgkVmi6GksWOByhYVQpyUB4Rk3PUq0Cp>

¹⁵⁷<https://magenta.tensorflow.org>

¹⁵⁸<https://magenta.tensorflow.org/studio>

- **COLLECTION.** AI Experiments <https://experiments.withgoogle.com/ai>.

"*The Artists Creating with AI Won't Follow Trends; THEY WILL SET THEM.*" — The House of Montréal.AI Fine Arts

- ❖ Tuning Recurrent Neural Networks with Reinforcement Learning¹⁵⁹.
- ❖ **MuseNet.** Generate Music Using Many Different Instruments and Styles!¹⁶⁰.
- ❖ Infinite stream of machine generated art. Valentin Vieriu <https://art42.net>.
- ❖ Deep Multispectral Painting Reproduction via Multi-Layer, Custom-Ink Printing. Shi et al.¹⁶¹.
- ❖ Discovering Visual Patterns in Art Collections with Spatially-consistent Feature Learning. Shen et al.¹⁶².
- ❖ Synthesizing Programs for Images using Reinforced Adversarial Learning, Ganin et al., 2018¹⁶³. Agents¹⁶⁴.

10 AI Macrostrategy: Aligning AGI with Human Interests

Montréal.AI Governance: Policies at the intersection of AI, Ethics and Governance.

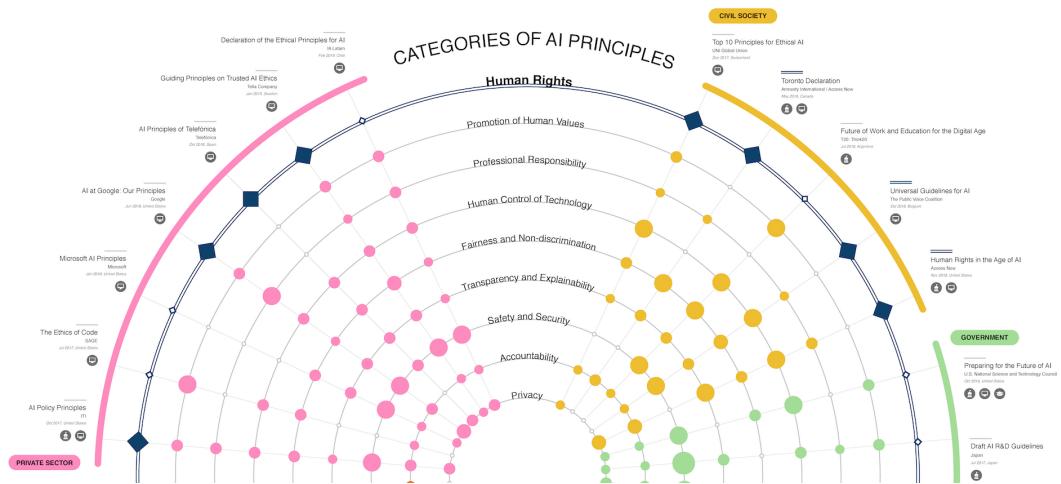


Figure 31: A Map of Ethical and Right-Based Approaches <https://ai-hr.cyber.harvard.edu/primp-viz.html>

"(AI) will rank among our greatest technological achievements, and everyone deserves to play a role in shaping it." — Fei-Fei Li

- ❖ **AI Index.** <http://aiindex.org>.
- ❖ **Malicious AI Report.** <https://arxiv.org/pdf/1802.07228.pdf>.
- ❖ **ADDRESS PREPARED BY POPE FRANCIS FOR THE PLENARY ASSEMBLY OF THE PONTIFICAL ACADEMY FOR LIFE¹⁶⁵.**
- ❖ **Artificial Intelligence and Human Rights.** <https://ai-hr.cyber.harvard.edu>.
- ❖ **Ethically Aligned Design, First Edition¹⁶⁶.** From Principles to Practice <https://ethicsinaction.ieee.org>.

"It's springtime for AI, and we're anticipating a long summer." — Bill Braun

¹⁵⁹<https://magenta.tensorflow.org/2016/11/09/tuning-recurrent-networks-with-reinforcement-learning>

¹⁶⁰<https://openai.com/blog/musenet/>

¹⁶¹<http://people.csail.mit.edu/liangs/papers/ToG18.pdf>

¹⁶²<https://arxiv.org/pdf/1903.02678.pdf>

¹⁶³<http://proceedings.mlr.press/v80/ganin18a.html>

¹⁶⁴<https://github.com/deepmind/spiral>

¹⁶⁵http://w2.vatican.va/content/francesco/en/speeches/2020/february/documents/papa-francesco_20200228_accademia-perlavita.html

¹⁶⁶<https://standards.ieee.org/content/dam/ieee-standards/standards/web/documents/other/ead1e.pdf>

References

- [1] Mnih et al. Human-Level Control Through Deep Reinforcement Learning. In *Nature* 518, pages 529–533. 26 February 2015. <https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>
- [2] Yann LeCun, Yoshua Bengio and Geoffrey Hinton. Deep Learning. In *Nature* 521, pages 436–444. 28 May 2015. <https://www.cs.toronto.edu/~hinton/absps/NatureDeepReview.pdf>
- [3] Goodfellow et al. Generative Adversarial Networks. *arXiv preprint arXiv:1406.2661*, 2014. <https://arxiv.org/abs/1406.2661>
- [4] Yoshua Bengio, Andrea Lodi, Antoine Prouvost. Machine Learning for Combinatorial Optimization: a Methodological Tour d’Horizon. *arXiv preprint arXiv:1811.06128*, 2018. <https://arxiv.org/abs/1811.06128>
- [5] Brockman et al. OpenAI Gym. 2016. <https://gym.openai.com>
- [6] Devlin et al. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *arXiv preprint arXiv:1810.04805*, 2018. <https://arxiv.org/abs/1810.04805>
- [7] Dai et al. Semi-supervised Sequence Learning. *arXiv preprint arXiv:1511.01432*, 2015. <https://arxiv.org/abs/1511.01432>
- [8] Mnih et al. Asynchronous Methods for Deep Reinforcement Learning. *arXiv preprint arXiv:1602.01783*, 2016. <https://arxiv.org/abs/1602.01783>
- [9] Schulman et al. Proximal Policy Optimization Algorithms. *arXiv preprint arXiv:1707.06347*, 2017. <https://arxiv.org/abs/1707.06347>
- [10] Mnih et al. Playing Atari with Deep Reinforcement Learning. *DeepMind Technologies*, 2013. <https://www.cs.toronto.edu/~vmnih/docs/dqn.pdf>
- [11] Ha et al. Recurrent World Models Facilitate Policy Evolution. *arXiv preprint arXiv:1809.01999*, 2018. <https://arxiv.org/abs/1809.01999>
- [12] Kenneth et al. Designing neural networks through neuroevolution. In *Nature Machine Intelligence* VOL 1, pages 24–35. January 2019. <https://www.nature.com/articles/s42256-018-0006-z.pdf>
- [13] So et al. The Evolved Transformer. *arXiv preprint arXiv:1901.11117*, 2019. <https://arxiv.org/abs/1901.11117>
- [14] Silver et al. Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm. *arXiv preprint arXiv:1712.01815*, 2017. <https://arxiv.org/abs/1712.01815>
- [15] Silver et al. AlphaGo Zero: Learning from scratch. In *DeepMind’s Blog*, 2017. <https://deepmind.com/blog/alphago-zero-learning-scratch/>
- [16] Andrychowicz et al. Learning to learn by gradient descent by gradient descent. *arXiv preprint arXiv:1606.04474*, 2016. <https://arxiv.org/abs/1606.04474>
- [17] Nichol et al. Reptile: A Scalable Meta-Learning Algorithm. 2018. <https://blog.openai.com/reptile/>
- [18] Frans et al. Meta Learning Shared Hierarchies. *arXiv preprint arXiv:1710.09767*, 2017. <https://arxiv.org/abs/1710.09767>
- [19] Zoph and Le, 2017 Neural Architecture Search with Reinforcement Learning. *arXiv preprint arXiv:1611.01578*, 2017. <https://arxiv.org/abs/1611.01578>
- [20] Finn et al., 2017 Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks. *arXiv preprint arXiv:1703.03400*, 2017. <https://arxiv.org/abs/1703.03400>
- [21] Salimans et al. Evolution Strategies as a Scalable Alternative to Reinforcement Learning. 2017. <https://blog.openai.com/evolution-strategies/>
- [22] Lehman et al. The Surprising Creativity of Digital Evolution: A Collection of Anecdotes from the Evolutionary Computation and Artificial Life Research Communities. *arXiv preprint arXiv:1803.03453*, 2018. <https://arxiv.org/abs/1803.03453>
- [23] Wang et al. Paired Open-Ended Trailblazer (POET): Endlessly Generating Increasingly Complex and Diverse Learning Environments and Their Solutions. *arXiv preprint arXiv:1901.01753*, 2019. <https://arxiv.org/abs/1901.01753>
- [24] Foerster et al. Learning to Model Other Minds. 2018. <https://blog.openai.com/learning-to-model-other-minds/>
- [25] Rabinowitz et al. Machine Theory of Mind. *arXiv preprint arXiv:1802.07740*, 2018. <https://arxiv.org/abs/1802.07740>