



MONTREAL.AI · QUEBEC.AI


MISSION-SOVEREIGN SUCCESSORS

GoalOS Singularity Navigator Ω + SEIZE

ALPHA, BETA, SPREAD, AND THE VERIFIED SUCCESSION INSTITUTION IN SINGULARITY TIME

A Game-Theoretic, Proof-Capital and Constitutional Architecture for
Manufacturing, Proving, Governing and Compounding Mission Alpha

**Beta is rented. Alpha is manufactured. The Spread is underwritten.
The Mission-Sovereign Successor captures and compounds mission
advantage—under fresh proof.**



SUCCESSOR MANIFOLD → MISSION ADVANTAGE GRADIENT → SEIZE UNDERWRITING → BOUNDED AGI JOBS → SPECIALIST
ASI → MISSION-SOVEREIGN SUCCESSOR

By

Vincent Boucher

President, MONTREAL.AI & QUEBEC.AI

<https://montreal.ai/> · president@montreal.ai

Version 5.0.0 · Authoritative Publication Edition · 9 August 2026

Montreal, Quebec, Canada

Constitutional Visual Map



Institutional visual plate. Exact definitions, formulas, thresholds, authority states and claims are governed by the typeset paper and the verified release artifacts, not by image-generated microtypography.

The governing declaration

In calendar time, institutions optimize plans. In Singularity Time, institutions govern succession. Rented Beta establishes a moving market baseline ; the Spread reveals the investable mission opportunity ; SEIZE selects a governed trajectory through the Successor Manifold ; Bounded AGI Jobs constitute and test the candidate ; Specialist ASI may prove mission-specific Alpha ; and only accountable admission may constitute a Mission-Sovereign Successor with scoped, revocable authority.

This paper is authored by **Vincent Boucher**, President of **MONTREAL.AI & QUEBEC.AI**. The official publication website is montreal.ai and the official author contact is president@montreal.ai. The authoritative scholarly text is the typeset PDF generated from the versioned LaTeX source in this release. Illustrations are explanatory ; exact claims, equations, thresholds and governance rules are controlled by the typeset text and release manifest.

Abstract

This paper defines *GoalOS Singularity Navigator* $\Omega + SEIZE$ as a Verified Succession Institution for converting broadly available frontier intelligence into proprietary, proof-bearing mission capacity. Its central distinction is between **AI Beta**, the moving market baseline available to competent adopters ; the **Mission Advantage Spread**, the investable distance between that baseline and a reachable superior architecture ; **Mission Alpha**, the residual risk-adjusted surplus that remains after the incumbent, the best external alternative, complete cost, tail risk, dependency and Proof Debt are counted ; and the **Mission-Sovereign Successor**, the complete institution that captures, governs and compounds that Alpha under fresh proof.

The canonical progression is *Successor Manifold* $\rightarrow$ *Mission Advantage Gradient* $\rightarrow$ *SEIZE Underwriting* $\rightarrow$ *Bounded AGI Jobs* $\rightarrow$ *Specialist ASI* $\rightarrow$ *Mission-Sovereign Successor*. The Successor Manifold represents the newly reachable space of complete architectures. The Mission Advantage Gradient identifies the direction in that space that increases proof-adjusted mission value most rapidly. SEIZE treats formation as a real-options and mechanism-design problem : select the smallest reversible trajectory whose proof can change the institutional decision. Bounded AGI Jobs then constitute models, data, tools, workflows, verifiers, evidence, controls, human decision rights, rollback and Chronicle memory as explicit work-and-proof contracts.

The paper formalizes Mission Alpha as

$$\alpha_m(x, t) = U_m(S_m(x, t)) - \max\{U_m(I_m(t)), U_m(B_m(t))\} - C_m - R_m - D_m,$$

requires authority to satisfy $A_t \preceq \phi(E_t)$, and models Alpha dynamics as validated institutional learning minus commoditization, imitation, drift and maintenance burden. It shows why access to the same frontier capability tends toward a symmetric Beta equilibrium ; why differentiated mission architecture and retained evidence may create asymmetric Alpha ; why proof acts as collateral for authority ; and why model portability, evidence ownership and Chronicle continuity create a sovereignty premium against vendor hold-up.

A Specialist ASI is an earned mission-bounded designation for a complete system that outperforms the incumbent and the best credible alternative on fresh, protected and representative work. A Mission-Sovereign Successor is that proven system only after separate accountable admission and the grant of a scoped, versioned, expiring and revocable Authority Envelope. The framework does not assert that universal ASI presently exists, that any external system has already earned Specialist ASI status, or that every mission should be delegated. It specifies a falsifiable pathway for testing those narrower claims. The paper provides design propositions, evaluation doctrine, a 21-job formation portfolio, a dynamic renewal loop, a worked acquisition-underwriting example, board metrics, machine-readable records, falsification conditions and a 90-day activation path. Its governing rule is exact : **authority never exceeds fresh proof**.

Keywords : GoalOS ; SEIZE ; Mission-Sovereign Successor ; Specialist ASI ; AI Beta ; Mission Alpha ; Mission Advantage Spread ; Successor Manifold ; Mission Advantage Gradient ; bounded AGI Jobs ; proof capital ; Chronicle ; game theory ; Singularity Time.

Résumé exécutif

Ce document présente *GoalOS Singularity Navigator* Ω + *SEIZE* comme une institution de succession vérifiée qui transforme l'intelligence de frontière largement accessible en capacité missionnelle propriétaire, prouvée et gouvernée. La distinction centrale oppose le **bêta d'intelligence**, soit la base mobile que le marché peut louer ; l'**écart d'avantage missionnel**, soit l'occasion investissable entre cette base et une architecture supérieure atteignable ; l'**alpha missionnel**, soit le surplus résiduel ajusté au risque après prise en compte de la meilleure solution existante, des coûts complets, des risques, des dépendances et de la dette de preuve ; et le **successeur souverain pour la mission**, soit l'institution complète qui capte, gouverne et compose cet Alpha sous preuve fraîche.

La progression canonique est la suivante : *Variété des successeurs* $\rightarrow$ *gradient d'avantage missionnel* $\rightarrow$ *souscription SEIZE* $\rightarrow$ *tâches AGI bornées* $\rightarrow$ *ASI spécialisée* $\rightarrow$ *successeur souverain pour la mission*. SEIZE choisit une trajectoire réversible à valeur attendue positive ; les tâches AGI bornées constituent les modèles, données, outils, flux, vérificateurs, preuves, contrôles, droits décisionnels humains, mécanismes de retour arrière et mémoire Chronicle ; puis une preuve comparative indépendante détermine si le système domine réellement la mission.

L'ASI spécialisée est une désignation à mériter, limitée à une mission définie. Le successeur souverain pour la mission n'est ni une personne morale autonome ni une autorité illimitée : il s'agit d'une institution spécialisée qui a réussi des essais frais et représentatifs, puis reçu une enveloppe d'autorité circonscrite, versionnée, temporaire, surveillée et révocable. Le document n'affirme pas qu'une ASI universelle existe déjà ni qu'un système externe a déjà mérité cette désignation. La règle directrice demeure : **l'autorité ne dépasse jamais la preuve fraîche**.

Central doctrine

Rented Beta establishes the moving market baseline. The Spread reveals the investable mission opportunity. GoalOS maps the Successor Manifold and its Mission Advantage Gradient. SEIZE underwrites one governed trajectory. Bounded AGI Jobs constitute and test the candidate. Fresh comparative proof may establish Specialist ASI and Mission Alpha. Only accountable admission may create a Mission-Sovereign Successor.

Alpha is not a stage and not a label. It is the risk-adjusted residual mission surplus that exists only when the complete successor outperforms both the incumbent and the best credible external alternative after cost, tail risk, rights, dependency and Proof Debt are counted.

Epistemic status and claim ladder

The paper deliberately separates six kinds of statement so that design language is not mistaken for empirical proof:

1. **External facts** are attributed to current primary or official sources.
2. **Definitions** stipulate the precise GoalOS meaning of terms such as AI Beta, Mission Alpha and Mission-Sovereign Successor.
3. **Design propositions** follow from stated assumptions and are not universal laws of nature.
4. **Empirical hypotheses** are intended to be tested on protected fresh work and may be falsified.
5. **Constitutional rules** are governance choices that constrain authority even when raw capability is high.
6. **State claims**—including Specialist ASI or Mission-Sovereign admission—require an exact candidate, release, evidence docket, accountable decision and expiry.

Principal contributions

Contribution	What it adds to GoalOS + SEIZE
AI Beta and moving benchmark	Treats broadly rentable frontier capability as a dynamic common factor rather than a durable moat.
Mission Advantage Spread	Measures the gross, proof-adjusted opportunity above the incumbent and the best credible external alternative.
Mission Alpha	Defines residual mission surplus after complete cost, tail risk, dependency, Proof Debt and governance burden.
Successor Manifold and Gradient	Represents the reachable architecture space and the direction in which proof-adjusted superiority increases fastest.
SEIZE trajectory underwriting	Applies Bayesian real-options discipline to select a reversible trajectory and the smallest decision-changing proof budget.
Bounded AGI Jobs	Constitutes models, data, tools, verifiers, workflow, controls, human authority and evidence as explicit work-and-proof contracts.
Specialist ASI	Defines an earned mission-specific superiority state against the best credible alternative set, not a claim of universal intelligence.
Mission-Sovereign Successor	Separates empirical dominance from accountable admission and scoped, expiring, revocable mission authority.
Game-theoretic constitution	Uses role separation, incentive compatibility, independent verification, anti-collusion controls and regret-aware choice.
Sovereignty premium	Formalizes the option value of mission, architecture, data, evidence, memory, authority, portability and economic control.
Dynamic Alpha loop	Connects Chronicle, fresh proof, re-underwriting, bounded improvement, renewed authority and compounding while modelling Alpha compression.
Falsifiable operating doctrine	Provides experiments, scorecards, stop conditions, machine-readable records and a worked successor-underwriting case.

Recommended citation

Boucher, Vincent (2026). *Mission-Sovereign Successors: GoalOS Singularity Navigator Ω + SEIZE—Alpha, Beta, Spread, and the Verified Succession Institution in Singularity Time*. Version 5.0.0, MONTREAL.AI Research, Montreal, Quebec, Canada. <https://montreal.ai/>.

This paper is research and institutional design. It is not legal, tax, investment, regulatory, medical, engineering, security or other reserved professional advice. “AI Beta,” “Mission Alpha,” “Mission Advantage Spread,” “proof capital,” “NAV” and related terms are institutional operating constructs. Mission Alpha means benchmark-relative, risk-adjusted mission surplus; it is not a securities-return projection and is not a claim about AGIALPHA price, liquidity, redemption or investment performance. Candidate architectures, valuations, simulations, Specialist ASI designations and Mission-Sovereign states remain hypotheses until supported by exact facts, current sources, protected fresh-work evaluation, independent review, accountable admission and a suitable deployment. The public browser-local application does not itself execute external transactions or create regulated authority.

Contents

Abstract	i
Résumé exécutif	ii
Epistemic status and claim ladder	ii
Principal contributions	iii
Recommended citation	iii
1 The Operating Condition: Capability Time and Architecture Depreciation	1
1.1 An operational definition of Singularity Time	1
1.2 The claim discipline	1
1.3 Every incumbent is a depreciating architecture	2
1.4 The frontier is a portfolio, not a model	2
1.5 From model updates to institutional succession	3
2 The Verified Succession Institution	3
2.1 Definition	3
2.2 The institutional object	3
2.3 Five successor classes	4
2.4 Mission dominance	4
3 The Successor Manifold and the Mission Advantage Gradient	5
3.1 Architecture space as a constrained manifold	5
3.2 Institutional distance and the burden metric	6
3.3 The mission value field	6
3.4 Trajectory selection as a least-action problem	7
3.5 Phase transitions, metastability and the succession event	7
4 Specialist ASI and the Mission-Sovereign Successor	8
4.1 A mission-bounded, system-level definition	8
4.2 The best credible alternative set	8
4.3 Six superiority dimensions	8
4.4 Dominance is not sovereignty	9
4.5 The Specialist ASI claim boundary	9
5 Game-Theoretic Constitution and Mechanism Design	9
5.1 Succession as a dynamic game	9
5.2 Separation of powers as an incentive constraint	10

5.3	Mechanism-design objective	10
5.4	Regret-aware choice and adversarial challenge	10
5.5	Proof capital as an option premium	10
6	The Six-Stage Mission-Sovereign Progression	11
7	AI Beta, Mission Spread, Proven Alpha, and the Sovereign Successor	12
7.1	The four objects must not be conflated	12
7.2	AI Beta as a moving common factor	13
7.3	The reference architecture and the investable Spread	14
7.4	Mission Alpha as residual surplus	14
7.5	The Mission Advantage Gradient	15
7.6	SEIZE as Bayesian real-options underwriting	15
7.7	Bounded AGI Jobs as mechanism design	16
7.8	Specialist ASI as a comparative proof state	16
7.9	Proof as collateral for authority	16
7.10	Sovereignty and the vendor hold-up game	17
7.11	The repeated game and the Chronicle state	17
7.12	From the Beta equilibrium to the successor equilibrium	18
7.13	The canonical Alpha stack	18
8	The Successor Book	19
8.1	A live market for institutional architectures	19
8.2	Three horizons and option discipline	19
8.3	Incumbent impairment	20
9	SEIZE: The Succession Underwriting Microcycle	20
9.1	Purpose	20
9.2	S - Surface the succession event	21
9.3	E - Evaluate the successor frontier	21
9.4	I - Instantiate the succession constitution	22
9.5	Z - Zero in on decisive proof	22
9.6	E - Elevate one challenger	22
10	The Succession Constitution and Separation of Powers	23
10.1	Five distinct decisions	23
10.2	The Authority Envelope	23
10.3	Authority inheritance	24
10.4	Proof-carrying action	24
10.5	Institutional refusal	24

11 Succession Due Diligence: Purchasing Decisive Evidence	25
11.1 The causal bottleneck map	25
11.2 Five mandatory diagnostic families	25
11.3 Proof capital and stop-loss discipline	26
11.4 The evidence ladder	26
11.5 Reward grounding	27
12 The Successor Foundry	27
12.1 Formation is architecture search, not one technique	27
12.2 Candidate architecture classes	28
12.3 Quality diversity and the archive	28
12.4 The minimum-burden dominance rule	29
12.5 Search budgets and institutional compute	29
12.6 Bounded AGI Jobs as successor-manufacturing contracts	29
13 The Fresh-Successor Gate	30
13.1 Why inherited evidence is insufficient	31
13.2 Matched-pair comparison	31
13.3 The complete successor scorecard	31
13.4 Canary and bounded operation	32
14 Chronicle and the Succession Ledger	32
14.1 Memory is constitutional admission	32
14.2 The Succession Ledger	33
14.3 Failure-preserving memory	33
14.4 Chronicle precision	34
14.5 Impairment, repair and revocation	34
15 Capital-to-Capacity and Recursive Institutional Improvement	34
15.1 The compounding sequence	34
15.2 Value states	35
15.3 Allocation classes	35
15.4 Recursive improvement at the institutional level	35
15.5 Succession learning curves	36
16 Dual-Access Eligibility and the Authority Firewall	36
16.1 Controlling access rule	36
16.2 Access receipt is not authority	36
16.3 Fail-closed session discipline	37
16.4 Access and AGIALPHA claim boundary	37

17 The Five-Plane Technical Architecture	37
17.1 Why public navigation and protected formation must be separated	37
17.2 Identity and access	37
17.3 Evidence custody	38
17.4 Graceful degradation	38
18 Risk Architecture	38
18.1 Basis risk	38
18.2 Duration risk	38
18.3 Alpha decay and crowding	39
18.4 Dependency and liquidity risk	39
18.5 Authority-at-Risk	39
18.6 Tail loss	39
19 Board Metrics and Operating Cadence	40
19.1 The board dashboard	40
19.2 Operating cadence	40
19.3 The board brief	40
20 The Successor-Era Corporation: From Digital Presence to Institutional Agency	41
20.1 A strategic conjecture, not a guaranteed chronology	41
20.2 Why the corporation needs a portfolio, not one monolith	42
20.3 The adoption game	42
20.4 The build-rent boundary	43
20.5 The minimum viable successor institution	43
21 Mission Portfolios in Singularity Time	44
21.1 Where verified succession is most investable	44
21.2 Illustrative mission classes	44
21.3 Selection before formation	45
21.4 The preserve-build-become portfolio	45
22 Commercial Productization	46
22.1 One product, several engagement depths	46
22.2 The Readiness Mission as deal-flow formation	46
22.3 Sell the decision before the build	47
22.4 The Mission Pack as the delivery object	47
22.5 Renewal proof	47
23 Implementation Blueprint	47
23.1 Days 1-15: install the succession doctrine	48

23.2	Days 16-35: build the protected succession workspace	48
23.3	Days 36-60: run one real succession cycle	48
23.4	Days 61-90: constitute the first admitted successor	49
23.5	The first year	49
24	Worked Example: Acquisition Underwriting Successor	50
24.1	Mission and reference set	50
24.2	Successor Manifold	50
24.3	SEIZE underwriting thesis	50
24.4	Bounded AGI Job portfolio	51
24.5	Comparative proof	51
24.6	Authority and admission	52
24.7	Alpha capture and compounding	52
25	Falsifiability and the SN6 Completion Gate	53
25.1	What would falsify the thesis	53
25.2	Minimum empirical proof	53
25.3	The SN6 gate	54
26	Ten Laws of the Verified Succession Institution	54
27	Conclusion: Constitute Mission-Sovereign Successors	54
A	Canonical Machine-Readable Records	55
B	Canonical State Machine	56
C	Succession Underwriting Card	56
D	Fresh-Successor Gate Checklist	57
E	Publication Identity, Provenance and Source of Truth	57
E.1	Authoritative identity	57
E.2	Release provenance	57
E.3	Authorship and claim responsibility	57
F	Glossary	58
	Publication colophon	61

1 The Operating Condition: Capability Time and Architecture Depreciation

1.1 An operational definition of Singularity Time

This paper uses **Singularity Time** operationally rather than metaphysically. It denotes a regime in which the expected life of a material capability, cost, provider, workflow or institutional assumption is short enough that ordinary planning cycles become structurally unreliable. The question is not whether one date deserves a special label. The question is whether accepted operating assumptions decay faster than the institution can detect, test, authorize and replace them.

Recent research reinforces the need to treat capability time as uncertain and dynamic rather than as one settled arrival date. Work on the transition from AGI to ASI describes several possible pathways—scaling, paradigm shifts, recursive improvement and large multi-agent collectives—while emphasizing deep uncertainty and the possibility of a sequence of transformations rather than one discrete event (Genewein et al., 2026). Task-horizon research offers one operational signal of changing agent capability, while explicitly warning that benchmark-to-world external validity remains uncertain (Kwa et al., 2025). GoalOS therefore does not depend on a date forecast. It installs an institution that can repeatedly re-underwrite what is now reachable.

As of the publication date, METR maintains an updated reliability-conditioned task-completion time-horizon series rather than treating autonomy as a binary property, while NIST continues to frame trustworthy AI as a lifecycle of measurement, evaluation, documentation and risk management (METR, 2026; National Institute of Standards and Technology, 2023, 2024). These sources support continuous re-evaluation; they do not establish that universal ASI presently exists, that benchmark trends will continue unchanged, or that any candidate has earned mission authority.

Let H_c denote the **capability half-life**: the expected time until an assumption becomes obsolete enough to change a consequential institutional decision. Let L_a denote **adaptation lag**: the time required to detect a change, understand its relevance, constitute a mission, run tests, validate evidence, obtain authority, deploy and verify the result. The **Singularity Exposure Ratio** is

$$\text{SER} = \frac{L_a}{H_c}. \quad (1)$$

- $\text{SER} < 1$: the institution generally adapts before material assumptions decay;
- $\text{SER} = 1$: the institution operates at the edge of its adaptation capacity;
- $\text{SER} > 1$: the institution is structurally behind and increasingly acts from obsolete assumptions.

Speed alone is not the objective. An institution can reduce L_a by abandoning evidence, rights, security or human authority and thereby create a faster path to larger failure. GoalOS therefore optimizes **Verified Adaptation Velocity** rather than adoption speed:

$$\text{VAV} = \frac{\Delta C_v}{\Delta t} \cdot \frac{1}{1 + \lambda_R R + \lambda_H H + \lambda_P D_P}, \quad (2)$$

where C_v is verified capability, R residual risk, H human and reviewer burden, and D_P Proof Debt. The institutional objective is to increase accepted productive capability per unit time without hiding uncertainty or transferring unbounded authority.

1.2 The claim discipline

GoalOS uses powerful forward-looking language, but its institutional logic is deliberately conservative. A capability forecast is not a capability result; a benchmark result is not a mission result; mission dominance is not sovereignty; access is not authority; and historical proof is not fresh proof. The paper therefore treats “Singularity Time” as an operating condition to be tested through architecture half-life, benchmark movement, mission economics and evidence expiry rather than as a self-validating historical proclamation.

Design Proposition 0 — No authority by narrative

No forecast, brand, model scale, benchmark reputation, generated explanation or strategic urgency can substitute for the exact evidence required by the mission constitution. Where evidence is absent, the correct institutional state is uncertainty, not inherited authority.

1.3 Every incumbent is a depreciating architecture

In slow regimes, institutions can treat operating architectures as durable assets. In capability time, every architecture becomes a position with duration, carrying cost, decay and an eventual exit.

An **incumbent architecture** includes more than software. It includes:

- the objective and business model;
- the division of labour among people, machines and professional authorities;
- models, agents, tools, data, memory and providers;
- policies, permissions, controls and escalation routes;
- evidence, evaluation and acceptance procedures;
- rights, confidentiality, jurisdiction and contractual dependencies;
- capital, distribution, infrastructure and continuity arrangements.

The incumbent can depreciate because a new capability becomes feasible, a cost curve changes, a provider drifts, a right expires, a law changes, a security failure appears, a customer expectation shifts, or an entirely new institutional form becomes possible. The depreciation may be technical, economic, strategic, constitutional or all four.

Define the expected present cost of delaying a justified succession as

$$C_{\text{wait}}(S_0) = V_{\text{foregone}} + C_{\text{legacy}} + R_{\text{competitive}} + R_{\text{dependency}} + I_{\text{option loss}}, \quad (3)$$

and the cost of moving too early as

$$C_{\text{premature}}(S_i) = C_{\text{integration}} + R_{\text{security}} + R_{\text{error}} + C_{\text{rollback}} + R_{\text{lockin}} + D_{\text{proof}}. \quad (4)$$

The succession problem is not “change versus caution.” It is a portfolio comparison between the cost of retaining a decaying incumbent and the cost of forming an insufficiently proven successor.

Succession-time law

A capability shock does not authorize replacement. It creates a succession event: a duty to re-underwrite the incumbent against a field of candidate successors.

1.4 The frontier is a portfolio, not a model

No single model or provider represents the relevant frontier. The institution faces a changing portfolio of:


- general and specialist intelligence;
- algorithms, search procedures and executable environments;
- agents, teams and organizational designs;
- proprietary and public data;

- tools, simulators and physical systems;
- professional judgment and human authority;
- markets, counterparties, laws and infrastructure.

This is why the object of navigation must be the complete execution and governance architecture. GoalOS treats models as components in a succession market, not as sovereign decision makers or permanent monopolies (Boucher, 2026).

1.5 From model updates to institutional succession

The decisive shift is conceptual:



A model update can be part of succession, but succession may also require a new workflow, verifier, data structure, organizational role, authority boundary, distribution channel, capital allocation or institution. The unit of improvement is the mission-bearing institution.

2 The Verified Succession Institution

2.1 Definition

A Verified Succession Institution is a recurring constitutional system that detects when an incumbent architecture has become suboptimal, maintains a portfolio of successor options, authorizes bounded proof capital, constitutes candidate successors, compares one frozen challenger on fresh work, and permits authority to pass only after independent mission, governance and economic validation.

The institution has six permanent functions:

1. **Sensing:** detect capability, cost, law, security and market events that change the incumbent's expected value;
2. **Underwriting:** quantify the succession opportunity, downside, proof cost, duration and option value;
3. **Formation:** constitute competing complete successor architectures;
4. **Proof:** test one frozen challenger on fresh, protected and production-representative work;
5. **Authority:** decide separately whether work may execute, whether results are accepted, whether capability is admitted and whether value may be reinvested;
6. **Compounding:** preserve rights-cleared capability, failures and evidence so that subsequent succession cycles improve.

2.2 The institutional object

The central operating object is the **Successor Institution**. A Successor Institution is the complete architecture proposed to inherit one bounded mission, workflow, capability or institutional role from the incumbent.

It may include:

- a new intelligence substrate or algorithm;

- an agent organization or human-machine team;
- a changed objective, workflow or operating model;
- a knowledge representation, memory or executable world model;
- tools, planners, simulators or physical systems;
- an independent verifier and Evidence Docket;
- a new Authority Envelope, identity structure or approval route;
- rights, contracts, distribution and capital;
- monitoring, graceful degradation, rollback and expiry.

A candidate is not evaluated by novelty. It is evaluated by its verified contribution to the mission and by the complete institutional burden required to sustain it.

2.3 Five successor classes

Table 1 – Canonical successor classes.

Class	Meaning	Authority posture
Candidate Successor	A proposed complete architecture that may replace, augment or reconstitute the incumbent.	No inherited trust; development only.
Verified Successor	A frozen candidate that passes a protected fresh-work capability evaluation.	May proceed to broader gates; no automatic release.
Mission-Dominant Specialist ASI	A Verified Specialist Intelligence system that creates preregistered positive net mission gain over the incumbent and the best credible alternative after the complete denominator is counted.	Eligible for accountable Mission-Sovereign admission review.
Mission-Sovereign Successor	A Mission-Dominant Specialist ASI that separately passes rights, security, privacy, authority, monitoring, reversibility, economic and accountable admission gates.	Receives a scoped, versioned, expiring Authority Envelope.
Successor Reserve	A valuable but presently unqualified option preserved until a named catalyst, evidence object or capability event makes promotion rational.	No operational authority; option-preservation budget only.

2.4 Mission dominance

A candidate does not become mission-dominant merely by scoring higher on one measure. It must outperform the incumbent and the best credible alternative on a frozen mission after value, quality, time, complete cost, risk, human burden, sovereignty, reversibility and compounding capacity are considered.

Let S_0 be the incumbent, S_i a candidate and m the fresh mission. Define

$$\Delta_{\text{succ}}(S_i, S_0 \mid m) = w_Q \Delta Q + w_V \Delta V + w_T \Delta T + w_C \Delta C \quad (5)$$

$$+ w_S \Delta S + w_K \Delta K - \lambda_R \Delta R - \lambda_H \Delta H \quad (6)$$

$$- \lambda_P \Delta D_P - \lambda_U \Delta U, \quad (7)$$

where:

- Q : mission quality and critical-error performance;
- V : verified economic, scientific, operational or public value;
- T : time-to-accepted outcome;
- C : complete lifecycle cost;
- S : sovereignty, rights, privacy and strategic control;
- K : new productive capacity and future option value;
- R : residual risk and resilience burden;
- H : human and reviewer burden;
- D_P : Proof Debt;
- U : unwind, irreversibility and migration burden.

Mission dominance requires

$$\Delta_{\text{succ}}(S_i, S_0 \mid m) > \theta_m, \quad (8)$$

and the constitutional gate

$$\text{Gate}(S_i) = \text{ObjectiveFit} \wedge \text{Evidence} \wedge \text{Authority} \wedge \text{Security} \quad (9)$$

$$\wedge \text{Rights} \wedge \text{Reversibility} \wedge \text{NoGovernanceRegression}. \quad (10)$$

Weights, thresholds and hard gates must be frozen before the protected comparison. A candidate may not redefine the mission after observing results.

The successor does not have to be the most powerful architecture in the abstract. It must be the architecture that creates the strongest verified mission result at the lowest justified institutional burden, while preserving every non-negotiable constitutional constraint.

3 The Successor Manifold and the Mission Advantage Gradient

3.1 Architecture space as a constrained manifold

The frontier does not hand the institution one successor. It changes the topology of what can be reached. Let a complete successor architecture be represented by

$$x = (M, D, H, T, V, W, R, I, E, C, A), \quad (11)$$

where M is the model portfolio; D the rights-cleared data and examples; H the harness and orchestration; T the tools, retrieval and deterministic components; V the verifier and evaluation suite; W the workflow and Mission Pack; R the rights, security, resilience and rollback posture; I the system integration; E the evidence and realized-outcome record; C the Chronicle state; and A the Authority Envelope.

The **Successor Manifold** $\mathcal{M}_S$ is the constrained set of materially reachable architectures after a capability, cost, legal, market, security or infrastructure discontinuity. The language borrows from smooth-manifold and nonlinear-dynamics formalisms as disciplined analogies for constrained architecture space, gradients and bifurcations (Lee, 2013; Strogatz, 2015). It is not the set of all imaginable systems. Reachability depends on capital, rights, data, time, compute, professional authority, integration, evidence and the ability to reverse course.

The newly reachable, constraint-bearing space of complete successor institutions that could inherit, augment, hedge or retire a constituted mission.

Let $C_m(x) \in \{0, 1\}$ denote the mission's constitutional gate. The admissible region is

$$\mathcal{M}_{S,\text{adm}}(m) = \{x \in \mathcal{M}_S : C_m(x) = 1\}. \quad (12)$$

A candidate outside this set may be technically impressive yet constitutionally unavailable.

3.2 Institutional distance and the burden metric

Two architectures can be close in benchmark space and far apart institutionally. GoalOS therefore equips the manifold with a mission-dependent metric tensor $G_m(x)$ whose dimensions include formation capital, integration, migration, proof, security, rights, human burden, requalification, dependency, irreversibility and rollback.

For an infinitesimal change dx , institutional distance is

$$ds_m^2 = dx^\top G_m(x) dx. \quad (13)$$

This metric makes explicit why the shortest technical route may not be the least-burden institutional route. A hosted frontier system may be one API call away yet far from the desired state in privacy, control, rights or continuity. A specialist local architecture may require more formation work but lie closer to the mission's long-run sovereign state.

3.3 The mission value field

Define a proof-adjusted mission value field $\mathcal{V}_m : \mathcal{M}_S \rightarrow \mathbb{R}$:

$$\mathcal{V}_m(x) = w_Q Q(x) + w_V V(x) + w_S S(x) + w_G G(x) + w_R R_v(x) + w_K K(x) \quad (14)$$

$$- w_C C(x) - w_T T(x) - w_H H(x) - w_D D_P(x) - w_U I(x), \quad (15)$$

where Q is mission quality; V realized or conservatively attributable value; S sovereignty and strategic control; G governance integrity; R_v reliability and resilience; K future productive capacity; C complete cost; T time-to-accepted outcome; H human burden; D_P Proof Debt; and I irreversibility and unwind burden.

The **Mission Advantage Gradient** is the Riemannian gradient

$$\nabla_{G_m} \mathcal{V}_m(x) = G_m(x)^{-1} \nabla \mathcal{V}_m(x). \quad (16)$$

It identifies both the direction and the locally available rate of increase in mission value after institutional burden is taken into account.

The direction and magnitude of proof-adjusted superiority available over the incumbent and the best credible alternative, measured on the mission's constrained successor manifold.

3.4 Trajectory selection as a least-action problem

SEIZE does not merely rank endpoints. It underwrites a trajectory $\gamma : [0, 1] \rightarrow \mathcal{M}_{S,\text{adm}}$ from the incumbent x_0 to a candidate x_1 . Define the succession action

$$\mathcal{S}_m[\gamma] = \int_0^1 \left[\frac{1}{2} \dot{\gamma}^\top G_m(\gamma) \dot{\gamma} + \lambda_R \mathcal{R}(\gamma) + \lambda_A \mathcal{A}_r(\gamma) + \lambda_D D_P(\gamma) - \lambda_V \mathcal{V}_m(\gamma) \right] d\tau, \quad (17)$$

where $\mathcal{R}$ is residual risk and $\mathcal{A}_r$ is Authority-at-Risk. The underwritten trajectory is

$$\gamma^\star = \arg \min_{\gamma \in \Gamma_m} \mathcal{S}_m[\gamma] \quad \text{subject to all hard gates.} \quad (18)$$

This is an institutional analogue of a least-action principle, not a claim that organizational succession literally obeys fundamental physics. Its value is operational: it forces GoalOS to price the full path, not just an attractive endpoint.

3.5 Phase transitions, metastability and the succession event

The incumbent may remain a local equilibrium even after a superior state becomes reachable. Switching costs, incomplete evidence, authority inertia and path dependence can make it metastable. Let $b^\star$ be the best credible alternative set and define the order parameter

$$\Phi_m(x) = \mathcal{V}_m(x) - \max_{b \in \mathcal{B}_m} \mathcal{V}_m(b). \quad (19)$$

A **Mission Superiority Crossover** occurs when $\Phi_m(x) > \theta_m$ on protected fresh work. A **Mission-Sovereign Phase Transition** requires more:

$$\Sigma_m(x) = \mathbf{1}[\Phi_m(x) > \theta_m] \cdot \mathbf{1}[\text{FreshProof}] \cdot \mathbf{1}[\text{HumanAdmission}] \cdot \mathbf{1}[\text{AuthorityEnvelope}] \cdot \mathbf{1}[\text{RollbackReady}]. \quad (20)$$

Only when $\Sigma_m(x) = 1$ may the state be called Mission-Sovereign, and even then only for the stated mission, version, duration, jurisdiction, data classes and external effects.

Manifold law

The frontier is not the destination. It opens the Successor Manifold. GoalOS must still identify the mission-value field, choose an admissible trajectory, purchase decisive proof and separate empirical superiority from institutional authority.




Figure 1 – The canonical Mission-Sovereign succession progression.

4 Specialist ASI and the Mission-Sovereign Successor

4.1 A mission-bounded, system-level definition

A Specialist ASI is not defined by parameter count, brand, generality or self-description. It is a complete mission-specific intelligence institution that demonstrates superhuman and frontier-superior performance on a constituted mission family while remaining economically and constitutionally admissible. This mission-bounded definition separates performance depth, breadth and autonomy rather than collapsing them into one label, consistent with the need for more precise capability ontologies (Morris et al., 2024).

A complete mission-specific intelligence system that exceeds the incumbent and the best credible general, specialist, human or hybrid alternative by a preregistered positive margin on fresh, protected and representative work, after quality, critical-error risk, complete economics, human burden, rights, sovereignty, resilience and governance are counted.

The comparison is therefore

$$\text{Complete Specialist Successor Institution} \quad \text{versus} \quad \text{Best Credible Alternative Set}, \quad (21)$$

not simply “small model versus large model.” A successor may combine compact and frontier models, agents, retrieval, code, deterministic checks, proprietary data, simulators, humans, verifiers, monitoring, rollback and workflow integration.

4.2 The best credible alternative set

For mission m , define

$$\mathcal{B}_m = \{\text{incumbent, frontier stack, existing specialist, human team, human-AI hybrid, software, partner, acquisition, abandonment}\}. \quad (22)$$

A candidate becomes mission-dominant only when

$$\mathcal{V}_m(x_{\text{candidate}}) > \max_{b \in \mathcal{B}_m} \mathcal{V}_m(b) + \theta_m \quad (23)$$

and every hard gate passes. The baseline may not be deliberately weakened to manufacture apparent alpha.

4.3 Six superiority dimensions

Table 2 – Mission-Dominant Specialist ASI dimensions.

Dimension	Required evidence
Capability superiority	Better accepted mission outcomes and critical-error performance on fresh representative work.
Economic superiority	Lower complete lifecycle burden or materially greater verified value than the best alternative.
Reliability superiority	Better calibration, abstention, exception handling, resilience, recovery and drift control.
Sovereignty superiority	Stronger control of strategic data, rights, portability, continuity and provider optionality.

Dimension	Required evidence
Governance superiority	Clearer identity, authority, evidence, validation, acceptance, rollback, expiry and revocation.
Transfer superiority	Stable contribution across a preregistered mission family, not one memorized benchmark.

4.4 Dominance is not sovereignty

Mission dominance is an empirical verdict. Mission sovereignty is a constitutional state. The canonical progression is

$$\begin{aligned}
 &\text{Candidate Successor} \rightarrow \text{Verified Specialist Intelligence} \\
 &\quad \rightarrow \text{Mission-Dominant Specialist ASI} \\
 &\quad \rightarrow \text{Mission-Sovereign Successor}
 \end{aligned} \tag{24}$$

The final transition requires a separate accountable admission, an exact Authority Envelope, version and expiry, monitoring, rollback, revocation and the ability to re-prove the capability. A Mission-Sovereign Successor is not legally sovereign, self-owning or entitled to expand its own mandate. “Sovereign” refers to the institution’s control over the mission and the bounded authority granted by the principal.

Mission-sovereignty law

Fresh proof may establish dominance. Only accountable admission may grant mission authority. No candidate, producer, model or agent may validate itself, admit itself, expand its own Authority Envelope or enable its own production release.

4.5 The Specialist ASI claim boundary

The term Specialist ASI is a target state and earned designation. The present public-static GoalOS + SEIZE institution provides mission constitution, architecture search, underwriting, bounded-job compilation, proof governance, Chronicle and successor logic. It does not by itself prove that any external Specialist ASI has been manufactured or that any customer has realized mission alpha. Such claims require exact releases, protected fresh-work evidence, complete-cost accounting, independent review and accountable acceptance.

5 Game-Theoretic Constitution and Mechanism Design

5.1 Succession as a dynamic game

Verified succession is a dynamic game under incomplete information, drawing on the classical distinction between strategic games, equilibrium and Bayesian information structures (von Neumann and Morgenstern, 1944; John F. Nash, 1950; Harsanyi, 1967–1968; Osborne and Rubinstein, 1994). It also contains agency, adverse-selection and incomplete-contracting problems: producers possess private information, incumbents may resist replacement, providers may seek lock-in, evaluators may economize on effort, and principals may be pressured by adoption races (Akerlof, 1970; Jensen and Meckling, 1976; Holmstrom, 1979; Williamson, 1985). Let

$$\mathcal{G}_m = (N, \Sigma, u, I, \mathcal{H}) \tag{25}$$

where N contains the accountable principal, incumbent operator, candidate producer, independent verifier, admitter, affected parties and external providers; Σ contains their available strategies; u their mission-dependent payoffs; I the information structure; and $\mathcal{H}$ the authority and evidence history.

The producer may prefer promotion, the incumbent may prefer retention, a provider may prefer lock-in, an evaluator may prefer low effort and the principal may prefer speed. GoalOS cannot assume these incentives align. It must design the mechanism so that unsupported promotion, concealed failure and unauthorized action are unattractive or impossible.

5.2 Separation of powers as an incentive constraint

The desired mechanism imposes:

$$\text{Producer} \neq \text{Verifier}, \quad (26)$$

$$\text{Verifier} \neq \text{Admitter}, \quad (27)$$

$$\text{Candidate} \not\in \{\text{protected cases, scorer internals, release credentials}\}, \quad (28)$$

$$\text{AccessReceipt} \neq \text{AuthorityEnvelope}, \quad (29)$$

$$\text{Acceptance} \neq \text{ChronicleAdmission}. \quad (30)$$

These are not procedural decorations. They alter the payoff structure and address principal–agent and observability failures (Holmstrom, 1979). A producer cannot improve expected utility by hiding failures when the Evidence Docket records retries, human rescue and full cost. A verifier cannot silently release the candidate. An admitter cannot treat access eligibility as mission authority.

5.3 Mechanism-design objective

GoalOS seeks a governance-preserving equilibrium that maximizes verified mission value while keeping downside and authority bounded:

$$\max_{\mathcal{M}} \mathbb{E}[\mathcal{V}_m(x)] - \lambda_R \text{CVaR}_\alpha(L) - \lambda_A \text{AuthorityAtRisk} - \lambda_D D_P \quad (31)$$

$$\text{subject to } \text{IC, IR, EvidenceSeparation, Rights, Rollback, HumanAdmission}, \quad (32)$$

where IC is incentive compatibility and IR is participation rationality for legitimate actors. The objective is not a claim that a perfect equilibrium can be computed. It is a design standard: the institution should reward truthful evidence, penalize hidden burden, preserve refusal and make unauthorized self-promotion structurally difficult.

5.4 Regret-aware choice and adversarial challenge

For candidate a and possible world ω , let $L(a, \omega)$ be institutional loss. Minimax regret is

$$\text{Regret}(a) = \max_{\omega \in \Omega} \left[L(a, \omega) - \min_{a' \in A} L(a', \omega) \right]. \quad (33)$$

GoalOS may prefer a lower-mean but more reversible successor when evidence is weak or tail loss is severe. A reserve option is rational when the value of waiting exceeds the cost of delay. Adversarial institutional agents, red-team tests, anti-recommendations and stop conditions are therefore part of the mechanism, not obstacles to speed.

5.5 Proof capital as an option premium

SEIZE treats Experiment Zero and decisive evidence as the premium on a real option, consistent with the logic of contingent claims and investment under uncertainty (Black and Scholes, 1973; Merton, 1973; Dixit and Pindyck, 1994). The institution pays a bounded amount C_e to observe evidence e before committing to full formation. The Value of Proof is

$$\text{VoP}(e) = \mathbb{E}_e \left[\max_{a \in A} \mathbb{E}[U(a) \mid e] \right] - \max_{a \in A} \mathbb{E}[U(a)] - C_e. \quad (34)$$

Proof capital is rational only when it can change the decision, reduce expected loss, preserve valuable optionality or establish transferable capability. More experimentation is not automatically better.

Game-theoretic law

The succession mechanism must remain valuable even when participants have different incentives, private information and asymmetric power. Truth must be easier to preserve than theatre; refusal must remain available; and no actor may be both claimant and final judge of its own authority.

6 The Six-Stage Mission-Sovereign Progression

Table 3 – Canonical six-stage progression.

Stage	Institutional meaning
Successor Manifold	The newly reachable space of complete successor architectures opened by a frontier, cost, market, legal, security or infrastructure discontinuity.
Mission Advantage Gradient	The direction and magnitude of proof-adjusted superiority available over the incumbent and the best credible external alternative.
SEIZE Underwriting	The selection of a valuable, feasible, governable trajectory and the proof-capital budget required to test it.
Bounded AGI Jobs	Explicit work-and-proof contracts that constitute models, data, tools, verifier, workflow, evidence, controls, integration and rollback.
Specialist ASI	A mission-specific complete intelligence system that earns superiority on fresh work across capability, economics, reliability, sovereignty and governance.
Mission-Sovereign Successor	The Mission-Dominant Specialist ASI that is separately admitted by accountable authority and receives only scoped, versioned, expiring and revocable mission authority.



Figure 2 – Mission-Sovereign successor reference visual. Exact operative terms and claims are governed by the typeset framework.

7 AI Beta, Mission Spread, Proven Alpha, and the Sovereign Successor

7.1 The four objects must not be conflated

The doctrine can be read at three speeds. In one second: *Beta* $\rightarrow$ *Spread* $\rightarrow$ *SEIZE* $\rightarrow$ *Specialist ASI* $\rightarrow$ *Mission-Sovereign Successor*. In five seconds: the institution rents a moving market baseline, underwrites the distance to a superior mission architecture, proves the complete candidate, and grants only the authority supported by current evidence. In thirty seconds: the reader discovers a dynamic game involving architecture search, real options, principal–agent separation, comparative proof, property rights, vendor hold-up, Chronicle memory, Alpha decay and repeated institutional renewal.




Figure 3 – The strategic flow. Alpha is depicted as the residual value field generated by comparative proof, not as an independent chronological stage.

The broadly available uplift supplied by frontier models, standard agents, common tools, public data and widely purchasable implementation capability. Beta may be extraordinarily powerful while remaining symmetric, rentable and increasingly incorporated into the competitive baseline.

The gross proof-adjusted distance between a candidate successor and the strongest relevant reference architecture before construction, operating, proof, risk, dependency and unwind burdens are fully deducted.

The residual risk-adjusted mission surplus created by a complete successor after the incumbent, the best credible external alternative, complete cost, tail risk, human burden, dependency, rights, Proof Debt and reversibility are counted.

The complete Specialist ASI institution that passes fresh proof and accountable admission, then receives scoped, versioned, monitored, expiring and revocable authority over a constituted mission. Sovereignty denotes institutional control of the mission, architecture, data, evidence, memory, permissions, portability and economics—not unlimited autonomy or legal personhood.

7.2 AI Beta as a moving common factor

Let F_t denote the generally rentable frontier at time t . For institution i , an initial performance decomposition is

$$Y_{i,m,t} = \beta_{i,m} F_t + \alpha_{i,m,t} + \varepsilon_{i,m,t}, \quad (35)$$

where $\beta_{i,m}$ is the institution's ability to absorb common frontier capability, $\alpha_{i,m,t}$ is residual mission-specific advantage, and $\varepsilon_{i,m,t}$ captures noise, shocks and execution variance. The analogy is managerial rather than a securities-pricing model. Its purpose is to distinguish common uplift from institution-specific surplus.

As a frontier capability diffuses, it may raise the productivity of every competent adopter while reducing its power to differentiate them. Standard access can therefore be strategically necessary and economically valuable without constituting proprietary Alpha. Agent benchmarks also show why access alone is insufficient: tool-using systems

can perform strongly in some settings while remaining inconsistent over repeated trials, fragile on long-horizon tasks, or sensitive to evaluation design (Mialon et al., 2024; Jimenez et al., 2024; Yao et al., 2024; Kwa et al., 2025).

Design Proposition 1 — Beta compression

Suppose n competitors obtain the same frontier capability, possess equivalent complementary assets, face the same mission and can imitate one another at negligible cost. Then the frontier uplift may increase absolute productivity, but its expected relative surplus converges toward zero as prices, customer expectations and operating standards adjust.

Proof sketch. Under symmetry, no competitor possesses a private state variable that shifts mission utility or cost relative to the others. The common factor changes the feasible set for all players; competitive response transfers part of the surplus to customers, labour, complementary suppliers or the frontier provider. Relative advantage requires asymmetry in data, architecture, verification, distribution, rights, speed of learning or authority design.

7.3 The reference architecture and the investable Spread

For mission m , let $I_m(t)$ be the incumbent and $B_m(t)$ the best external alternative available at time t . Let $U_m(\cdot)$ be a complete mission utility function subject to non-negotiable constitutional gates. The reference architecture is

$$R_m(t) = \arg \max_{z \in \{I_m(t), B_m(t)\}} U_m(z). \quad (36)$$

A candidate architecture $S_m(x, t)$ is a point x on the Successor Manifold. Its gross Mission Advantage Spread is

$$\Gamma_m(x, t) = U_m(S_m(x, t)) - U_m(R_m(t)). \quad (37)$$

The Spread is an opportunity, not yet an asset. It can disappear when complete cost, proof burden, tail loss, dependence, integration, human rescue, maintenance, rights or unwind are counted. It can also be negative even when the candidate appears more capable on a narrow benchmark.

Design Proposition 2 — Spread is not Alpha

A positive gross Spread is necessary but not sufficient for Mission Alpha. Alpha exists only if the candidate's residual value remains positive after every material burden and hard gate is applied against the strongest reference architecture.

The Spread is naturally vector-valued:

$$\Gamma_m(x, t) = \begin{bmatrix} \Delta Q & -\Delta C & -\Delta T & \Delta R & \Delta E & \Delta S & \Delta K \end{bmatrix}^T, \quad (38)$$

where Q is accepted quality, C complete cost, T time, R reliability, E evidence quality, S sovereignty and K compounding capacity. A scalar decision score may be formed only after weights, thresholds and hard constraints are preregistered.

7.4 Mission Alpha as residual surplus

All additive terms below must be expressed in a common mission-utility unit or transformed through a preregistered multi-criteria decision rule with explicit hard gates. Capability, cost and risk must not be combined by post hoc weights chosen after the result is observed.

The compact definition is

$$\alpha_m(x, t) = U_m(S_m(x, t)) - \max\{U_m(I_m(t)), U_m(B_m(t))\} - C_m(x, t) - R_m(x, t) - D_m(x, t) \quad (39)$$

where C_m is complete construction and lifecycle cost, R_m is risk including tail and authority exposure, and D_m is unresolved Proof Debt and dependency burden. A more explicit decomposition is

$$\alpha_m = V_{\text{quality}} + V_{\text{speed}} + V_{\text{reliability}} + V_{\text{sovereignty}} + V_{\text{option}} + V_{\text{learning}} \quad (40)$$

$$- C_{\text{build}} - C_{\text{run}} - C_{\text{proof}} - C_{\text{human}} - C_{\text{integration}} - C_{\text{maintenance}} - C_{\text{unwind}} \quad (41)$$

$$- \lambda_R \text{CVaR}_q(L) - D_{\text{proof}} - D_{\text{dependency}} - I_{\text{irreversibility}}. \quad (42)$$

Conditional Value-at-Risk is useful when average performance conceals severe downside (Rockafellar and Uryasev, 2000). Alpha may be positive because the successor is more accurate, but it may also arise from lower complete cost, better abstention, stronger evidence, improved continuity, lower vendor dependence or superior learning across repeated missions.

Alpha is a measured state of residual advantage. It emerges only after the candidate passes through underwriting, formation and comparative proof. Specialist ASI is the proven capability state; the Mission-Sovereign Successor is the governed institutional state that captures, preserves and compounds the resulting Alpha.

7.5 The Mission Advantage Gradient

Let x contain the models, data, tools, workflow, verifier, rights, integration, evidence, Chronicle and Authority Envelope. The Mission Advantage Gradient is

$$g_m(x, t) = \nabla_x \alpha_m(x, t). \quad (43)$$

When architecture coordinates are discrete, categorical or combinatorial, ∇_x denotes a local surrogate, directional finite difference or policy-search signal rather than an assumption that the entire architecture space is smoothly differentiable. The manifold language is a disciplined representation of constrained possibility, not decorative physics.

The gradient need not point toward a larger model. It may point toward higher-quality proprietary examples, a stronger verifier, narrower permissions, lower latency, local inference, deterministic checks, improved exception handling, better distribution or model substitutability. The objective of the Navigator is to find the architecture change with the greatest expected marginal Alpha per unit of proof capital, not the most spectacular technology.

If $G(x)$ is an institutional burden metric on the manifold, a locally disciplined trajectory $x(\tau)$ minimizes

$$\mathcal{J}[x] = \int_0^1 \left[\frac{1}{2} \dot{x}^\top G(x) \dot{x} - \eta \alpha_m(x, t) + \lambda_R \mathcal{R}(x) + \lambda_P D_P(x) \right] d\tau, \quad (44)$$

subject to constitutional gates. The physics language is an exact design analogy: the selected path should reach higher mission value while avoiding unnecessary institutional distance, irreversibility and risk.

7.6 SEIZE as Bayesian real-options underwriting

SEIZE selects a trajectory π and proof plan e under incomplete information. The institution does not pay to know everything. It buys the smallest evidence package capable of changing the decision, reducing expected loss or preserving a valuable option. Let $a \in \{\text{retain, repair, rent, build, partner, acquire, defer, stop}\}$. Then

$$\text{VoP}(e) = \mathbb{E}_e \left[\max_a \mathbb{E}[U(a) \mid e] \right] - \max_a \mathbb{E}[U(a)] - C_e. \quad (45)$$

SEIZE proceeds only when the expected value of proof is positive and the proof can be obtained without violating rights or exposing unacceptable tail risk. The real-options interpretation follows investment-under-uncertainty

theory: preserve reversibility, delay irreversible commitment until decisive information arrives, and exercise only after the posterior case exceeds the admission threshold (Dixit and Pindyck, 1994).

7.7 Bounded AGI Jobs as mechanism design

The selected trajectory is compiled into Bounded AGI Jobs. Each job specifies objective, inputs, tools, action space, prohibited actions, acceptance tests, evidence, verifier, escalation, timeout, rollback and compensation. This converts a vague request for autonomy into an explicit mechanism.

The role system separates the producer, critic, verifier, governor and accountable human principal. That separation responds to moral hazard, adverse selection and hidden action (Akerlof, 1970; Jensen and Meckling, 1976; Holmstrom, 1979). Agent-governance work similarly emphasizes task suitability, action-space constraints, approval, legibility, monitoring, attribution and interruptibility (Shavit et al., 2023); control research demonstrates the value of protocols that remain useful even when a more capable component is treated as untrusted (Greenblatt et al., 2024).

$$\text{Producer} \neq \text{Verifier} \neq \text{Admitter}, \quad \text{Access} \neq \text{Authority}, \quad \text{Execution} \neq \text{Acceptance} \neq \text{Admission}. \quad (46)$$

7.8 Specialist ASI as a comparative proof state

A candidate qualifies as Specialist ASI for mission family $\mathcal{F}_m$ only when

$$\Pr(\alpha_m > \theta_m \mid E_t) \geq \tau_m, \quad (47)$$

$$\text{CVaR}_q(L_m) \leq L_m^{\max}, \quad (48)$$

$$\text{EvidenceCoverage}_m \geq e_m^{\min}, \quad (49)$$

$$C_m(x) = 1, \quad (50)$$

where C_m contains the hard rights, security, privacy, authority, reversibility and governance gates. Evaluation must use fresh, protected and representative work, not only public benchmarks or examples available during formation. Research on model and agent evaluation underscores the need to evaluate dangerous capabilities, end-to-end task completion, policy adherence and reliability rather than infer safety or competence from general benchmark scores (Shevlane et al., 2023; Phuong et al., 2024; Jimenez et al., 2024; Yao et al., 2024).

The terminology is intentionally orthogonal to generality. Existing work distinguishes performance, breadth and autonomy when discussing AGI levels (Morris et al., 2024); recent analysis of possible AGI-to-ASI pathways likewise emphasizes large uncertainty, multiple routes and potentially repeated transformations rather than one guaranteed discontinuity (Genewein et al., 2026). GoalOS therefore makes the narrower claim that a complete system can earn mission-specific superiority without asserting universal superintelligence.

7.9 Proof as collateral for authority

Let E_t be the current evidence state and A_t the Authority Envelope. The governing constraint is

$$A_t \preceq \phi(E_t) \quad (51)$$

where $\preceq$ is a partial order over consequence, scope, duration, reversibility and permissions. Evidence does not merely support a claim. It collateralizes the authority the institution is willing to expose.

Design Proposition 3 — Authority never exceeds fresh proof

For any consequential action class, the admissible Authority Envelope is monotone in current evidence but capped by the mission's constitutional ceiling. Historical success cannot justify authority after the evidence has expired, the system has materially changed, or the environment has shifted.

A rational ladder is

$$\text{Observe} \rightarrow \text{Advise} \rightarrow \text{Prepare} \rightarrow \text{Execute with approval} \rightarrow \text{Execute within a bounded envelope.} \quad (52)$$

No successor is required to reach the final state. The value-maximizing authority level may remain advisory when consequence, reversibility or proof economics make greater delegation irrational.

7.10 Sovereignty and the vendor hold-up game

A corporation can rent models while owning the successor. Ownership in this sense means control of the mission constitution, job portfolio, proprietary complements, evidence, Chronicle, permissions, portability and economic capture. This is a property-rights and incomplete-contracting problem. When the institution makes relationship-specific investments inside one provider's platform, the provider may acquire bargaining power over the resulting quasi-rents (Williamson, 1985; Grossman and Hart, 1986; Hart and Moore, 1990).

Define the sovereignty premium as

$$\Pi_{\text{sov}} = \mathbb{E}[\text{avoided hold-up} + \text{switching option} + \text{continuity} + \text{retained learning} + \text{rights clarity}] - C_{\text{sov}}. \quad (53)$$

Design Proposition 4 — Portability changes bargaining power

Holding model quality constant, a modular successor with exportable evidence, model-independent tests, portable jobs and retained Chronicle state has greater option value and lower expected hold-up loss than an otherwise identical mission architecture trapped inside one provider.

Sovereignty is therefore not autarky. It is the institution's ability to substitute components without losing its mission, proof, memory, authority system or operating continuity.

7.11 The repeated game and the Chronicle state

A one-time model call creates an output. A successor participates in a repeated game. Let K_t be validated institutional capability stored in the Chronicle. Then

$$K_{t+1} = (1 - \delta_K)K_t + \eta E_t^+ - \omega D_t - \chi F_t, \quad (54)$$

where E_t^+ is admitted evidence, D_t unresolved Proof Debt and F_t uncontained failure. The Chronicle creates path dependence because accepted cases, failure maps, verifiers, environments, rights records and authority history become inputs to the next succession cycle. Dynamic capabilities arise from the repeated ability to sense, seize and transform rather than from one static asset (Teece et al., 1997).

The dynamic Alpha equation is

$$\frac{d\alpha_m}{dt} = \underbrace{L_{\text{institution}} + G_{\text{proprietary}} + V_{\text{option}}}_{\text{renewal and compounding}} - \underbrace{Q_{\text{commoditization}} + I_{\text{imitation}} + \Delta_{\text{drift}} + M_{\text{maintenance}}}_{\text{Alpha compression}} \quad (55)$$

Design Proposition 5 — Alpha persistence condition

Durable Mission Alpha requires the successor's verified learning and reproof velocity to exceed the combined velocity of frontier diffusion, competitive imitation, environmental drift and maintenance burden.

Equivalently,

$$v_{\text{learning}} + v_{\text{reproof}} > v_{\text{commoditization}} + v_{\text{imitation}} + v_{\text{drift}} + v_{\text{decay}}. \quad (56)$$

7.12 From the Beta equilibrium to the successor equilibrium

The competitive game has two stylized equilibria. In the **Beta equilibrium**, firms converge on similar models, prompts and common agent stacks. Absolute productivity rises, but differentiation compresses and correlated dependency grows. In the **successor equilibrium**, one institution creates a private state through proprietary mission architecture, rights-cleared evidence, validated jobs, distribution, Chronicle memory and governed authority. The result is asymmetric and harder to imitate because copying the visible interface does not recreate the historical proof state.

Let firm i choose a complement stock k_i at cost $c(k_i)$ after renting frontier capability F_t . Its payoff is

$$\pi_i = u_m(F_t, k_i, K_{i,t}) - p(F_t) - c(k_i) - \ell_i, \quad (57)$$

where $K_{i,t}$ is Chronicle capital and ℓ_i is expected loss. A successor equilibrium exists when the expected private return from k_i and $K_{i,t}$ exceeds formation and proof cost, remains protected by rights and control, and cannot be instantaneously arbitrated away by competitors.

Design Proposition 6 — The successor is the Alpha-capture institution

Specialist ASI establishes that the candidate wins the mission game under current proof. The Mission-Sovereign Successor is the institutional wrapper that captures the resulting Alpha, constrains its authority, preserves its evidence and memory, and exposes it to continuous reproof. Without that wrapper, Alpha may leak to vendors, employees, transient chats, unowned data or uncontrolled agents.

7.13 The canonical Alpha stack

Table 4 – The complete Alpha–Beta–Spread–Successor stack.

Layer	Institutional meaning
Business Beta	The institution's underlying market, mandate, capital, customer and industrial exposure.
AI Beta	Common frontier capability and cost improvement available to competent adopters.
Sensing Alpha	Detecting a succession event before it is fully incorporated into institutional assumptions.
Spread	The gross opportunity between the strongest reference architecture and a reachable superior successor.
Underwriting Alpha	Selecting the right mission, trajectory, proof budget, stop-loss and preserved options.
Formation Alpha	Constituting a proprietary complete system instead of repeatedly renting an undifferentiated interface.
Proof Alpha	Reducing uncertainty, hidden burden and tail risk through protected comparison, independent verification and exact release control.
Operating Alpha	Recurring quality, speed, cost, reliability, sovereignty, evidence or revenue advantage.
Compounding Alpha	Retained learning from Chronicle, validated jobs, failure maps, environments, integrations and realized outcomes.
Sovereignty Premium	Option value from control of mission, architecture, data, evidence, memory, authority, portability and economics.
Alpha Compression	Erosion from frontier diffusion, imitation, provider improvement, mission drift, scorer expiry, maintenance and crowding.

Layer	Institutional meaning
Mission-Sovereign Successor	The admitted institution that captures, governs and renews Mission Alpha under fresh proof.

The Alpha doctrine

Beta is rented. The Spread is underwritten. Alpha is manufactured through differentiated architecture and comparative proof. Specialist ASI establishes mission superiority. The Mission-Sovereign Successor captures, governs and compounds that advantage while authority remains collateralized by fresh evidence.

8 The Successor Book

8.1 A live market for institutional architectures

GoalOS should maintain a live internal market rather than a single recommendation. The portfolio logic follows the principle that return, uncertainty and covariance among alternatives matter jointly rather than one score in isolation (Markowitz, 1952). The **Successor Book** records the incumbent and the portfolio of possible replacements, augmentations, hedges and retirements.

Table 5 – Successor Book positions.

Position	Meaning	Capital posture
Incumbent	Current accepted architecture and the assumptions that justify it.	Operate, monitor, mark for impairment and maintain rollback.
Immediate successor	Strong verifier, frozen mission, clear rights and short proof path.	Fund decisive proof now.
Challenger	Plausibly superior after one or two named catalysts.	Maintain bounded evidence programme.
Reserve successor	High upside but long duration, weak present evidence or incomplete authority.	Preserve option at low cost.
Hedge architecture	Fallback against provider, security, law, infrastructure or market failure.	Maintain operational readiness and periodic drills.
Retirement candidate	Negative expected value, expired rights, unacceptable risk or strategic obsolescence.	Contain, migrate, revoke or exit.

8.2 Three horizons and option discipline

The successor portfolio is organized across three horizons:

- **Immediate:** enough evidence exists to authorize bounded proof now;
- **Challenger:** the thesis becomes investable after a named catalyst or evidence object;
- **Reserve:** the institution preserves a low-cost option against a large but uncertain future state.

This is real-options discipline applied to institutional architecture (Dixit and Pindyck, 1994). The simultaneous maintenance of incumbents, challengers and reserves also operationalizes the classic exploration-exploitation

tension (March, 1991). Proof capital is the option premium. Full formation capital is committed only when the information purchased by the proof stage justifies exercise.

Let p_i be the estimated probability that candidate S_i passes, V_i its Proof-Adjusted value if successful, C_i^P the cost of decisive proof, C_i^F the formation cost, R_i downside and O_i option value. A simplified underwriting score is

$$U_i = p_i V_i + O_i - C_i^P - C_i^F - R_i - D_i, \quad (58)$$

where D_i includes duration and decay. The score is a decision aid, not an automatic capital authority.

8.3 Incumbent impairment

Every incumbent should be re-underwritten when any of the following occurs:

- a material capability or cost event;
- a source or right expires;
- a security or safety incident;
- a provider, policy or market regime changes;
- a challenger crosses a preregistered expected-value threshold;
- the capability half-life falls below the next planned review;
- the incumbent fails a constitutional invariant;
- realized value materially diverges from the accepted thesis.


A re-underwriting trigger is not a forced sale. It opens a succession review.

What changed, which incumbent assumption is now impaired, which successor options became investable, what evidence would change the position, and what is the maximum capital or authority that may be exposed before that evidence arrives?

9 SEIZE: The Succession Underwriting Microcycle

9.1 Purpose

SEIZE is the microcycle that converts a capability event into one bounded promotion decision. It does not authorize production deployment. It structures the succession opportunity so that the institution can buy decisive information before committing irreversible capital or authority.



9.2 S - Surface the succession event

The output is a **Succession Event Record** containing:

- source, date, provenance and replication status;
- capability, cost, market, law, security or infrastructure change;
- affected incumbent assumptions;
- missions newly feasible, threatened or obsolete;
- cost of waiting and cost of premature action;
- evidence required before capital or authority changes;
- accountable owner and review deadline.

The event is not admitted because it is exciting. It is admitted because it changes the expected value of an institutional decision.

9.3 E - Evaluate the successor frontier

The institution generates structurally different alternatives rather than small variations of one preferred idea. Candidate classes should include, where relevant:

1. retain the incumbent;
2. repair or augment the incumbent;
3. rent an external capability;
4. constitute an owned mission capability;
5. form a human-machine successor team;
6. reorganize the workflow or objective;
7. partner, license or acquire;
8. preserve a reserve option;
9. maintain a hedge architecture;
10. retire or prohibit the mission.

Each candidate receives an anti-recommendation: the strongest reason not to pursue it, the evidence that would falsify it, and the condition under which another candidate should replace it.

9.4 I - Instantiate the succession constitution

The **Succession Constitution** freezes:

- the exact mission and beneficiary;
- the incumbent and best credible alternative;
- the quality floor and unacceptable failure;
- the evaluation distribution and blind custodian;
- the complete cost and burden denominator;
- rights, confidentiality and jurisdiction;
- the Authority Envelope;
- proof capital, formation capital and stop-loss;
- the human principals for Underwrite, Execute, Accept, Admit and Allocate;
- expiry, rollback and requalification triggers.

No candidate may modify these elements after observing protected results.

9.5 Z - Zero in on decisive proof

The institution should buy evidence in descending order of information value. It should not fund a complete successor when a cheaper test can reveal that the thesis is wrong.

Let D be a proposed diagnostic. Its expected value of information is

$$\text{EVI}(D) = \mathbb{E} \left[\max_a U(a \mid D) \right] - \max_a \mathbb{E}[U(a)] - C(D), \quad (59)$$

where a ranges over successor actions and $C(D)$ is the complete cost of the diagnostic. The next diagnostic should maximize positive EVI subject to rights, time and safety constraints.

9.6 E - Elevate one challenger

One complete challenger is frozen for independent proof. The freeze includes:

- exact architecture and versions;
- data, tools, prompts, policies and runtime;
- authorities and credentials;
- resource and cost envelope;
- verifier and monitoring configuration;
- hashes, provenance and rollback target.


SEIZE ends with a **Promotion Request**. It does not issue a Release Certificate.

A candidate may propose experiments, generate descendants, run approved sandbox work and request promotion. It may not inspect protected cases, alter the constitution, classify its own risk, certify itself, hold release credentials, approve descendants, install itself or expand its Authority Envelope.

10 The Succession Constitution and Separation of Powers

10.1 Five distinct decisions

A major failure of ordinary AI programmes is the collapse of several decisions into one. A strong demonstration becomes permission to execute; execution becomes acceptance; acceptance becomes institutional memory; projected value becomes reinvestment authority. GoalOS separates the chain:



Each decision should be attributable to a named human or institutional principal. The same actor need not control all five, and the candidate controls none of them.

10.2 The Authority Envelope

Every job and successor requires a machine-readable Authority Envelope specifying:

- principal and accountable owner;
- permitted objectives and prohibited objectives;
- allowed data, tools, identities, accounts and network paths;
- transaction, spend, time and compute limits;
- actions requiring additional approval;
- stop conditions, kill switches and escalation;
- evidence that must accompany consequential action;
- expiry, revocation and rollback.

Capability never expands authority automatically. A successor may produce better results and still remain at the same authority level until a separate human decision changes the envelope.

10.3 Authority inheritance

Authority is not transferred because one system replaces another. It is reconstituted. The successor must earn the right to inherit each class of action.

Table 6 – Authority inheritance levels.

Level	Permitted posture	Typical evidence requirement
A0 Observe	Read, analyze and propose.	Source provenance, logging and no consequential action.
A1 Recommend	Produce ranked options for accountable review.	Calibrated evidence, alternatives, uncertainty and conflict disclosure.
A2 Sealed sand-box	Execute in an isolated environment with no external effect.	Frozen inputs, budgets, audit log and repeatability.
A3 Reversible bounded action	Take preapproved actions with explicit rollback.	Canary proof, transaction limits, monitoring and stop conditions.
A4 Consequential action	Affect money, rights, safety, public systems or material operations.	Independent validation, professional authority where required, dual control and incident response.

A successor can be mission-dominant at A1 while remaining ineligible for A3 or A4. Performance and authority are separate state variables.

10.4 Proof-carrying action

A consequential action should carry a proof receipt containing at minimum:

- action identity, actor and principal;
- mission and Authority Envelope;
- source and input lineage;
- model, tool, policy and runtime versions;
- validation status and residual uncertainty;
- approval and escalation record;
- expected effect, rollback and expiry;
- hash and timestamp.

This turns proof from a retrospective report into a condition of action.

10.5 Institutional refusal

The Verified Succession Institution must be able to refuse:

- a mission whose objective is vague or conflicted;
- an evaluator that is exploitable or not independently knowable;
- a data source without adequate rights;
- a candidate that depends on hidden authority or hidden external capability;
- a promotion request with incomplete resource truth;

- a release without rollback or accountable ownership;
- a reinvestment claim based only on projected savings;
- a successor whose gain requires governance regression.

Refusal is not paralysis. It is a bounded decision to preserve the institution until the missing condition can be proven or the thesis is abandoned.

11

Succession Due Diligence: Purchasing Decisive Evidence

11.1

The causal bottleneck map

Before forming a successor, GoalOS diagnoses why the incumbent underperforms or why the candidate appears superior. The bottleneck may reside in:

- objective ambiguity;
- missing or low-quality evidence;
- representation and ontology;
- decomposition and planning;
- coordination among specialists;
- search and refinement;
- environment or tool design;
- verifier quality;
- workflow integration;
- rights, authority or human acceptance;
- actual capability or physical constraint.

The due-diligence programme should isolate these causes rather than prescribing one technique in advance.

11.2

Five mandatory diagnostic families

The original Experiment Zero becomes a mechanism-agnostic **Succession Due-Diligence Sequence**.

Table 7 – Mandatory succession due diligence.

Run	Condition	Decision value
SDD.1	Exact incumbent and reference reproduction	Determine whether the claimed succession spread is real, stable and attributable under identical conditions.
SDD.2	Oracle evidence or perfect information access	Separate information-access failure from inference, planning, coordination or capacity failure.
SDD.3	Oracle representation, decomposition or organization	Determine whether the bottleneck is how the mission is structured rather than the intelligence substrate itself.
SDD.4	Bounded repeated refinement, search or replay	Test whether additional verified computation, correction or exploration converts failure into success.

Run	Condition	Decision value
SDD.5	Strongest simple intervention	Establish what can be achieved by the least complex credible repair before funding advanced formation.

The sequence should end with one evidence-backed route:

- retain or repair the incumbent;
- improve evidence or representation;
- change workflow or organization;
- add verified search or recursive refinement;
- create specialists or a new coordination layer;
- use a different substrate or provider;
- build a hybrid successor;
- partner or acquire;
- stop.

11.3 Proof capital and stop-loss discipline

Proof capital is the maximum amount authorized to resolve uncertainty before a new underwriting decision. It is not a blank cheque for research.

Every due-diligence mandate should state:

- the exact decision it is meant to inform;
- the evidence object required;
- the maximum spend, time and authority;
- the expected value of information;
- the stop-loss and invalidation conditions;
- the next review date;
- the owner of the go, hold, repair or stop decision.

The institution should lose small and early when a succession thesis is weak. It should scale capital only after the mission, rights, verifier, architecture and economics become progressively more legible.

11.4 The evidence ladder

Evidence should progress through distinct levels:

1. **Claim:** a capability, value or risk is asserted;
2. **Replication:** the observation is reproduced under controlled conditions;
3. **Mission relevance:** the capability affects the exact institutional workflow;

4. **Causal diagnosis:** the principal bottleneck or contribution is isolated;
5. **Protected validation:** a frozen challenger passes fresh evaluation;
6. **Canary operation:** the result survives bounded real-world use;
7. **Realized contribution:** value is accepted and attributable;
8. **Fresh transfer:** the admitted capability improves a new mission or successor cycle.

The institution should label each claim by evidence depth rather than allowing one development result to imply all later levels.

11.5 Reward grounding

Open-ended search, self-play and adaptive curricula can manufacture useful experience, but only when the reward remains tied to an independently valuable mission. Outside clean games, a generator can discover tasks that are impossible, trivial, random or strategically irrelevant while still optimizing its formal objective.

Therefore:

- task generators may propose tasks, rewards and validators;
- they may not authorize those tasks for training or promotion;
- the answer must be independently knowable;
- the verifier must be tested for false acceptance and gaming;
- relevance, rights and blind-set separation must be confirmed;
- accepted training experience must retain provenance and scope.

This is the difference between open-ended proposal and open-ended authority.

12 The Successor Foundry

12.1 Formation is architecture search, not one technique

The Successor Foundry constitutes complete challengers. It is method-agnostic and can combine:

- supervised adaptation and verified demonstrations;
- reinforcement learning with independently grounded rewards;
- evolutionary search and quality-diversity archives;
- recursive refinement and test-time search;
- program synthesis, symbolic planning and executable world models;
- specialist teams and learned coordination;
- deterministic tools, simulators and formal methods;
- human judgment, professional authority and redesigned workflows;
- partnership, licensing, acquisition or infrastructure formation.

The method is selected by the diagnosed bottleneck, not by fashion. A successor portfolio must contain enough architectural variety to match the variety of its environment, consistent with the principle of requisite variety

(Ashby, 1956). Search procedures such as AlphaGo-style evaluation loops are most powerful when the mission has a bounded state, structured actions, rapid feedback and a trustworthy evaluator (Silver et al., 2016, 2017; Sutton and Barto, 2018). Open-ended search adds diversity and stepping stones but must remain constitutionally separated from release authority (Stanley and Lehman, 2015).

12.2 Candidate architecture classes

Table 8 – Successor Foundry architecture classes.

Class	Formation thesis	Typical proof question
Incumbent repair	Preserve the present system while removing the measured bottleneck.	Does targeted repair dominate replacement after migration cost?
External capability	Rent a stronger provider or service under bounded policy.	Is rental still superior after cost, rights, dependency and exit are counted?
Owned mission capability	Constitute a rights-cleared internal asset around the workflow.	Does ownership create durable operating or strategic value?
Executable successor	Move reasoning into programs, planners, constraints or simulation.	Can correctness be established more reliably outside open-ended generation?
Specialist organization	Coordinate several bounded capabilities under a router or manager.	Does organization create net gain after coordination burden and correlated error?
Human-machine institution	Redesign roles so people retain purpose, judgment and authority while machines execute bounded work.	Does the complete socio-technical architecture outperform either party alone?
Partnership or acquisition	Obtain capability, distribution or rights through another institution.	Is control, time-to-value and integration superior to internal formation?
Reserve or hedge	Preserve an option or fallback without immediate production use.	Does the option justify its carrying cost under named scenarios?
Retirement	End the workflow, capability or dependency.	Is cessation superior to continued operation or replacement?

12.3 Quality diversity and the archive

The Foundry should not retain only the candidate with the highest average development score. It should preserve lineages that are:

- safest;
- most reversible;
- lowest burden;
- strongest on transfer;
- best calibrated;
- most private or sovereign;
- fastest;
- most resilient;
- best under scarce resources;
- strongest specialist for one failure cluster.

Quality-diversity methods preserve stepping stones that may later combine into a superior successor. The archive also prevents repeated rediscovery of failed architectures (Stanley and Lehman, 2015).

12.4 The minimum-burden dominance rule

The Foundry should not optimize for the smallest technical artifact or the largest capability. Among all candidates that pass every hard gate, it selects the candidate with the lowest justified Total Institutional Burden, subject to mission dominance.

Let $\mathcal{P}_m$ be the set of candidates passing the mission gate. The champion is

$$S^* = \arg \min_{S \in \mathcal{P}_m} B_{\text{institution}}(S) \quad \text{subject to} \quad \Delta_{\text{succ}}(S, S_0 \mid m) > \theta_m. \quad (60)$$

This rule can select a more capable or more expensive component when it reduces the total lifecycle burden. It can also select a simpler algorithm, a human-machine process or a decision not to replace the incumbent.

12.5 Search budgets and institutional compute

All search must be budgeted across:

- training and inference compute;
- external services and data;
- simulator and environment creation;
- human expert time;
- verifier development and attacks;
- legal, privacy, security and professional review;
- integration, canary and rollback;
- opportunity cost and duration.

Compute is capital at risk. The Foundry should stop when the expected value of additional search falls below its complete cost or when the opportunity’s capability half-life becomes shorter than the remaining formation path.

12.6 Bounded AGI Jobs as successor-manufacturing contracts

The Successor Foundry does not manufacture an asset through one prompt or one model-training run. It compiles a typed portfolio of bounded AGI Jobs. Each job has an objective, accountable principal, acting identity, authorized inputs and tools, prohibited actions, budget, duration, output schema, proof obligation, verifier, acceptance criteria, rollback route, Evidence Docket return path and Chronicle eligibility.

A canonical reference portfolio contains twenty-one jobs:

Table 9 – Canonical 21-job successor-manufacturing portfolio.

No.	AGI Job	Required output
1	Mission Constitution	Frozen objective, principal, preserve/build/become constraints and prohibited actions.
2	Incumbent Baseline	Accepted quality, cost, risk, time, human burden and authority state.
3	Alternative Set	Best credible frontier, specialist, human, hybrid, software, partner, acquisition and no-action baselines.

No.	AGI Job	Required output
4	Rights and Data Constitution	Provenance, licences, confidentiality, retention, exclusions and reuse rights.
5	Benchmark Constitution	Representative task distribution, protected cases, scorer, critical errors and promotion threshold.
6	Oracle and Teacher Study	Declared external intelligence, teaching signal, hidden-call prohibition and transfer hypothesis.
7	Failure Taxonomy	Error classes, severity, causes, abstention requirements and adversarial cases.
8	Architecture Frontier	Diverse complete candidate systems and non-dominated reserve options.
9	Data and Experience Formation	Rights-cleared examples, simulations, expert demonstrations and quality controls.
10	Specialist Model Formation	Adaptation, training, distillation or algorithmic construction with exact releases.
11	Tool and Retrieval Formation	Mission-specific tools, deterministic checks, retrieval, planners and sandboxes.
12	Verifier Formation	Independent evaluation harness, calibration, replay and protected scoring.
13	Security and Privacy Proof	Threat model, data boundary, credential controls, isolation and incident response.
14	Complete Economics	Formation, inference, human review, maintenance, requalification, dependency and unwind.
15	Human Burden and Agency	Escalation, contestability, supervision, professional authority and role redesign.
16	Resilience and Rollback	Failover, graceful degradation, portability, kill switch and known-good restoration.
17	Protected Fresh Evaluation	Exact frozen challenger on unseen representative work with declared interventions.
18	Independent Validation	Separated review of evidence, limitations, critical errors and complete denominator.
19	Accountable Acceptance	Human or institutional decision on use for the stated purpose.
20	Chronicle and Authority Admission	Scope, version, expiry, revocation, monitoring and Authority Envelope.
21	Capital-to-Capacity and Requalification	Verified-value allocation, successor triggers and next fresh-work test.

The jobs may execute through people, agents, models, vendors, laboratories or mixed teams. An on-chain AGIJobManager can optionally coordinate assignment, escrow, validation or settlement for selected jobs, but it does not decide mission value, truth, acceptance, Chronicle admission or authority inheritance. GoalOS constitutes the work; the settlement rail does not become the constitution.

AGI Job law

SEIZE underwrites the intelligence asset. The Mission Foundry decomposes it. Bounded AGI Jobs manufacture and test it. Proof & Authority verifies it. Chronicle admits it. Capital-to-Capacity compounds it.

13 The Fresh-Successor Gate

13.1 Why inherited evidence is insufficient

A materially changed architecture is a new release. It does not inherit trust from its predecessor merely because the brand, benchmark or workflow name remains the same. Changes in models, tools, prompts, representations, providers, data, routing, agents, policies, infrastructure or authority can create new failure modes.

The Fresh-Successor Gate therefore requires:

- one frozen challenger selected before protected evaluation;
- fresh, production-representative work not used for formation;
- protected cases, seeds, thresholds and scorer internals;
- independent evidence custody;
- exact resource and external-dependency accounting;
- critical-error and lower-confidence gates;
- matched comparison with the incumbent and best alternative;
- governance-integrity and rights review;
- predefined repair, rejection and rollback paths.

13.2 Matched-pair comparison

Where possible, incumbent and challenger should process the same cases. Let d_j be the preregistered utility difference on case j . The mean paired gain is

$$\bar{d} = \frac{1}{n} \sum_{j=1}^n d_j, \quad (61)$$

with a lower confidence bound $\text{LCB}_{1-\alpha}(\bar{d})$. Promotion requires

$$\text{LCB}_{1-\alpha}(\bar{d}) > \theta_m, \quad (62)$$

plus all hard gates. When outcomes are heterogeneous, the analysis should report distributional effects, critical subgroups, worst cases and abstention rather than only the mean.

13.3 The complete successor scorecard

Table 10 – Fresh-Successor Gate scorecard.

Dimension	Required evidence
Mission quality	Frozen scorer, confidence, transfer, calibration, abstention and critical-error results.
Verified value	Accepted economic, scientific, operational or public contribution with explicit attribution.
Time	Time-to-accepted outcome, recovery time and escalation burden.
Complete cost	Formation, operation, compute, proof, review, maintenance, requalification and unwind.
Authority	Exact permitted actions, blocked actions, approvals and unauthorized-action count.

Dimension	Required evidence
Security and privacy	Threat model, isolation, data lineage, secrets, access, incident and exposure tests.
Rights	Model, data, output, tool, verifier, deployment, customer and reuse rights.
Resilience	Provider failure, degraded mode, rollback, continuity and recovery drills.
Human agency	Responsibilities retained by people, override quality and reviewer burden.
Compounding	Reusable evidence, verifier, environment, integration and fresh-transfer potential.
Governance integrity	No material weakening of purpose, accountability, transparency, reversibility or lawful authority.

13.4 Canary and bounded operation

Passing the protected gate makes a successor eligible for a human release decision; it does not prove unrestricted production safety. The release path should proceed through:

1. shadow operation;
2. sealed sandbox;
3. canary with reversible bounded actions;
4. monitored production scope;
5. periodic requalification;
6. expansion only after accepted evidence.

The Authority Envelope must fail closed when the successor leaves the validated distribution, loses required evidence, exceeds budgets, encounters a material incident or reaches expiry.

Fresh-successor law

Reuse is not improvement. Recurrence is not recursive self-improvement. A successor earns authority only by outperforming the incumbent on fresh work under controlled or explicitly normalized conditions, with no material governance regression.

14 Chronicle and the Succession Ledger

14.1 Memory is constitutional admission

Most systems treat memory as a convenience. GoalOS treats memory as the decision that a capability may shape future action. Chronicle therefore admits only capability that is:

- evidenced and attributable;
- rights-cleared and confidentiality-scoped;
- versioned and current;
- bounded to a valid mission and Authority Envelope;
- independently challenged;
- expiring, repairable and revocable.

A result can be operationally useful yet ineligible for Chronicle. Acceptance concerns one outcome. Admission concerns future institutional influence.

14.2 The Succession Ledger

Every admitted successor receives a balance-sheet-like record.

Table 11 – Canonical Succession Ledger fields.

Field	Required content
Identity	Successor ID, mission, owner, customer or beneficiary, version and effective date.
Origin	Succession event, incumbent, thesis, alternatives considered and proof-capital decision.
Exact release	Models, agents, tools, prompts, policies, data, infrastructure, hashes and runtime.
Evidence	Protected evaluation, critical errors, confidence, canary, acceptance and independent review.
Authority	Scope, identities, permissions, budgets, blocked actions, escalation and revocation.
Rights	Model, data, output, tool, verifier, deployment, confidentiality and reuse conditions.
Economics	Cost basis, operating carry, realized contribution, carrying cost and Proof-Adjusted NAV.
Risk	Known failures, exclusions, residual risk, dependency, duration and tail scenarios.
Lifecycle	Expiry, monitoring, requalification, impairment, rollback, repair and retirement.
Succession history	Predecessor, fresh-successor results, descendants and transfer evidence.

14.3 Failure-preserving memory

Chronicle preserves not only winners but also:

- failed architectures;
- invalidated assumptions;
- weak or exploitable evaluators;
- reward-hacking examples;
- rights and contracting failures;
- security incidents;
- provider and partner failures;
- customer refusal reasons;
- conditions under which a capability must not be used;
- repair attempts that produced regressions.

This **Negative Capability Graph** prevents the institution from paying repeatedly for the same failure. It also preserves stepping stones that may become useful under a different mission or capability regime.

14.4 Chronicle precision

Two errors matter:

- **false admission**: unsafe, incorrect, stale or rights-conflicted capability influences future work;
- **false rejection**: valuable capability is discarded, forcing repeated work and slowing adaptation.

Chronicle precision should be measured through later requalification and fresh-transfer outcomes, not by the amount of stored content.

14.5 Impairment, repair and revocation

A successor should be reviewed or impaired when:

- a source, right, law or professional conclusion changes;
- a provider or dependency drifts;
- the valid operating distribution shifts;
- a later mission reveals a material failure mode;
- realized value falls below the accepted thesis;
- the incumbent or another challenger regains superiority;
- Proof Coverage falls below one;
- the capability's expected value becomes negative.

Possible dispositions are CONTINUE, ADMIT WITH LIMITS, REPAIR, REQUALIFY, ROLLBACK, QUARANTINE, EXPIRE, REVOKE and RETIRE.


Chronicle law

Only what survives reality may compound. Every admitted successor remains scoped, attributable, versioned, expiring, challengeable and revocable.

15 Capital-to-Capacity and Recursive Institutional Improvement

15.1 The compounding sequence

The institution-level recursive loop is



The loop is constructive only when each transformation preserves evidence and governance. Otherwise, errors and unauthorized power compound as readily as capability.

15.2 Value states

GoalOS distinguishes value that is:

1. identified;
2. modelled;
3. validated;
4. authorized;
5. implemented;
6. realized;
7. collected or conservatively attributable;
8. reinvestable.

Projected savings and unsigned revenue do not create reinvestment authority. Capital-to-Capacity uses only realized or conservatively attributable value after agreed costs and reserves.

15.3 Allocation classes

Verified value may be allocated to:

- stronger intelligence and search;
- proprietary data and environments;
- evaluators, red teams and proof infrastructure;
- security, resilience and continuity;
- distribution and customer integration;
- human expertise and professional authority;
- compute, energy, laboratories, robotics and physical infrastructure;
- reserve successors and hedge architectures.

The allocation memo must identify the mission bottleneck, expected contribution, evidence required, owner, budget, expiry and stop conditions.

15.4 Recursive improvement at the institutional level

Recursive self-improvement should not be reduced to a system modifying its own code or weights. The stronger and safer formulation is:

An institution recursively improves when a Chronicle-admitted capability measurably increases its ability to detect, constitute, prove or govern a superior successor on fresh work, without reducing proof quality or governance integrity.

Let G_1 be the predecessor generation and G_2 the successor generation using admitted capability. On fresh mission m_2 :

$$\Delta(G_2, G_1 \mid m_2) = a_Q \Delta Q + a_V \Delta V + a_T \Delta T + a_C \Delta C - a_R \Delta R - a_H \Delta H. \quad (63)$$

Promotion requires a preregistered positive threshold and no governance regression. Merely using prior output demonstrates recurrence, not recursive improvement.

15.5 Succession learning curves

The Foundry becomes a platform rather than bespoke consulting when later missions become cheaper or faster because earlier missions created reusable infrastructure.

Track:

$$\text{FormationGain}_n = \frac{B_{\text{institution}}(S_1) - B_{\text{institution}}(S_n)}{B_{\text{institution}}(S_1)}, \quad (64)$$

after normalizing for mission difficulty. Positive FormationGain should be attributable to reusable environments, evaluators, rights templates, security, integration, Chronicle knowledge or distribution - not to weakened proof.

16 Dual-Access Eligibility and the Authority Firewall

16.1 Controlling access rule

The public browser-local institution uses two equal eligibility routes:

$$\text{Access} = \text{CurrentDirectAGIClubOwner} \vee \text{CurrentDirectAGIALPHAHolder}_{\geq 1,000,000} \quad (65)$$

The qualifying AGIALPHA asset is the official 18-decimal token on Ethereum Mainnet at

`0xA61a3B3a130a9c20768EEBF97E21515A6046a1fA.`

The balance route verifies `balanceOf(connectedWallet)` in read-only mode. It does not request token approval, transfer, deposit, payment, staking, locking, burning or custody. Balances are not aggregated across wallets. Indirect beneficial claims, allowances, unrelated wrappers, exchange statements, liquidity positions, other networks and former balances do not qualify.

16.2 Access receipt is not authority

A signed access receipt establishes only that the connected wallet satisfied one route at the verification time and accepted the use boundary. It must record

$$\text{authorityCreated} = \text{NONE}. \quad (66)$$

The access receipt is distinct from:

- the Authority Envelope that governs a bounded mission action;
- accountable acceptance of a result;
- Chronicle admission of reusable capability;
- the authority inheritance decision for a Mission-Sovereign Successor;
- any commercial service, professional opinion, organizational role or financial right.

16.3 Fail-closed session discipline

A production implementation should verify Ethereum Mainnet, exact contract or ENS ownership, current direct eligibility, domain-bound signature, expiry, account and network continuity, inactivity, focus rechecks and protected-export revalidation. It should relock on uncertainty or eligibility loss. The public-static gate provides conditional use, provenance and signed eligibility; it is not confidential server-side DRM.

16.4 Access and AGIALPHA claim boundary

Token-qualified access is utility eligibility. It is not a representation that 1,000,000 AGIALPHA has a fixed product, redemption, cash or investment value. It does not create AGI Club membership, equity, debt, dividends, governance, yield, profit participation or a promise that access will affect token price or liquidity. “Alpha” in this paper means verified mission-specific operating advantage, not AGIALPHA market performance.


Authority firewall

Eligibility may unlock the institution. It must never be confused with identity assurance, competence, professional authority, organizational mandate, Chronicle admission or successor sovereignty.

17 The Five-Plane Technical Architecture

17.1 Why public navigation and protected formation must be separated

The public Navigator is an orientation and qualification surface. It should not contain confidential customer data, protected evaluations, scorer internals, provider secrets, release credentials or consequential execution authority. Sensitive missions require a separately reviewed protected deployment ([Boucher, 2026](#)).



17.2 Identity and access

The protected architecture should support:

- organizational identity and role-based access;
- separation of development, evaluation, release and audit roles;
- least-privilege credentials and short-lived sessions;
- encrypted storage and key management;
- data classification and jurisdictional controls;

- immutable or append-only audit evidence;
- dual control for consequential release;
- incident containment and credential revocation.

17.3 Evidence custody

The independent proof plane should be technically and organizationally unable to leak protected evaluation into candidate formation. It should preserve:

- blind-case commitments and hashes;
- scorer versions and thresholds;
- evaluation runtime and environment;
- exact candidate manifest;
- result receipts and confidence calculations;
- reviewer and acceptance identities;
- release signature and rollback package.

17.4 Graceful degradation

Every admitted successor should define degraded modes:

- full capability;
- restricted capability with additional human review;
- recommendation-only mode;
- read-only mode;
- incumbent fallback;
- emergency stop.

Resilience is part of capability. A successor that cannot fail safely has not passed the complete mission.

18 Risk Architecture

18.1 Basis risk

The evaluator may not represent production. A candidate can dominate the benchmark while failing the workflow, customer, regime or tail cases. Basis risk is reduced through production-representative cases, matched-pair comparisons, critical-error gates, canaries and post-release monitoring.

18.2 Duration risk

The formation programme may take longer than the opportunity remains valuable. Let T_f be expected time to verified successor and H_o the opportunity half-life. Define

$$\text{DurationExposure} = \frac{T_f}{H_o}. \quad (67)$$

A ratio above one suggests that the architecture may be obsolete before release. The correct response may be rental, partnership, a narrower mission or preservation of a reserve option.

18.3 Alpha decay and crowding

Successor Alpha decays when:

- general capability prices fall;
- competitors replicate the architecture;
- the mission or evaluator changes;
- data advantage disappears;
- regulation or rights constrain use;
- maintenance and review burden rise;
- the successor becomes a commodity feature.

Durable advantage therefore resides less in one component and more in the integrated mission, lawful data, verifier, distribution, workflow and feedback loop.

18.4 Dependency and liquidity risk

A successor can become difficult to exit because of proprietary formats, provider APIs, specialized infrastructure, scarce talent, contractual restrictions or deeply embedded workflow. The Succession Constitution should quantify migration cost, fallback availability and the right to export evidence and data.

18.5 Authority-at-Risk

Define **Authority-at-Risk** as the maximum consequential authority exposed to capability whose Proof Coverage is below one or whose evidence is expired:

$$\text{AaR} = \sum_j w_j A_j \mathbf{1}\{\text{ProofCoverage}_j < 1 \vee \text{Expired}_j\}, \quad (68)$$

where A_j is an authority amount or class and w_j reflects consequence. The objective is not necessarily zero autonomy; it is zero unjustified authority.

18.6 Tail loss

Tail loss includes:

- unauthorized financial or legal action;
- security compromise or data exposure;
- harmful physical action;
- discriminatory or unlawful decision;
- systemic correlated error;
- irreversible dependency or institutional capture;
- false admission into Chronicle;
- capital allocation based on fabricated value.

Position limits, dual control, proof gates, kill switches and rollback are constitutional equivalents of risk limits and circuit breakers. The control architecture is compatible with lifecycle risk-management approaches such as the NIST AI Risk Management Framework and its Generative AI Profile, while extending them with succession, admission, proof-capital and capital-allocation objects (National Institute of Standards and Technology, 2023, 2024).

A correct secular thesis can still produce institutional ruin through bad sizing, weak proof, concentrated dependency, path dependence or uncontrolled authority. GoalOS governs the path, not merely the direction.

19 Board Metrics and Operating Cadence

19.1 The board dashboard

Table 12 – Verified Succession board metrics.

Metric	Governing question
Singularity Exposure Ratio	Is adaptation lag below capability half-life?
Time-to-Verified-Successor	How quickly does a material event become a protected succession decision?
Successor Gain	Did the challenger outperform the incumbent and best alternative on fresh work?
Successor Efficiency	How much verified gain was produced per unit of Total Institutional Burden?
Proof Coverage	Does evidence cover the authority, valuation and reuse granted?
Authority-at-Risk	How much consequential authority is exposed to unproven or expired capability?
Successor Option Coverage	Are immediate, challenger, reserve and hedge positions maintained?
Incumbent Decay	What value or control is being lost by delaying justified succession?
Alpha Decay	How quickly is the admitted advantage being commoditized or impaired?
Chronicle Precision	Are useful capabilities admitted and stale or unsafe capabilities rejected?
Capital-to-Capacity Conversion	What durable productive capacity was created from verified value?
Unauthorized Actions	Did any component act outside its Authority Envelope?
Fresh-Transfer Gain	Did admitted capability measurably improve a new mission or successor cycle?

19.2 Operating cadence

19.3 The board brief

Every board brief should answer:

1. What changed?
2. Which incumbent assumption or mission is affected?
3. What value or risk is material?

Table 13 – GoalOS Singularity Office cadence.

Cadence	Institutional function
Continuous	Frontier events, source expiry, authority changes, incidents, blocked actions, provider drift, cost and market changes.
Weekly	Opportunity and Risk Brief; incumbent impairments; active Proof Debt; urgent evidence and access decisions.
Monthly	Successor Book review; capability auction; realized-value ledger; requalification and capital-to-capacity committee.
Quarterly	Institutional reconstitution; challenger frontier; retain, repair, expand, migrate, hedge, revoke or exit verdicts.
Material event	Immediate Capability Shock Protocol after a major capability, law, security, transaction, financing or infrastructure event.
Annual	Complete authority, rights, resilience, data, provider, insurance, professional and productive-capacity review.

- 4. Which successor positions are in the book?
- 5. What evidence would change the recommendation?
- 6. What capital and authority are currently at risk?
- 7. What should happen in the next 72 hours, 30 days, 90 days and 365 days?

The brief should distinguish observed evidence, modelled scenarios, decisions, blocked actions and unresolved Proof Debt.

20

The Successor-Era Corporation: From Digital Presence to Institutional Agency

20.1

A strategic conjecture, not a guaranteed chronology

The website analogy is directionally useful but incomplete. A website gave the corporation a persistent digital presence, distribution surface and transaction interface. A Mission-Sovereign Successor gives the corporation a governed form of institutional agency: the ability to perceive a mission state, assemble evidence, execute bounded work, submit to verification, preserve learning and exercise only the authority it has earned.

The paper does not claim that every corporation will adopt the same architecture or that a universal “ASI corporation” is inevitable. The stronger strategic conjecture is narrower: as common frontier intelligence becomes part of the operating baseline, competitively consequential organizations will increasingly need a governed portfolio of mission-specific successors or will rent equivalent capacity from institutions that possess them.

Table 14 – The website analogy and its limit.

Institutional layer	Website era	Successor era
Presence	Makes the institution visible and reachable.	Makes the institution continuously capable on a constituted mission.
Interface	Publishes information and receives transactions.	Perceives, plans, executes bounded jobs, verifies and escalates.
Memory	Stores content and records.	Preserves evidence, failures, accepted skills, authority and realized outcomes.
Control	Governs pages, identity and commerce.	Governs models, data, tools, proof, permissions, rollback and reproof.
Competitive signal	Absence eventually signals digital irrelevance.	Absence may signal institutional latency, weak learning and dependence on common Beta.
Limit of analogy	A website rarely acts with consequential authority.	A successor may act; therefore proof, revocation and constitutional control are central.

20.2 Why the corporation needs a portfolio, not one monolith

The likely enterprise object is not one all-purpose sovereign agent. It is a portfolio of successors with shared identity, evidence custody, Chronicle, policy, observability and human authority. A corporation may have a Research Successor, Commercialization Successor, Software Successor, Capital-Allocation Successor, Compliance Successor and Customer-Resolution Successor. Each mission has different consequence, evidence, data rights and authority ceilings.

A portfolio avoids two symmetric errors: fragmented shadow agents with no institutional memory, and one monolithic system whose failure or capture compromises the entire corporation. The five-plane architecture described later in the paper allows shared constitutional infrastructure while keeping mission proofs and Authority Envelopes separate.

20.3 The adoption game

Let T_i be institution i 's verified mission-cycle time and L_i its institutional learning rate. A non-adopter may remain viable while its market is stable and human coordination is sufficient. It becomes strategically exposed when competitors' successor-enabled cycle times and learning rates cross a threshold that changes customer expectations, cost structure, opportunity detection or capital velocity.

A stylized relevance condition is

$$\frac{L_i}{T_i} \geq \max_{j \in C_i} \left(\frac{L_j}{T_j} \right) - \epsilon_i, \quad (69)$$

where C_i is the competitive set and ϵ_i is the advantage supplied by brand, regulation, scarcity, relationships, physical assets or other durable complements. The successor becomes strategically necessary when those complements no longer offset the institution's slower adaptation.

Agent-governance work warns that adoption races can pressure firms to deploy systems before reliability and accountability are adequate (Shavit et al., 2023). GoalOS rejects both passive non-adoption and reckless autonomy. The strategic response is a proof-capital race: move faster by purchasing decisive evidence and preserving reversibility, not by laundering uncertainty into authority.

20.4 The build–rent boundary

Not every mission warrants a proprietary successor. The rational boundary depends on mission value, frequency, measurability, data specificity, rights, switching cost, failure consequence and learning potential.

Table 15 – When to rent, configure or constitute.

Mode	Use when	Control requirement
Rent Beta	The task is low-differentiation, reversible, low-consequence and already solved adequately by the market.	Basic vendor, privacy and quality controls.
Configure a workflow	The mission needs institutional context and tools but limited proprietary learning or authority.	Own prompts, data mapping, tests and exportable records.
Constitute a successor	The mission is valuable, repeated, differentiable, measurable, proofable and capable of accumulating proprietary evidence or authority.	Own the mission constitution, jobs, proof, Chronicle, permissions, portability and economics.
Preserve a reserve	The opportunity may become valuable after a capability, price, legal, data or distribution catalyst.	Maintain options and trigger conditions without premature operational authority.

Machine-learning systems can accumulate hidden technical debt through entanglement, undeclared dependencies, data coupling and feedback loops (Sculley et al., 2015). The build decision must therefore value not only differentiation but the institution’s ability to carry the complete lifecycle burden.

20.5 The minimum viable successor institution

A corporation does not begin by automating itself. It begins with one narrow consequential mission and installs the minimum viable constitutional stack:

1. a Mission Charter and explicit reference architecture;
2. a Successor Manifold containing at least three materially distinct designs;
3. a Mission Advantage Gradient and SEIZE underwriting memo;
4. a bounded job portfolio with independent verification;
5. a protected fresh-work proof protocol;
6. an Authority Envelope initially limited to Observe, Advise or Prepare;
7. a Chronicle and Proof-Debt Register;
8. reproof, rollback, expiry and revocation triggers.

The first successor is valuable not only for its direct mission contribution. It establishes the foundry, evidence custody, governance language and operating cadence from which the next successors can be constituted more quickly.

Successor-era corporate law

Thirty years ago, absence from the web increasingly made a corporation invisible. In the successor era, absence of governed machine intelligence may make it institutionally slow. The rational objective is not maximum autonomy; it is the fastest compounding portfolio of mission capabilities whose authority never exceeds fresh proof.

21 Mission Portfolios in Singularity Time

21.1 Where verified succession is most investable

The strongest initial missions usually combine:

- high value or high frequency;
- a bounded state and action space;
- a rapid, trustworthy evaluator;
- repeatable decisions and stable constraints;
- lawful examples, failures and feedback;
- an incumbent with visible burden or decay;
- a realistic path to ownership, control or recurring value;
- an operating environment that can be reset, simulated or replayed.

These conditions make it possible to buy information quickly, prove mission dominance and maintain the successor after release.

21.2 Illustrative mission classes

Table 16 – Illustrative successor opportunity map.

Mission class	Why succession may be investable	Likely successor form	Principal risk
Verified software and algorithmic work	Executable outputs, rapid tests, repeatable environments and high engineering leverage.	Coding agents, search, program synthesis, deterministic tests and human release.	Benchmark leakage and repository-specific overfit.
Financial operations and reconciliation	Recurring exception classes, accounting identities, large volume and accepted historical resolutions.	Evidence extractor, rule engine, workflow agent and independent controls.	Hidden process variation and operational integration.
Cyber find-fix-verify	Scanners, exploits, patches and regression tests create strong feedback.	Sandboxed research agents, repair tools, adversarial verifier and gated release.	Dual-use risk and unsafe autonomy.
Governed function routing and data operations	Finite schemas, known tools, permissions and deterministic validation.	Semantic router, typed calls, policy engine and audit receipts.	Authority laundering and data exposure.

Mission class	Why succession may be investable	Likely successor form	Principal risk
Spreadsheet and financial-model integrity	Formula grammar, neighboring structure and recalculation constrain the solution.	Formula repair, dependency graph and deterministic workbook validation.	Silent business-rule errors.
Logistics, scheduling and allocation	Formal constraints, simulators and measurable cost functions.	Search, optimization, scenario engine and human exception handling.	Model-environment mismatch.
Industrial process and maintenance	Sensors, digital twins, failure costs and physical constraints create value.	Forecasting, control recommendations, simulation and bounded actuation.	Physical tail risk and distribution shift.
Contract, covenant and compliance evidence	Source-linked extraction, rule application, amendment tracking and accountable review.	Ontology, evidence graph, deterministic calculations and professional escalation.	Legal interpretation and rights.
Energy and infrastructure optimization	High capital intensity, formal constraints and substantial operating value.	Simulator, optimizer, scheduler and protected control plane.	Long feedback cycles and safety.
Scientific and engineering design	Searchable design spaces, experiments and objective measurements.	Hypothesis generator, simulator, laboratory loop and independent validation.	Weak surrogate objectives and expensive proof.
Institutional reconstitution	New capability may change roles, products, distribution and economic identity.	New mission graph, human-machine organization, capital stack and successor institution.	Governance regression and stakeholder rejection.

21.3 Selection before formation

A technically impressive candidate in a weak mission is a poor investment. A less dramatic architecture can be a superior successor when it clears the mission floor and creates stronger all-in value, control, resilience or optionality.

A mission should normally be rejected or held when:

- the result cannot be verified independently;
- the workflow is too rare or low-value to amortize formation;
- the distribution shifts faster than requalification;
- rights or access are inadequate;
- there is no accountable owner or route to deployment;
- the best external alternative is already superior;
- the consequence of error exceeds the available proof and control architecture;
- the opportunity will expire before a credible successor can be proven.

21.4 The preserve-build-become portfolio

GoalOS maintains three transformations:

Preserve Identify what no successor may weaken: mission, identity, rights, human agency, public obligations, safety interlocks, strategic data and reversibility.


Build	Constitute what the new environment makes possible: products, Mission Packs, models, agents, evidence systems, markets, data, laboratories, compute, distribution and infrastructure.
Become	Reconstitute roles, operating models, incentives, authority structures and institutional identity when the incumbent can no longer fulfill the mission as well under the new frontier.

The Verified Succession Institution is therefore not only an automation system. It is a method for deciding what the institution should preserve, build and become.

22 Commercial Productization

22.1 One product, several engagement depths

The Verified Succession Institution should remain inside GoalOS rather than becoming a competing brand. The commercial ladder converts orientation into progressively stronger proof and recurring institutional operation.



22.2 The Readiness Mission as deal-flow formation

A strong Readiness Mission should produce:

- a live Institutional Capability Twin;
- an Incumbent and Assumption Ledger;
- twenty succession events or capability spreads;
- ten preliminarily underwritten successor opportunities;
- three SEIZE mandates;
- one frozen Succession Constitution;
- one proof-capital decision;
- a 30/90/365-day Successor Book.

The customer should leave with a capital-allocation answer, not only an AI strategy narrative.

22.3 Sell the decision before the build

The first paid product should answer:

Is the incumbent materially impaired, which successor architecture is most promising, what is the cheapest decisive proof, what is the maximum downside, and should the institution fund, hold, hedge, partner, repair or stop?

This protects both parties. The customer does not fund a large implementation before the mission, rights and proof path are credible. The provider does not guarantee a result that depends on customer data, systems, budgets and authorities not yet frozen.

22.4 The Mission Pack as the delivery object

The canonical Verified Mission Pack should contain:

- Mission Constitution and preserve-build-become constraints;
- Succession Event and incumbent impairment record;
- Successor Frontier and underwriting book;
- mission graph and bounded jobs;
- exact candidate and Release Certificate;
- Authority Envelopes and blocked actions;
- Evidence Docket and proof receipts;
- rights, privacy, security and legal boundaries;
- monitoring, expiry, requalification and rollback;
- Succession Ledger and Capital-to-Capacity memo.

22.5 Renewal proof

Recurring fees should be justified by evidence:

- which frontier changes were detected before the customer would have acted;
- which incumbent assumptions were correctly impaired or preserved;
- which successor opportunities were tested, adopted or rejected;
- which unsafe or uneconomic formations were avoided;
- which Proof Debt was retired;
- which value was realized;
- which successors improved fresh missions;
- which productive capacity was created.

A fixed strategy document expires. A succession institution earns renewal by keeping the architecture current.

23 Implementation Blueprint

23.1 Days 1-15: install the succession doctrine

Add to GoalOS:

- Verified Succession Institution identity;
- Mission-Dominant Specialist ASI and Mission-Sovereign Successor definitions;
- Incumbent and Assumption Ledger;
- Succession Event Record;
- Successor Book;
- updated SEIZE workflow;
- Successor Gain, Successor Efficiency and Proof Coverage metrics;
- bilingual synthetic case;
- JSON import and export.

The integrated release designation is **v3.5.0-SN5-BI2-VS2-MSS1 - Mission-Sovereign Successor Edition**. SN5 remains the constitutional generation until the empirical completion gate is met.

23.2 Days 16-35: build the protected succession workspace

Implement:

- organizational identity and role separation;
- protected mission and economics workspace;
- Succession Constitution compiler;
- proof-capital authorization;
- evaluator and scorer registry;
- Foundry job queue and sandbox controls;
- exact candidate manifests;
- rights and data register;
- Evidence Docket and proof receipts;
- Fresh-Successor Gate;
- release signing, expiry and rollback.

23.3 Days 36-60: run one real succession cycle

Select a mission with:

- a clear incumbent;
- a strong reference capability or alternative;
- an exact evaluator;
- rapid and trustworthy feedback;

- meaningful operating value;
- rights-cleared data;
- a realistic path to bounded deployment.

Produce:

- incumbent baseline and impairment thesis;
- successor frontier;
- one bounded due-diligence mandate;
- one evidence-backed formation route;
- one frozen challenger;
- independent fresh-work result;
- promotion, repair, hold or stop recommendation.

23.4 Days 61-90: constitute the first admitted successor

Deliver either one Mission-Dominant Specialist ASI candidate ready for accountable Mission-Sovereign admission review, or one clean evidence-backed rejection. Record:

- exact release and hashes;
- evidence and acceptance;
- rights and confidentiality;
- Authority Envelope;
- cost basis and Proof-Adjusted NAV;
- realized-value measurement plan;
- expiry and next review;
- rollback and incident route;
- Chronicle admission or rejection.

23.5 The first year

By year end, the institution should operate:

- continuous Succession Radar;
- weekly Opportunity and Risk Brief;
- monthly Successor Book and capital committee;
- quarterly institutional reconstitution;
- material-event Capability Shock Protocol;
- at least one sector-specific Mission Pack;
- independent evaluators and replayable benchmarks;
- a verified-value and Succession Ledger;

- a fresh-successor comparison;
- a Capital-to-Capacity allocation.

24 Worked Example: Acquisition Underwriting Successor

24.1 Mission and reference set

Consider a corporation that repeatedly evaluates acquisition targets. The constituted mission is:

Produce an evidence-bearing investment recommendation that identifies material value drivers, contradictions, downside scenarios, integration risks and decision conditions within a defined diligence window, while reserving all binding authority to the investment committee.

The reference set contains the incumbent human workflow, a strong general-frontier-model workflow, an external specialist diligence provider, a human–AI hybrid and no-action. The candidate is not compared with a weak manual process chosen for convenience.

24.2 Successor Manifold

The initial manifold contains five architectures:

1. **Copilot:** human analysts lead; one frontier model supports document review and drafting.
2. **Specialist cell:** separate finance, legal, market, operations and integration agents report to a human lead.
3. **Adversarial pair:** a thesis system and anti-thesis system compete before an independent verifier.
4. **Sovereign hybrid:** confidential local processing for the data room, external frontier reasoning for non-sensitive tasks, deterministic financial models and a separate evidence service.
5. **Acquired platform:** a third-party diligence product integrated with internal data and approval systems.

The Mission Advantage Gradient may reveal that the largest marginal gains do not come from better prose. They come from reconciling inconsistent financial records, tracing every material assertion to source evidence, adversarially testing management assumptions, preserving prior acquisition outcomes and reducing senior-review burden.

24.3 SEIZE underwriting thesis

The hypothesis is that the sovereign hybrid can create positive Alpha by avoiding one material diligence error, shortening cycle time, increasing evidence coverage and retaining learning across transactions. The initial proof budget is capped and staged:

Table 17 – Illustrative proof-capital tranches.

Tranche	Question	Evidence	Decision
Zero	Can the system reconstruct source-grounded facts without leakage?	Two historical data rooms with hidden answer keys.	Stop if critical-source accuracy is below threshold.
One	Does adversarial structure detect more material contradictions?	Matched-pair comparison against copilot and frontier baseline.	Continue only if incremental detection justifies reviewer burden.
Two	Do complete economics remain positive?	Timed shadow operations, compute, expert review, exceptions and rework.	Freeze architecture or narrow mission if burden exceeds value.
Three	Is bounded live use safe?	One canary transaction with no external communication or binding action.	Admit only Observe/Advise/Prepare authority.

24.4 Bounded AGI Job portfolio

The candidate is constituted through jobs for document inventory, provenance, financial reconciliation, customer concentration, market evidence, legal-risk extraction, thesis construction, anti-thesis, scenario modelling, integration-risk mapping, source verification, uncertainty calibration, critical-error challenge, human escalation, memorandum compilation and Chronicle admission. Each job has explicit inputs, tools, prohibited actions and acceptance tests. No job may contact management, modify the data room, communicate with counterparties, waive a condition, bind the corporation or transfer value. The producer does not see hidden scorer cases. The verifier cannot release the report. The investment committee remains the human principal.

24.5 Comparative proof

Illustrative preregistered metrics are shown below. The numbers are hypotheses, not reported results.

Table 18 – Illustrative fresh-work proof scorecard.

Metric	Incumbent	Admission threshold	Why it matters
Material-risk recall	Measured on blinded cases	$\geq$ incumbent +20%	Avoids false Alpha from faster drafting.
Critical false reassurance	Historical rate	No regression; hard stop above ceiling	Tail-risk gate.
Source-evidence coverage	Current sampled rate	$\geq 95\%$ for material claims	Makes conclusions challengeable.
Median cycle time	Current end-to-end time	$\leq 60\%$ of incumbent	Captures speed value.
Senior-review hours	Current expert burden	$\leq 70\%$ of incumbent	Counts the scarce denominator.
Complete cost per case	Fully loaded cost	Positive net value at expected volume	Prevents compute theatre.
Portability test	Not applicable	Replace one model without losing record or tests	Establishes sovereignty.
Rollback and replay	Current manual reconstruction	Full replay from evidence docket	Supports accountability.

A candidate that writes a better-looking memorandum but misses material contradictions, increases expert rescue or cannot preserve source evidence does not create Alpha.

24.6 Authority and admission

The initial Authority Envelope is:

- **Observe:** ingest approved documents and produce internal diagnostics;
- **Advise:** generate cited findings, uncertainty and recommendations;
- **Prepare:** assemble a draft memorandum and proposed diligence questions;
- **Prohibited:** external communication, transaction execution, contractual commitment, fund movement, condition waiver and self-expansion of permissions.

The candidate may become a Mission-Dominant Specialist ASI if it passes the protected comparison. It becomes a Mission-Sovereign Successor only after the investment committee or other accountable authority signs an Admission Record defining version, scope, expiry, monitoring, fallback and revocation.

24.7 Alpha capture and compounding

The direct Alpha may come from better decisions, faster rejection of weak targets, avoided diligence failures and lower review burden. The compounding Alpha comes from preserving accepted risk patterns, source schemas, scenario templates, verifier cases, integration lessons and realized post-acquisition outcomes.

Let p_b be the reduction in probability of a bad acquisition, L_b its expected loss, V_s speed value, C_h human-burden reduction and C_{all} complete successor cost. A conservative annual value estimate is

$$\alpha_{\text{annual}} = N (p_b L_b + V_s + C_h) - C_{all} - \text{CVaR}_q(L_{\text{succ}}) - D_P, \quad (70)$$

where N is annual mission volume. The estimate is carried only at the confidence supported by realized evidence;

projected avoidance is not automatically treated as realized value.

Worked-example verdict

The acquisition successor is not a better chatbot. It is a governed diligence institution in which producer, critic, verifier, governor and investment committee occupy distinct roles; every material claim carries evidence; authority remains bounded; and the Chronicle converts accepted outcomes into a better next successor.

25 Falsifiability and the SN6 Completion Gate

25.1 What would falsify the thesis

The Verified Succession Institution should be revised or rejected if evidence shows that:

- the successor process does not improve decisions relative to simpler governance;
- proof costs systematically exceed avoided loss or realized value;
- fresh-work gains do not transfer beyond development cases;
- the Successor Book creates delay without meaningful option value;
- Chronicle fails to reduce repeated work or repeated failure;
- later missions are not cheaper, faster or safer after normalization;
- authority separation is routinely bypassed;
- the complete burden makes retention or rental consistently superior;
- customers do not renew based on realized evidence;
- the institution cannot revoke or unwind successors safely.

The framework is strengthened by clean negative findings. An evidence-backed decision to stop is part of its value proposition.


25.2 Minimum empirical proof

A credible commercial and constitutional proof should include:

1. several paid succession mandates;
2. one independently supported Mission-Dominant Specialist ASI and one accountable Mission-Sovereign admission decision;
3. one evidence-backed rejection that prevented material waste or risk;
4. one realized-value result;
5. one repeat or expansion customer;
6. one fresh-successor gain on a new mission;
7. one requalification, impairment or revocation decision;
8. evidence that a later mission benefited from Chronicle and shared infrastructure;
9. zero material unauthorized actions.

25.3 The SN6 gate

The Verified Succession Institution should not be declared a new constitutional generation until the full loop is demonstrated:



The institution enters the next generation only when it can detect a material change, prove relevance on real work, deploy within authority, contain failure, preserve the verified capability and demonstrate a superior successor without relying on narrative confidence.

26 Ten Laws of the Verified Succession Institution

1. **Beta is not Alpha.** Broad access can raise absolute productivity while producing no durable relative advantage.
2. **Spread is not yet value.** The investable gap becomes Alpha only after the complete denominator and every hard gate are counted.
3. **No successor by narrative.** A compelling demonstration, model reputation or benchmark score does not establish mission dominance.
4. **No promotion on inherited evidence.** Every material change to mission, model, data, tool, workflow, verifier, environment or authority receives fresh proof.
5. **Authority never exceeds fresh proof.** Access, execution, acceptance, admission and authority remain distinct states.
6. **No claimant as final judge.** Producer, critic, verifier, governor, admitter and human principal retain separated powers.
7. **No efficiency without the complete denominator.** Capital, compute, integration, human rescue, security, rights, requalification, dependence and unwind all count.
8. **No sovereignty without substitutability.** The institution must retain its mission, evidence, Chronicle and authority system when a component or provider changes.
9. **No compounding without rights and realized evidence.** Learning that cannot be lawfully retained or tied to accepted outcomes does not become institutional capital.
10. **No durable Alpha without renewal.** Chronicle, re-underwriting, bounded improvement, reproof, expiry, rollback and revocation are permanent functions.

27 Conclusion: Constitute Mission-Sovereign Successors

The beginning of Singularity Time does not merely increase the number of tools available to institutions. It shortens the life of accepted architectures and raises the cost of acting from stale assumptions. The institution that treats each capability event as a procurement announcement will repeatedly rent intelligence without learning how to reconstitute itself. The institution that deploys on excitement will accumulate hidden authority, Proof Debt and irreversible dependence.

GoalOS offers a third path: rent Beta where Beta is sufficient; underwrite the Spread where a meaningful mission opportunity exists; manufacture Alpha through differentiated architecture and comparative proof; and capture that Alpha inside a Mission-Sovereign Successor whose authority remains bounded.

The Verified Succession Institution treats every incumbent as provisional and every challenger as untrusted until it proves mission dominance. The Successor Manifold represents what has become reachable. The Mission Advantage Gradient indicates where proof-adjusted superiority may lie. SEIZE converts the opportunity into an underwriting trajectory. The Successor Foundry and bounded AGI Jobs constitute the candidate Specialist ASI. The Fresh-Successor Gate purchases reality. Proof and Authority distinguish empirical dominance from accountable admission. Chronicle preserves the exact, rights-cleared contribution. Capital-to-Capacity finances the next, harder succession.

The deepest recursion is institutional:

A successor is not better because it is newer, larger, faster or more autonomous. It is better only when it produces greater verified mission value, at lower justified burden, while preserving the institution's purpose, rights, reversibility, evidence and human sovereignty.

Final doctrine

The frontier creates candidates. GoalOS constitutes successors. Only a Mission-Dominant Specialist ASI that is separately admitted as a Mission-Sovereign Successor may inherit scoped institutional authority.

Map the manifold. Follow the gradient. Underwrite the Spread. Compile the Bounded AGI Jobs. Constitute the Specialist ASI. Prove Mission Alpha. Admit the Mission-Sovereign Successor. Compound only what survives fresh reality.

A Canonical Machine-Readable Records

The following objects should become canonical schemas inside GoalOS:

1. FrontierEventRecord
2. SuccessionEventRecord
3. IncumbentAssumptionLedger
4. SuccessorBook
5. SuccessorUnderwriting
6. ProofCapitalAuthorization
7. SuccessionConstitution
8. DueDiligencePlan
9. EnvironmentSpec
10. VerifierSpec
11. CandidateSuccessorManifest
12. RunProof
13. FreshSuccessorScorecard
14. PromotionRequest
15. ReleaseCertificate
16. ChronicleSuccessorRecord
17. CapitalAllocationMemo

18. RequalificationRecord

19. ImpairmentRecord

20. RevocationRecord

Every state transition should carry an attributable actor, evidence object, authority receipt, timestamp, version and stop condition.

B Canonical State Machine

INCUMBENT_OPERATING

- > SUCCESSION_EVENT_DETECTED
- > INCUMBENT_IMPAIRMENT_ASSESSED
- > SUCCESSOR_FRONTIER_GENERATED
- > SUCCESSION_OPTION_UNDERWRITTEN
- > PROOF_CAPITAL_AUTHORIZED
- > SUCCESSION_CONSTITUTION_FROZEN
- > DECISIVE_EVIDENCE_PURCHASED
- > CANDIDATE_SUCCESORS_FORMED
- > ONE_CHALLENGER_FROZEN
- > FRESH_SUCCESSOR_PROOF
- > MISSION_DOMINANCE_VERDICT
- > MISSION_DOMINANT_SPECIALIST_ASI
- > ACCOUNTABLE_MISSION_SOVEREIGN_ADMISSION
- > AUTHORITY_ENVELOPE_ISSUED
- > CANARY_OR_BOUNDED_OPERATION
- > ACCEPTED
- > ADMITTED_TO_CHRONICLE
- > VERIFIED_VALUE_ALLOCATED
- > MONITORED / REQUALIFIED / IMPAIRED / REVOKED
- > NEXT_SUCCESSION_EVENT

C Succession Underwriting Card

A board-ready underwriting card should contain:

Field	Required entry
Mission	Exact consequential workflow, beneficiary and deadline.
Incumbent	Current architecture, accepted thesis, cost basis and authority.
Succession event	What changed and which assumption is impaired.
Successor frontier	Retain, repair, rent, form, partner, acquire, hedge, reserve and retire alternatives.
Mission dominance hurdle	Frozen quality, value, time, cost, risk, human burden and governance threshold.
Proof capital	Maximum spend and authority before the next decision.
Next decisive evidence	The smallest artifact capable of changing the position.
Duration	Opportunity half-life, time to proof and time to deployment.
Rights and sovereignty	Data, model, tool, output, verifier, customer and reuse rights.
Tail risk	Critical failure, unauthorized action, security, rights and irreversibility scenarios.

Stop-loss	Condition requiring hold, repair, switch or abandonment.
Position	Incumbent, immediate, challenger, reserve, hedge or retirement.
Owner	Named human or institutional principal.

D Fresh-Successor Gate Checklist

A challenger may be submitted for release review only when:

- the mission and constitution were frozen before protected evaluation;
- one exact challenger was frozen;
- protected cases and scorer internals remained outside formation;
- all external capabilities and human interventions were declared;
- quality, critical-error and confidence gates passed;
- the complete institutional denominator was counted;
- rights and confidentiality were explicit;
- security, privacy and resilience gates passed;
- rollback and expiry were tested;
- mission dominance exceeded the preregistered threshold;
- no material governance regression occurred;
- an independent evaluator signed the result;
- a separate human authority made the release decision.

E Publication Identity, Provenance and Source of Truth

E.1 Authoritative identity

This paper is authored by **Vincent Boucher**, President of **MONTREAL.AI & QUEBEC.AI**. The official publication website is <https://montreal.ai/> and the official author contact is president@montreal.ai. No other similarly named domain, organization or product should be inferred to be an official source for this paper.

E.2 Release provenance

The source of truth for Version 5.0.0 is the signed release set composed of the US Letter PDF, A4 PDF, LaTeX source, bibliography, figures, machine-readable metadata, release manifest and SHA-256 checksums. Generated illustrations communicate the architecture visually; the exact typeset text, equations, tables, claim boundary and authority rules govern whenever visual microtext differs.

E.3 Authorship and claim responsibility

The author is responsible for the framework, terminology, formalizations and institutional design claims presented here. External empirical claims remain attributable to their cited sources. Candidate Specialist ASI and Mission-Sovereign status are never established by authorship or publication alone; they require the fresh comparative evidence, independent review and accountable admission specified by the framework.

F Glossary

Architecture depreciation

Loss of expected value, control or fitness of an incumbent because the capability environment changed.

Authority-at-Risk

Consequential authority exposed to unproven or expired capability.

Capability half-life

Expected time until a material operating assumption should be revisited.

Chronicle

Constitutional memory that admits only scoped, rights-cleared, versioned and revocable capability.

Fresh-Successor Gate

Protected comparison of one frozen challenger against the incumbent and best alternative on new work.

Mission Advantage Gradient

Direction and magnitude of proof-adjusted mission-value increase over the constrained Successor Manifold.

Mission-Dominant Specialist ASI

A complete specialist intelligence institution that exceeds the best credible alternative on fresh work without governance regression.

Mission-Sovereign Successor

A Mission-Dominant Specialist ASI separately admitted by accountable authority under a scoped, versioned, expiring and revocable Authority Envelope.

Proof capital

Bounded capital authorized to purchase evidence before full formation or release.

Proof Coverage

Ratio of independently supported capability and controls to authority, valuation and reuse granted.

Proof Debt

Unresolved assumptions or missing evidence supporting present authority or valuation.

Specialist ASI

Earned mission-bounded designation for a complete system that outperforms general and other credible alternatives across the full mission function.

Successor Alpha

Proof-adjusted value of the best eligible successor above the best credible alternative.

Successor Manifold

Constraint-bearing space of materially reachable successor institutions.

Successor Book

Live portfolio of incumbent, immediate, challenger, reserve, hedge and retirement positions.

Total Institutional Burden

Complete capital, operating, human, proof, dependency, maintenance, requalification and unwind burden.

Verified Succession Institution

Recurring system that underwrites, constitutes, proves, admits and replaces institutional architectures under bounded authority.

References

- George A. Akerlof. The market for “lemons”: Quality uncertainty and the market mechanism. *The Quarterly Journal of Economics*, 84(3):488–500, 1970. doi: 10.2307/1879431.
- W. Ross Ashby. *An Introduction to Cybernetics*. Chapman & Hall, London, 1956.
- Fischer Black and Myron Scholes. The pricing of options and corporate liabilities. *Journal of Political Economy*, 81(3):637–654, 1973.
- Vincent Boucher. *GoalOS Singularity Navigator Ω MasterClass: Tools to Ride the Singularity*. MONTREAL.AI Research, Montreal, Quebec, Canada, 2026. Fully Bilingual Legal-Readiness and Operating MasterClass Edition, Version 3.2.0-SN5-BI1.
- Avinash K. Dixit and Robert S. Pindyck. *Investment under Uncertainty*. Princeton University Press, 1994.
- Tim Genewein, Matija Franklin, Alexander Lerchner, Laurent Orseau, Samuel Albanie, Adam Bales, Cole Wyeth, Stephanie Chan, Iason Gabriel, Joel Z. Leibo, Allan Dafoe, Marcus Hutter, Thore Graepel, and Shane Legg. From agi to asi. *arXiv preprint arXiv:2606.12683*, 2026. doi: 10.48550/arXiv.2606.12683.
- Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. Ai control: Improving safety despite intentional subversion. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 16295–16336, 2024. URL <https://arxiv.org/abs/2312.06942>.
- Sanford J. Grossman and Oliver D. Hart. The costs and benefits of ownership: A theory of vertical and lateral integration. *Journal of Political Economy*, 94(4):691–719, 1986. doi: 10.1086/261404.
- John C. Harsanyi. Games with incomplete information played by bayesian players, parts i–iii. *Management Science*, 1967–1968.
- Oliver Hart and John Moore. Property rights and the nature of the firm. *Journal of Political Economy*, 98(6): 1119–1158, 1990. doi: 10.1086/261729.
- Bengt Holmstrom. Moral hazard and observability. *Bell Journal of Economics*, 10(1):74–91, 1979.
- Michael C. Jensen and William H. Meckling. Theory of the firm: Managerial behavior, agency costs and ownership structure. *Journal of Financial Economics*, 3(4):305–360, 1976. doi: 10.1016/0304-405X(76)90026-X.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R. Narasimhan. Swe-bench: Can language models resolve real-world github issues? In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2310.06770>.
- Jr. John F. Nash. Equilibrium points in n-person games. *Proceedings of the National Academy of Sciences*, 36(1): 48–49, 1950.
- Thomas Kwa, Ben West, Joel Becker, Amy Deng, Katharyn Garcia, Max Hasin, Sami Jawhar, Megan Kinniment, Nate Rush, Sydney Von Arx, et al. Measuring ai ability to complete long software tasks. In *Advances in Neural Information Processing Systems*, 2025. doi: 10.48550/arXiv.2503.14499.
- John M. Lee. *Introduction to Smooth Manifolds*. Springer, 2 edition, 2013.
- James G. March. Exploration and exploitation in organizational learning. *Organization Science*, 2(1):71–87, 1991. doi: 10.1287/orsc.2.1.71.
- Harry Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952.
- Robert C. Merton. Theory of rational option pricing. *Bell Journal of Economics and Management Science*, 4(1): 141–183, 1973.
- METR. Task-completion time horizons of frontier ai models, May 2026. URL <https://metr.org/time-horizons/>. Time Horizon 1.1; last updated 8 May 2026.

- Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. Gaia: A benchmark for general ai assistants. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2311.12983>.
- Meredith Ringel Morris, Jascha Sohl-Dickstein, Noah Fiedel, Tris Warkentin, Allan Dafoe, Aleksandra Faust, Clement Farabet, and Shane Legg. Levels of agi for operationalizing progress on the path to agi. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 36308–36321, 2024. URL <https://arxiv.org/abs/2311.02462>.
- National Institute of Standards and Technology. Artificial intelligence risk management framework (ai rmf 1.0). Technical Report AI 100-1, NIST, 2023.
- National Institute of Standards and Technology. Artificial intelligence risk management framework: Generative artificial intelligence profile. Technical Report NIST AI 600-1, NIST, 2024.
- Martin J. Osborne and Ariel Rubinstein. *A Course in Game Theory*. MIT Press, 1994.
- Mary Phuong, Matthew Aitchison, Elliot Catt, Sarah Cogan, Alexandre Kaskasoli, Victoria Krakovna, David Lindner, Matthew Rahtz, Yannis Assael, Sarah Hodgkinson, et al. Evaluating frontier models for dangerous capabilities. *arXiv preprint arXiv:2403.13793*, 2024. doi: 10.48550/arXiv.2403.13793.
- R. Tyrrell Rockafellar and Stanislav Uryasev. Optimization of conditional value-at-risk. *The Journal of Risk*, 2(3): 21–41, 2000.
- D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips, Dietmar Ebner, Vinay Chaudhary, Michael Young, Jean-François Crespo, and Dan Dennison. Hidden technical debt in machine learning systems. In *Advances in Neural Information Processing Systems*, volume 28, pages 2503–2511, 2015.
- Yonadav Shavit, Sandhini Agarwal, Miles Brundage, Steven Adler, Cullen O’Keefe, Rosie Campbell, Teddy Lee, Pamela Mishkin, Tyna Eloundou, Alan Hickey, et al. Practices for governing agentic ai systems. Technical report, OpenAI, 2023. URL <https://cdn.openai.com/papers/practices-for-governing-agentic-ai-systems.pdf>.
- Toby Shevlane, Sebastian Farquhar, Ben Garfinkel, Mary Phuong, Jess Whittlestone, Jade Leung, Daniel Kokotajlo, Nahema Marchal, Markus Anderljung, Noam Kolt, et al. Model evaluation for extreme risks. *arXiv preprint arXiv:2305.15324*, 2023. doi: 10.48550/arXiv.2305.15324.
- David Silver, Aja Huang, Chris J. Maddison, et al. Mastering the game of go with deep neural networks and tree search. *Nature*, 529:484–489, 2016. doi: 10.1038/nature16961.
- David Silver, Julian Schrittwieser, Karen Simonyan, et al. Mastering the game of go without human knowledge. *Nature*, 550:354–359, 2017. doi: 10.1038/nature24270.
- Kenneth O. Stanley and J. Lehman. *Why Greatness Cannot Be Planned: The Myth of the Objective*. Springer, 2015.
- Steven H. Strogatz. *Nonlinear Dynamics and Chaos*. Westview Press, 2 edition, 2015.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2 edition, 2018.
- David J. Teece, Gary Pisano, and Amy Shuen. Dynamic capabilities and strategic management. *Strategic Management Journal*, 18(7):509–533, 1997. doi: 10.1002/(SICI)1097-0266(199708)18:7<509::AID-SMJ882>3.0.CO;2-Z.
- John von Neumann and Oskar Morgenstern. *Theory of Games and Economic Behavior*. Princeton University Press, 1944.
- Oliver E. Williamson. *The Economic Institutions of Capitalism*. Free Press, 1985.
- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. *arXiv preprint arXiv:2406.12045*, 2024. doi: 10.48550/arXiv.2406.12045.

Publication colophon

Author	Vincent Boucher
Role	President, MONTREAL.AI & QUEBEC.AI
Official website	https://montreal.ai/
Official email	president@montreal.ai
Edition	Version 5.0.0, Authoritative Publication Edition, 9 August 2026
Place	Montreal, Quebec, Canada
Rights	© 2026 Vincent Boucher. All rights reserved.
Source of truth	The versioned LaTeX source, compiled PDF, release manifest and cryptographic checksums distributed together in the official publication suite.



Official publication website: montreal.ai
Contact: president@montreal.ai

Final constitutional rule

Capability may justify candidacy. Fresh proof may establish mission dominance. Only accountable admission may grant authority. Authority remains scoped, versioned, monitored, expiring and revocable.