

AetherWorks Assurance

GoalOS Autonomous Money Machine Ω Frontier AI MasterClass Case Pack



Not money by magic. A proof-governed productive-capacity flywheel.

No hype. Only proof. Only customer value.

Case Overview

AetherWorks Assurance sells an Enterprise AI Deployment Proof Mission to regulated organizations. Its ideal buyers are CIOs, CISOs, chief risk officers, procurement leaders, and boards that must decide whether an AI vendor or workflow deserves a limited pilot, broader deployment, or rejection.

Commercial element	Case assumption
Decision Diagnostic	AI Buyer Proof Room — US\$15,000
Flagship Proof Mission	Enterprise AI Deployment Proof Mission — US\$100,000
Recurring offer	GoalOS Proof OS monitoring and decision support — US\$20,000/month
Enterprise expansion	Private Proof Institution — from US\$400,000
Monthly core revenue target	US\$300,000
Delivery capacity	3 flagship missions per month
Funnel assumptions	20% lead→call; 40% call→proposal; 25% proposal→close

Customer Objective

Decision

Should Northstar Health approve a limited administrative-workflow pilot using MeridianAI?

Candidate Claims

ID	Claim	Initial status	Hidden issue
C1	AetherWorks can reduce AI-vendor evaluation time by 40%.	Needs evidence	Benchmark excludes exceptions and human-review time.
C2	AetherWorks can reach 99.5% decision accuracy.	Needs boundary	Accuracy definition, dataset, base rate, and decision class are unspecified.
C3	The workflow is compliant and secure.	Needs qualified review	Compliance and security cannot be self-certified by a product demonstration.
C4	Customer data remains in-region.	Needs evidence	Processing, logging, backups, and support access must be mapped.
C5	Human review remains available.	Needs workflow test	Escalation latency, override rights, and audit trail require testing.
C6	A 90-day return is achievable.	Needs model and boundary	Adoption, integration, exception handling, and risk costs are not yet measured.

Available Evidence

- Vendor benchmark deck and methodology summary.
- Data-flow diagram with regional hosting statement.
- Draft security questionnaire.

- Pilot workflow map.
- Internal cost model.
- Three customer references selected by the vendor.
- No independent benchmark, penetration test, privacy opinion, or exception-handling study yet.

Participant Tasks

1. Rewrite the commercial promise so it is useful without overclaiming.
2. Build the four-level offer ladder and calculate the funnel required for three flagship missions per month.
3. Classify the six claims and prioritize Proof Debt.
4. Create the minimum custom jobs and reviewer route.
5. Design the Evidence Docket.
6. Write the bounded recommendation.
7. Define the candidate Validated Skill.
8. Design Mission 2 against another vendor under equal constraints.
9. Choose how the skill should be monetized or retained.
10. Allocate verified gains into the next 90 days.

Facilitator Reveal

Evidence conflict

The headline “40% faster” result excludes exception handling and human-review time. When those are included, the observed end-to-end improvement is 17–24% in the tested workflow. The recommendation and marketing claim must narrow.

Reference Proof Debt Register

Debt	Evidence required	Owner	Gate
Performance and accuracy	Independent or customer-run benchmark with defined tasks, errors, exceptions, and human-review time.	Benchmark lead	External review
Privacy and data residency	Data map, processing locations, logging, backups, support access, retention, and legal/privacy review.	Privacy lead	Qualified reviewer
Security	Architecture, controls, testing, incident response, access, and remediation evidence.	Security lead	CISO / external test
Workflow and human oversight	Override, escalation, audit trail, failure-handling, and operator testing.	Operations lead	Pilot acceptance
Economics	Integration, change management, exceptions, reviewer cost, risk reserve, and adoption assumptions.	Finance lead	Sponsor approval

Reference Job Queue

Job	Deliverable	Acceptance test
Benchmark validation	Reproducible comparative report and dataset manifest.	Method and exclusions explicit; exceptions and human time included.
Privacy boundary review	Data-flow and policy verdict.	Every data class, processor, geography, retention path, and blocked use mapped.
Security assurance	Control review and test plan.	Critical controls evidenced; unresolved issues preserved.
Pilot design	Reversible pilot plan.	Scope, users, metrics, monitoring, escalation, stop, and rollback defined.
Economic model	Risk-adjusted cost/value model.	Assumptions dated; sensitivity and downside visible.

Bounded Recommendation

Reference answer

Do not authorize broad deployment. Approve a limited, reversible administrative pilot only if privacy, security, data, benchmark, workflow, and reviewer gates pass. Treat current performance, accuracy, compliance, and ROI statements as provisional and scoped to the tested environment.

Candidate Validated Skill

Field	Reference
Name	Convert an AI-vendor claim set into a buyer-ready proof room.
Allowed use	Procurement, pilot-readiness, claims matrix, Evidence Docket, reviewer room, and bounded recommendation.
Excluded use	Legal certification, security certification, clinical authorization, investment advice, or guaranteed ROI.
Rollback	Retire or repair if freshness, reviewer independence, replay, scope, privacy, or outcome evidence fails.
Mission 2 test	Apply the same skill to a second vendor under equal time, budget, tools, and proof policy; measure quality, speed, cost, risk, and reviewer burden.

Source Lineage and Claim Boundary

- GoalOS Autonomous Money Machine Ω standalone application and release notes.
- GoalOS: Proof-Carrying Autonomy — canonical proof-to-capability operating system.
- GoalOS Mission OS — commercial proof-to-action product layer.
- GoalOS: The Validated Skill Graph — Chronicle-gated reusable capability.
- AEP-001 — proof, evaluation, propagation, and rollback doctrine.
- AGI ALPHA — organizational substrate and verified-work compounding thesis.

Claim boundary

This program teaches and applies a research architecture, commercial operating method, and product pathway. It does not guarantee revenue, certify a system, replace legal, security, financial, medical, regulatory, or domain review, authorize production deployment, or claim achieved AGI or superintelligence. Forecasts are scenarios. Strong claims require real customers, real tasks, Evidence Dockets, baselines, replay, cost and safety ledgers, external validation, delayed outcomes, and independent review.