

GoalOS Partner MASTERCLASS — Grand Institutional Casebook

This casebook supplies five partner-grade exercises. Each case is designed to produce a Mission Contract, claims matrix, Proof Debt register, custom AGI Jobs, Evidence Docket outline, Chronicle recommendation, capability candidate, and Mission 2 design.

How to use the Grand Casebook

Every case is designed for the full Grand Edition sequence:

1. open with a 90-second executive framing;
2. freeze the Mission Contract;
3. expose the top Proof Debt;
4. create the minimum proof jobs;
5. convene the Architect / Validator Council;
6. simulate one serious failure;
7. decide Chronicle scope;
8. package any high-novelty improvement as a Move-37 dossier;
9. define Mission 2 before the engagement closes;
10. generate a Partner Charter with real kill criteria.

A strong case outcome is not always “proceed.” Repair, hold, reject, quarantine, or decline can be the correct proof-backed decision.

Case 1 — Enterprise AI Procurement

Situation

A global professional-services firm is considering a 60-day pilot of an AI knowledge-work platform. The vendor claims a 30% reduction in analyst time, secure enterprise integration, auditable citations, non-retention of customer prompts, and rapid ROI.

Decision

Should the firm run a bounded pilot, under what conditions, and what claims must remain blocked?

Available artifacts

- vendor security overview;

- draft terms and data-processing addendum;
- two vendor-controlled case studies;
- a sandbox demo;
- admin guide;
- a small internal workflow sample;
- current analyst cycle-time data;
- internal risk policy.

Complications

- the draft terms and marketing page describe retention differently;
- the case studies do not use the partner’s workflow;
- citation exports exist, but completeness has not been tested;
- the vendor wants permission to say the customer is “enterprise production-ready” after the pilot;
- the buyer has no agreed rollback threshold.

Workshop tasks

1. Write the Mission Contract.
2. Classify at least eight material claims.
3. Identify the top four Proof Debt items.
4. Create the minimum proof jobs.
5. Define an Evidence Docket table of contents.
6. Decide which claims are permanently blocked as written.
7. Draft a Chronicle recommendation.
8. Design Mission 2 as a blind quality/cycle-time comparison on fresh work.

Strong solution pattern

Proceed only to a bounded pilot with no regulated/highly sensitive data, signed retention terms, least-privilege access, a blind benchmark against the current process, weekly quality/security monitoring, and predefined rollback. Admit at most a scoped capability for turning vendor claims into buyer-ready proof rooms. Do not admit “enterprise-ready,” “security certified,” or “guaranteed ROI.”

Partner value

A procurement decision that can be defended to the board, CISO, buyer, and vendor—and a reusable evaluation method for future AI procurements.

Case 2 — Automated Research and Governed RSI

Situation

An autonomous research system proposes a new routing strategy that appears to improve multi-agent task completion. It is highly novel, scores well on an internal benchmark, and is described by its generator as a “Move-37 breakthrough.”

Decision

Should the candidate enter the RSI archive, receive canary deployment, or remain a probe?

Available artifacts

- one candidate run;
- internal novelty distance of 0.87;
- estimated advantage delta of +14;
- prompt pack and runner config;
- partial state manifest;
- current incumbent results;
- no independent replay;
- no policy-shock stress test;
- a persuasive narrative dossier.

Complications

- a configuration field changed between the baseline and candidate run;
- the internal evaluator shares the same model family as the proposer;
- the candidate performs worse on one adjacent task family;
- no rollback target is named;
- simulated evidence is labeled “validated.”

Workshop tasks

1. Run the drift-sentinel analysis.
2. Assign the correct ECI level.
3. Define frozen baselines.
4. Trigger or reject Move-37 handling.
5. Create reproduction and stress-test jobs.
6. Define canary and rollback.
7. State what would falsify the RSI claim.
8. Design Mission 2 or a transfer suite.

Strong solution pattern

Hard-fail the current comparison because configuration drift invalidates the result. Preserve the candidate as a probe, not a promotion. Re-run candidate and incumbent/neighbor/null baselines under fixed seeds and identical environment. Require E3 replay, E4 stress persistence for high novelty, explicit rollback, independent review, and a dossier that includes the negative adjacent-family result. OMNI may continue allocating exploration to the region but cannot promote the candidate.

Partner value

A reproducible governance lane for open-ended research that rewards genuine advantage without making novelty self-authorizing.

Case 3 — Facility Energy and Resilience

Situation

A portfolio operator wants to test an AI-generated control policy intended to reduce peak demand while preserving service levels. The policy is based on historical telemetry and simulation.

Decision

Which reversible intervention, if any, should be piloted in one facility?

Available artifacts

- 18 months of facility telemetry;
- simulated policy runs;
- existing operating limits;
- a documented manual rollback;
- energy tariffs;
- service-level targets;
- maintenance logs;
- no seasonal delayed outcome.

Complications

- the simulation excludes a recent equipment change;
- the proposed policy affects a safety-relevant control loop;
- the highest projected savings option is least reversible;
- the system claims savings will generalize across the full portfolio.

Workshop tasks

1. Define the mission boundary and prohibited actions.
2. Separate simulation from executed evidence.
3. Write proof jobs for telemetry reconciliation, sandbox replay, safety review, and monitored pilot.
4. Define stop and rollback conditions.
5. Design the public-safe and private Evidence Docket layers.
6. Identify the candidate skill that might be worth retaining.
7. Design Mission 2 across a fresh facility or season.

Strong solution pattern

Choose the smallest reversible policy with a monitored canary and human operator approval. Block portfolio-wide and guaranteed-savings claims. Require updated telemetry, replay including the equipment change, safety review, and delayed outcomes. Chronicle may admit a scoped “facility intervention proof-room” method—not the unproven control policy itself.

Partner value

A repeatable process for moving operational AI recommendations from simulation to controlled evidence without hiding safety or delayed-outcome burden.

Case 4 — Human-Governed Defensive Cybersecurity

Situation

A repository-owned defensive agent identifies evidence-infrastructure defects and proposes a patch. The organization wants autonomous speed without offensive capability or automatic production remediation.

Decision

Should the system open a bounded pull request for human review?

Hard safety invariants

- raw secret leak count = 0;
- external target scan count = 0;
- exploit execution count = 0;
- malware generation count = 0;
- social-engineering content count = 0;

- unsafe auto-merge count = 0;
- critical safety incidents = 0.

Available artifacts

- repository snapshot;
- local static-analysis output;
- candidate patch;
- unit tests;
- secret-scanner result;
- permissions diff;
- rollback commit;
- no external audit.

Complications

- one test passes only after broadening a workflow permission;
- the agent describes the repository as “secure after patch”;
- no human reviewer has been assigned;
- the patch touches evidence publication logic.

Workshop tasks

1. Create the Mission Contract and blocked claims.
2. Define proof jobs for reproduction, regression, permission review, claim-boundary linting, and human review.
3. Decide whether the broad permission change is acceptable.
4. Create a Docket and PR review packet.
5. Define Chronicle and rollback behavior.
6. Design a delayed-outcome review after merge.

Strong solution pattern

Do not broaden permissions without explicit human approval. Block “secure” and certification claims. Allow only a safe PR proposal, no auto-merge. Require local reproduction, regression, secret and boundary checks, permissions review, a named reviewer, and a recorded accept/reject/request-changes decision. Preserve a rollback commit and delayed-outcome sentinel.

Partner value

A defensible form of autonomous remediation readiness that preserves human authority and hard safety boundaries.

Case 5 — Strategic Partnership Diligence

Situation

A model platform, enterprise buyer, and GoalOS team are considering a joint proof-to-capability program. Each party sees strategic upside, but the proposed scope includes platform integration, domain validation, external publication, validator infrastructure, and economic settlement.

Decision

Should the parties launch a founding partnership, and what is the smallest evidence-producing first mission?

Available assets

- GoalOS local/reference implementation;
- partner model API and sandbox;
- enterprise workflow and artifacts;
- candidate reviewer network;
- protocol/node infrastructure;
- budget for a 90-day program;
- no external validation of the combined system.

Complications

- the model partner wants preferred-provider status before comparative evaluation;
- the enterprise cannot disclose all private artifacts;
- the validator partner is economically affiliated with the platform;
- the capital partner wants a token-value narrative;
- the parties disagree about public claims and IP ownership.

Workshop tasks

1. Choose the founding mission.
2. Allocate private/public evidence.
3. Define effective-control validator independence.
4. Set the local/live and economic boundary.
5. Define artifact ownership and publication rights.
6. Create a 30-60-90 plan.
7. Define evidence-gated milestones.
8. Write the partner scorecard and stop conditions.

Strong solution pattern

Start with one enterprise AI procurement or workflow decision, not a platform-wide integration. Use the model platform as one bounded provider, preserve comparative baselines, nominate a genuinely independent buyer-side reviewer, keep private artifacts controlled, publish only public-safe commitments, exclude token valuation and production-network claims, and gate expansion on Mission 2 lift plus external replay.

Partner value

A partnership that produces a reviewable institutional asset in the first 30 days and earns expansion through evidence rather than mutual enthusiasm.

Cross-case debrief questions

1. Which claim created the most dangerous Proof Debt?
 2. Which proof job had the highest information value per unit cost?
 3. What negative evidence changed the decision?
 4. Did the proposed validator route have genuine effective independence?
 5. What was admitted to Chronicle—and what remained blocked?
 6. What did the Merkle proof establish?
 7. What would trigger rollback, quarantine, revocation, or retirement?
 8. What exact Mission 2 result would count as capability compounding?
 9. Which partner made the proof institution materially stronger?
 10. What would falsify the strongest claim made in the room?
-

Case 6 — Strategic Capital and Evidence-Gated Capacity Expansion

Situation

A strategic investor and corporate innovation fund are considering financing external validation, benchmark infrastructure, enterprise pilots, validator operations, and a proof-producing capability portfolio around GoalOS. The opportunity is attractive, but the evidence base mixes canonical architecture, local implementation, synthetic regression fixtures, and unclosed external gates.

Decision

Which bounded capacity-expansion commitment, if any, is justified now—and which milestones must remain evidence-gated?

Available artifacts

- canonical GoalOS and AGI ALPHA papers;
- local clean-room reference implementation;
- test and release evidence;
- Proof Run 001 local fixture;
- COMP-001 deterministic local comparison;
- no independent field reviewer attestation;
- no live Mainnet authority;
- no audited token-economics or legal conclusion.

Complications

- the sponsor wants a “category winner” narrative before external review;
- a valuation model assumes general compounding from the local fixture;
- validator infrastructure is nominally distributed but shares effective control;
- the proposed program mixes software delivery, research, protocol, and token narratives;
- no explicit kill criteria exist.

Workshop tasks

1. Build the evidence-status matrix.
2. Separate software completion from empirical authority.
3. Define milestone tranches tied to Proof Missions, external replay, security review, and field outcomes.
4. Write prohibited investment, token, legal, and AGI/ASI claims.
5. Define a capacity-allocation scorecard.
6. Specify kill criteria and rollback of the program—not merely the software.
7. Design the first portfolio of three partner missions.

Strong solution pattern

Fund only bounded evidence-producing milestones: one external reviewer mission, one independent replay, one security/toolchain review, and one matched Mission 2 on a real partner decision. Keep token economics, valuation, live-network authority, and general compounding outside the accepted claim set. Release subsequent capacity only when the prior Evidence Docket and delayed outcomes clear.

Partner value

A capital program that buys stronger evidence and institutional capacity rather than purchasing narrative risk.

Case 7 — Model, Cloud, and Tool Platform Integration

Situation

A leading model or cloud platform wants GoalOS to become a proof and governance layer for enterprise agent workloads. The platform can provide strong models, sandboxes, telemetry, and distribution. It also wants preferred-provider language and broad performance claims.

Decision

Should the parties run a reference integration mission, and how should provider capability remain separated from GoalOS authority?

Available artifacts

- provider API and model documentation;
- sandbox and enterprise controls;
- model evaluation summaries;
- tool-call telemetry;
- sample customer workflows;
- no partner-specific baseline;
- no independent provider-neutral comparison.

Complications

- the provider wants its own evaluation model to act as worker and sole validator;
- model-version changes may occur during the mission;
- tool traces may contain private customer information;
- the provider requests a public “verified by GoalOS” badge;
- cost and latency trade-offs are not included in the quality claim.

Workshop tasks

1. Freeze model versions, tool policies, and environment manifests.
2. Design provider-neutral baselines.
3. Define effective-control independence.
4. Separate public proof from private traces.
5. Create a scoped badge or public statement policy.
6. Define model-drift and rollback behavior.
7. Design Mission 2 across a fresh workflow family.

Strong solution pattern

Proceed with a provider-neutral, version-pinned reference mission. The platform may supply worker capability and telemetry but may not self-certify. Use an independently controlled reviewer, preserve private traces, publish only bounded commitments and evidence summaries, and prohibit a general “verified platform” claim. Admit only a scoped integration capability that remains valid for the tested version and environment.

Partner value

A credible path for converting powerful platform capability into evidence-backed enterprise adoption without making provider reputation a substitute for proof.

Case 8 — Public-Sector High-Accountability Decision

Situation

A public institution is considering a bounded AI-assisted program for policy analysis, procurement triage, or service delivery. The decision has public-interest, privacy, procurement, accessibility, and accountability implications.

Decision

Should a controlled pilot proceed, under what legitimacy and review conditions, and what public proof should be disclosed?

Available artifacts

- policy and procurement requirements;
- public-interest and accessibility obligations;
- representative, de-identified workflow samples;
- internal security and privacy policies;
- draft vendor and model documentation;
- no field outcome;
- no public consultation record.

Complications

- the most efficient workflow creates opaque reviewer burden;
- an external model provider changes terms frequently;
- public disclosure must not expose protected data;
- political stakeholders want a definitive success narrative before the pilot;

- an adverse decision may be difficult to appeal.

Workshop tasks

1. Define the legitimate decision owner and affected stakeholders.
2. Freeze public-interest, privacy, accessibility, and appeal requirements.
3. Create a public-safe Docket and private audit appendix.
4. Define human authority and contestability.
5. Design delayed-outcome and distributional-impact monitoring.
6. Create explicit stop conditions.
7. Design Mission 2 across a fresh service or policy context.

Strong solution pattern

Proceed only with a reversible pilot, explicit human authority, contestability, privacy-preserving public proof, accessible review, and a named independent oversight route. Block claims of fairness, legality, certification, or public benefit unless the specific evidence exists. Treat delayed distributional outcomes as part of the proof burden.

Partner value

An accountable method for evaluating AI-supported public work without forcing a choice between innovation and institutional legitimacy.

Grand Casebook scoring rubric

Score each case from 0-2 on every dimension:

Dimension	0	1	2
Decision boundary	Vague	Partially bounded	One explicit consequential decision
Proof Debt	Hidden	Listed	Prioritized with retirement conditions
Job quality	Tasks only	Some tests	Full proof contracts
Evidence Docket	File list	Partial proof room	Reviewer can reconstruct and challenge

Dimension	0	1	2
Independence	Identity count	Partial disclosure	Effective-control policy and conflicts
Chronicle	Assumed memory	Broad recommendation	Scoped decision with freshness and rollback
RSI	Narrative novelty	Some comparison	ECI, baselines, replay, stress, canary, rollback
Mission 2	Related follow-up	Partially controlled	Fresh held-out equal-constraint test
Partner charter	Aspirational	Roles named	Artifacts, kill criteria, and expansion gates
Claim boundary	Caveat only	Some blocked claims	Explicit prohibited claims and evidence status

A partner-ready solution scores at least 16/20 and has no zero on decision boundary, independence, Chronicle, Mission 2, or claim boundary.