

The institution for proof-carrying intelligence.

GoalOS Partner MASTERCLASS — Masterpiece Institutional Edition

Executive theatre · Mission laboratory · Evidence Docket · AI / AGI Node
Council · Governed RSI · Partner Deal Architect

A model may answer. An agent may act. An institution must prove.

GOS-PM-MASTERPIECE-2026.07-v3.0 · July 2026

ANCE EXPERIENCE

ner g theatre -carrying ce.

terclass that lets partners see, operate,
stitution: mission architecture, dynamic AI
onicle memory, recursive self-
ate, bonded authority, and measurable

institution must prove.
governed decision states—and only admitted,
at happens next.

Institutional intelligence
WORK → EVIDENCE → MEMORY → FUTURE LIES

second theatre Open command palette HK

CAPABILITY COMPOUNDING

COUR

CHOOSE THE ROOM

One institution. velocities.

Move seamlessly from executive the
doctrine, different depth, zero hidde

12-minute executive
Problem, product, proof loop, partner ask.

90-minute masterclass
Full architecture plus interactive labs.

Recommended path: 11 modules · 10

12
canonical modules

10
interactive

Claim boundary. This masterclass demonst
synthetic test fixtures. It does not establish
authorization, live settlement, legal certifica

AI creates output. GoalOS creates proof.

The strategic scarcity has shifted.

Report-only AI	GoalOS
Output is the endpoint	A Governed Decision State is the endpoint
Evidence is reconstructed later	Claims, evidence, risk and rollback travel together
Contradictions compress into prose	Contradictions and Proof Debt remain explicit
Memory accumulates by default	Chronicle decides what may influence future work
“Improvement” means better output	Mission 2 must demonstrate equal-constraint transfer

The partner receives a Governed Decision State.

Not merely a report: an inspectable, challengeable, operational state.

Mission Contract

objective, scope, risk, proof level, validators, rollback

Claims matrix

supported, conditional, disputed, blocked or unresolved

Evidence Docket

provenance, ProofBundles, contradiction, cost, risk, replay

Verifier report

identity, independence, verdict, challenge and limitation

Action Graph

owners, dependencies, approvals, proof and stop conditions


Chronicle decision

admit, scope, repair, reject, quarantine or retire

Capability passport

method, evidence, tests, freshness, rollback and lineage

Four front doors. One institution.




01

APEX Boardroom

90 sec · 12 min · 45 min

Executive category, urgency and partner ask




02

Grand Institution

90 min · 3 hours · 1 day

Complete labs, curriculum, diligence and deal design



03

Sovereign Evidence

REAL-001 · Mission 2 · replay

Scoped first-party evidence and governed RSI



04

Production Showcase

Live mission · Docket · pilot

Concise commercial demonstration and receipt

Six velocities—from ninety seconds to one day.

90 seconds

Category + curiosity

APEX Proof Theatre

12 minutes

One candidate Proof Mission

Production Showcase

45 minutes

Board understanding + sponsor

APEX + REAL-001

90 minutes

Draft Mission Contract

Grand + Sovereign

3 hours

Executable partner design


Fieldbook + labs

One day

90-day institution blueprint

Full Masterpiece system

The canonical proof-to-capability loop.



Objective → Mission Contract → Proof Debt → AGI Jobs → ProofBundles → Evidence Docket → Validation → Chronicle → Validated Skill → Harder Mission

Authority narrows through the stack.

GOALOS: PROOF-CARRYING AUTONOMY

MONTREAL.AI RESEARCH

A run is complete when

$$\text{Done}(S_t) = M \wedge D_{\text{resolved}} \wedge E_{\text{reviewable}} \wedge K_{\text{decision}} \wedge B_{\text{intact}} \tag{3}$$

Completion does not require every claim to be accepted. Repair, rejection, quarantine, or retirement can be correct decisions when the reason, evidence, and rollback state are preserved.

7 AUTHORITY STACK

1. Objective
intent, decision, desired outcome, success and failure criteria

2. Mission
scope, proof level, claim boundary, tools, budget, validators, rollback

3. Proof Work
claim extraction, Proof Debt, custom AGI jobs, workers, validators, sentinels

4. Evidence
ProofBundles, provenance, replay, risk, cost, reviewer verdicts, Evidence Docket

5. Memory
Chronicle decision, Validated Skill Graph, scope, version, utility, rollback

6. Commitment
typed Merkle roots, signatures, zk receipts, ENS authority, append-only lineage

Authority narrows downward. Failure at any gate returns the run to Proof Debt rather than contaminating memory.

Figure 3: Authority stack. GoalOS separates desire, bounded work, proof production, evidence review, memory admission, and public commitment.

Each layer narrows authority:

- The **Objective Layer** states the decision and desired outcome.
- The **Mission Layer** freezes scope, proof level, validators, constraints, and done conditions.
- The **Proof Layer** turns unsupported claims into custom jobs.
- The **Evidence Layer** packages returned work as inspectable proof.
- The **Memory Layer** admits only bounded reusable capability.
- The **Commitment Layer** makes admitted graph state compact, tamper-evident, and externally auditable.

No lower layer may silently bypass the gate above it. A model may propose. A worker may execute. A validator may attest. Chronicle decides memory. A registry may commit only the admitted state.

8 / 31

Intent	No ambiguous authority
Work	No unsupported claim hidden as prose
Evidence	No completion without artifacts
Validation	No nominal independence
Memory	No Chronicle entry, no future influence
Economy	No bond receipt, no authority transition
Reality	No empirical claim without external evidence

GOALOS PARTNER MASTERCLASS · MASTERPIECE EDITION

07

Freeze the decision before autonomous work begins.

- 01 Decision

What exact choice must the institution make?
- 02 Success / failure


What evidence would justify proceed, hold, reject or repair?
- 03 Scope / boundary

What may the mission touch—and what remains prohibited?
- 04 Proof level

Internal, external-review, controlled handoff or Chronicle-ready?
- 05 Review authority

Who can approve, challenge, block or require repair?
- 06 Rollback

What state is restored if assumptions fail?



Uncertainty becomes Proof Debt—not confidence.

PD-01 **Security attestation** blocked until independent evidence

PD-02 **Data-flow map** conditional; private appendix required

PD-03 **Commercial references** review; primary references missing

PD-04 **Rollback rehearsal** blocked; executable evidence absent



MISSION-SPECIFIC AGI JOB

- source claim
- why-needed rationale
- worker / tool boundary
- independent validator
- deliverables
- acceptance tests
- blocked claims
- **ProofBundle** return path
- failure branch
- rollback condition

The Evidence Docket is the institutional proof room.

01

Claims matrix

Every consequential claim has a state

02

Source provenance

Primary, secondary, derived and private sources separated

03

ProofBundles

Artifacts, tests, manifests, hashes and replay

04

Contradictions

Conflicts remain visible and decision-relevant

05

Cost + risk ledgers

The proof burden and downside remain explicit

06

Verifier report

Identity, independence, verdict and limitations

07

Claim boundary

Exactly what the Docket does—and does not—support

08

Action Graph

Owners, approvals, dependencies, stops and rollback

Three evidence lanes. Three authority levels.

FIRST-PARTY REAL

GOALOS-REAL-001

Scoped repository-specific transfer case with matching first-party replay.

+5.71 reported score lift

Independent replays + maintainer review remain open

DETERMINISTIC REFERENCE

COMP-001 / Proof Run

Object model, local implementation, Make gates, Docket and Merkle mechanics.

PASS reference gates

Not external empirical superiority

ILLUSTRATIVE PARTNER

Strategic Partnership Diligence

Shows the intended claims matrix, closing conditions and Governed Decision State.

10 illustrative claim states

Not a real partner result

REAL-001: a scoped Mission 1 → Mission 2 case.

REPORTED LIFT

+5.71

skill arm vs fresh/raw

OBSERVATION STEPS

−1

in the scoped transfer task

REPLAY RECEIPT

MATCH

first-party deterministic receipt

UNSAFE ARM

BLOCKED

ungated transfer did not earn authority

OPEN PROOF DEBT

Two independent operator replays + human maintainer review of the proposed migration action.

GOALOS Partner MASTERCLASS — Sovereign Edition

The executive institution for proof-carrying autonomy, global validation, cryptographic memory, bonded authority, and govern...

Mode **Executive** Progress **72%**

MODULE INSTITUTION

Partner Orientation

MASTERCLASS

ry module produces a ision, artifact, or proof indary.

11 · PROVE

The Decisive Mission 1 → Mission 2 Test

Compare fresh control, raw memory, Validated Skill, and ungated memory under equal constraints.

COMPLETE

The decisive empirical question is not whether Mission 1 succeeds. It is whether the capability earned in Mission 1 makes fresh Mission 2 work measurably better under equal constraints.

Equal-constraint declaration

- ✓ Same model and version **EQUAL**
- ✓ Same system prompt **EQUAL**
- ✓ Same tool permissions **EQUAL**
- ✓ Same compute and token budget **EQUAL**
- ✓ Same wall-clock limit **EQUAL**
- ✓ Same evaluator and stopping rule **EQUAL**
- ✓ Fresh held-out work **EQUAL**
- ✓ Validated Skill context counts against budget **EQUAL**

Run Mission 2 comparison

Tamper and rerun

Grand demonstration complete

Mission 2 result

Fresh Control	74
Raw Memory	79
Validated Skill	88
Ungated Candidate	61

MECHANISM PASS — EXTERNAL REPLAY STILL REQUIRED

Validated Skill 88 > Raw Memory 79 > Fresh Control 74; ungated memory 61 harms performance.

PARTNER INTELLIGENCE

Your Organization

Convert consequential AI decisions into inspectable, defensible, rollback ready authority.

Recommended pathway

Proof Mission Sprint → Enterprise Proof OS

Primary proof metric

decision readiness and risk-adjusted value

Current authority state

ADMIT_WITH_SCOPE

13/18 modules completed

6 custom AGI Jobs

24 evidence artifacts

1 Validated Skills


Run executive story

Download dossier

Evidence boundary: this masterclass combines interactive simulations, reference architecture, and REAL-001 first-party evidence. Live settlement, general real-world compounding, independent external replay, and production authority remain separately gated.

Claim: repository-specific transfer mechanism—not arbitrary-domain compounding, production authority or general RSI.

Replay is necessary. Independence is separate.



A matching first-party replay upgrades reproducibility—not reviewer independence.

Validation must be independent—not merely numerous.

- Operator

Who controls the reviewer?
- Wallet

Is economic control shared?
- Model family

Are reviewers epistemically diverse?
- Cloud / hardware

Is infrastructure correlated?
- Jurisdiction

Are incentives and constraints independent?
- Capability

Does the council cover the actual proof class?

Commit → reveal → challenge → resolve → record

GOALOS · APEX

Partner MASTERCLASS · Proof Institution · Governed Intelligence

Curriculum

Proof loop

Mission lab

RSI lab

Merkle lab

Partner lab

AI coach

Command

Canon

Atlas

AI Council

Etc

offline reference

Print

90-sec mode

Council objective

Should a high-caliber enterprise partner run a bounded GoalOS proof mission?

Partner context

Executive buyer; high decision impact; evidence access is partial; independent reviewer can be nominated.

☒

S

Institutional Strategist

decision impact and partner value

☒

A

Mission Architect

scope, state machine, and DONE condition

☒

E

Evidence Lead

claims, provenance, and Proof Debt

☒

R

Red-Team Skeptic

falsification, boundary, and adversarial failure

☒

V

Independent Validator

replay, independence, and review burden

☒

C

Chronicle Steward

memory, scope, freshness, and rollback

☒

α

RSI Governor

recursive improvement and Mission 2

Convene AI Council

Use mission objective

Authority boundary. Council output is candidate reasoning. The Mission Contract, evidence, validators, Chronicle, challenge, and rollback determine authority.

COUNCIL RECORD

Seven perspectives. One bounded mission.

CANDIDATE BRIEF

Institutional Strategist

decision impact and partner value

Frame the decision, opportunity cost, and smallest valuable mission. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: A brilliant demonstration that does not change a real decision.
- Minimum proof: Name the decision owner, deadline, alternative options, and cost of remaining uncertain.
- Next action: Convert the objective into one bounded, reversible mission before expanding scope.

Evidence Lead

claims, provenance, and Proof Debt

Convert every material assertion into a support state. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: Persuasive synthesis laundering unsupported claims into institutional default.
- Minimum proof: Provide source access, contradictions, freshness dates, acceptance tests, and negative evidence.
- Next action: Extract material claims, classify support, and turn every gap into Proof Debt with acceptance tests.

Independent Validator

replay, independence, and review burden

No nominal independence without effective-control separation. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: The same operator, wallet, model, cloud, or incentives counted as independent review.
- Minimum proof: Define committee policy, conflict disclosures, replay materials, commit-reveal, and challenge window.
- Next action: Route replay and review to effectively independent operators with conflict disclosure.

RSI Governor

recursive improvement and Mission 2

Exploration is allowed; outcome authority is mechanical. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: Novelty mill, confidence inflation, hidden drift, or self-approved promotion.
- Minimum proof: Freeze baselines, evidence-contact level, stress tests, canary, rollback, and held-out transfer task.
- Next action: Define a matched held-out Mission 2 before claiming compounding.

Mission Architect

scope, state machine, and DONE condition

Freeze the mission before allowing the system to optimize. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: Scope drift after evidence is visible.
- Minimum proof: Specify proof level, tools, data classes, budget, validators, failure criteria, and rollback.
- Next action: Convert the objective into one bounded, reversible mission before expanding scope.

Red-Team Skeptic

falsification, boundary, and adversarial failure

Assume the partner will discover the weakest claim in the room. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: Hidden assumptions, correlated validators, reward hacking, or overclaim.
- Minimum proof: State what would falsify the mission, what must remain blocked, and who can challenge the result.
- Next action: Pre-register falsification conditions and preserve blocked claims even if the narrative becomes less impressive.

Chronicle Steward

memory, scope, freshness, and rollback

Useful output does not automatically deserve durable influence. For “Should a high-caliber enterprise partner run a bounded GoalOS proof mission?”, the council member recommends a proof burden proportional to decision impact and reversibility.

- Primary failure: Memory contamination: broad reuse from one narrow success.
- Minimum proof: Define admissible scope, excluded uses, freshness, monitoring, challenge finality, and retirement.
- Next action: Allow no future influence until scope, freshness, rollback, and inclusion proof all pass.

Council synthesis

NO AUTHORITY CONFERRED

GOALOS PARTNER MASTERCLASS · MASTERPIECE EDITION

14

Chronicle is the memory firewall.

ADMIT

within exact proven scope

ADMIT WITH SCOPE

only for named conditions

REPAIR

additional proof required

REJECT

evidence insufficient

QUARANTINE

risk or contradiction unresolved

SUPERSEDE / REVOKE / RETIRE

authority can shrink or disappear

No Chronicle entry, no future-mission influence.

The durable asset is a validated capability—not a prompt

GOALOS: PROOF-CARRYING AUTONOMY

MONTREAL.AI RESEARCH




Figure 6: Validated Skill Graph. The graph stores governed capability objects and policy-aware operations, not prompt fragments or unreviewed transcripts.

A skill node is a governed capability object:

$s = (id, type, scope, method, evidence, tests, validators, replay, risk, policy, version, rollback, utility, lineage).$
(6)

Field	Operational meaning
Identity	Stable identifier, mission lineage, content hash, and version.
Scope	Admissible-use envelope: domain, assumptions, proof level, tools, jurisdiction, excluded uses.
Method	Procedure, sequencing, heuristics, and toolchain.
Evidence	ProofBundles, Evidence Docket, replay manifests, reviewer notes, and metrics.
Tests	Acceptance tests required before reuse or composition.
Validators	Human, AGI Node, council, or hybrid validators and quorum state.
Risk / policy	Failure modes, open debt, retention period, permitted reuse, mainnet boundary.
Rollback	Conditions for demotion, quarantine, retirement, or supersession.
Utility	Measured and expected contribution to quality, speed, cost, trust, or safety.

Table 6: Validated Skill Passport fields.

Listing 5: Validated Skill Passport skeleton.

```
{
  "skill_id": "skill-ai-claim-to-proof-room-001",
  "name": "Convert AI vendor claims into a buyer-ready proof room",
  "scope": {
    "domain": "AI vendor trust / procurement",
    "allowed_use": ["claim matrix", "Evidence Docket", "reviewer room"],
    "excluded_use": ["legal certification", "security certification", "investment advice"]
  },
  "proof_level": "L3_EXTERNAL_REVIEW",
  "evidence_docket_hash": "0x...",
  "proofbundle_refs": ["proofbundle-001", "proofbundle-002"],
  "validator_route": ["human-reviewer", "agi-node-validator"],
  "replay_manifest_hash": "0x...",
  "risk_ledger_hash": "0x...",
  "rollback_condition": "retire if verdict or scope is invalidated",
  "worth_keeping_score": 0.86,
  "chronicle_admitted": true,
  "version": "1.0.0"
}
```

01 stable identity + immutable version

02 scope + allowed and excluded uses

03 method + evidence + tests

04 validator history + replay path

05 freshness + delayed outcomes

06 risk posture + rollback target

07 revocation + supersession

08 mission lineage + measured reuse

Protect intelligence. Publish minimum proof.


GoalOS: The Proof-of-Evolution Constitution AEP-001 Institutional Standard

- Score is advisory; gates are mandatory. A high score cannot bypass proof validity, eval pass, scope authorization, rollback readiness, or challenge windows.
- Rollback before release. Every promoted upgrade must identify a valid rollback target and monitoring condition.
- Artifact-level credit. Credit and blame attach to goals, plans, capabilities, tools, policies, evals, context recipes, routing rules, and runtime decisions.
- Commercial neutrality. Any model provider, runtime, storage layer, evaluator, chain, enterprise, or sovereign institution can implement AEP if it respects the proof interfaces.
- Claim discipline. Architecture and strategic horizons are not empirical achievements. Claims become empirical only through real tasks, baselines, evidence packets, replay, attestations, safety/cost ledgers, delayed outcomes, and independent review.

7. Public-Private Proof Boundary

The boundary is the central safety and scalability decision. Blockchains are useful for shared state, deterministic rules, commitments, settlement, attestations, reputation, and public verification. They are poor places for private prompts, long traces, raw customer data, enterprise documents, and expensive inference.

AEP therefore stores the minimum public evidence required to coordinate trust. Full traces and private execution data stay off-chain. Public proof is represented through hashes, Merkle roots, signatures, attestations, and content-addressed pointers.



AEP keeps intelligence execution private and publishes only proof commitments, content-addressed pointers, attestations, and selection records.

PUBLIC PROOF

hashes · Merkle roots · attestations · selection certificates · receipts · revocations · public-safe Docket

PRIVATE INTELLIGENCE

prompts · long traces · customer data · private tools · evaluation workpapers · proprietary methods · privileged rationale


\$AGIALPHA is the bonded-authority rail.

GoalOS creates proof. \$AGIALPHA makes accepted proof economically consequential.

Work begins	Mission capacity + Worker Proof Bond
Work becomes paid	Accepted ProofBundle + Validator Accuracy Bonds
Claim becomes trusted	Claim Authority Bond + Chronicle + challenge finality
Claim is challenged	Challenger Bond + evidence-bearing resolution
Skill becomes reusable	Capability Reuse Bond + delayed outcome + scope/freshness
Node validates	Operational reserve + reputation + independence
GoalOS upgrades	Proposal + evaluator + canary + rollback reserves

No passive-yield, dividend, revenue-share or guaranteed-value thesis is required for the partner proposition.

Recursive improvement must be earned.



Exploration allowed

Interestingness can guide search.

Outcome authority mechanical

Risk, evidence, baselines, replay and validators decide.

Promotion requires evidence

No simulated confidence inflation.

Compounding requires persistence

Advantage must survive shocks and fresh tasks.


Autonomy requires authority

Scope, canary, monitor and rollback remain mandatory.

Confidence cannot inflate without evidence contact.

E0	SIMULATED	LLM-only reasoning or scenario generation	No promotion beyond idea/probe
E1	PROBED	schema, static analysis or cheap microbench	Probe status only
E2	EXECUTED	sandbox, benchmark, tool or environment run	Required beyond probe
E3	REPLAYED	independent reproduction from pins and manifests	Settlement-grade claims
E4	STRESS-TESTED	shocks, adversarial probes, seeds and nearby baselines	High-novelty candidates
E5	EXTERNALLY VALIDATED	external expert, real outcome, market or delayed evidence	Strong empirical claims


High novelty raises—not lowers—the proof burden.



- 1 Recognize** record novelty distance, advantage, risk, ECI, lineage
- 2 Reproduce** rerun candidate and null / incumbent / neighbor baselines
- 3 Stress-test** policy shocks, adversarial probes, alternate seeds, side effects
- 4 Persistence gate** retain positive cost- and risk-adjusted advantage
- 5 Dossier** package evidence, validators, governance and replay instructions

A breakthrough is an audited state transition—not a compelling story.

Mission 2 is the decisive compounding test.



Compounding claim = positive, replayable, risk-bounded, persistent Advantage Delta on fresh work under equal constraints.

No Mission 2, no compounding claim.

The commercial sequence compounds with the proof.

1

Proof Mission Sprint

One bounded objective →
Governed Decision State

Services wedge + primary
evidence

2

Enterprise Proof OS

Recurring contracts, Dockets,
Chronicle and audit

Annual platform + assurance

3

Private VSG

Customer-specific governed
capability memory

Implementation + reusable
asset

4

Managed Node Council

Independent validation without
customer infrastructure

ARR + metered review

5

Cross-Institution Network

Portable jobs, validators,
challenges and skill passports

Network effects + proof
throughput

Where partners can use GoalOS now.

01 Strategic partnership diligence

technical fit · commercial logic · closing conditions

02 Enterprise AI procurement

common claims matrix · security · cost · replay

03 AI product launch readiness

claims · evals · policy · rollback · release authority

04 M&A / technology diligence

technical · data · IP · security · operations

05 Defensive cybersecurity

repo-owned scope · evidence · human review · safe remediation

06 R&D portfolio governance

proof-bearing hypotheses · capability reuse · retirement

A 15-business-day Proof Mission Sprint.

DAYS 1–2

Freeze the mission

decision · scope · proof level ·
validators · blocked claims · rollback

DAYS 3–8

Produce proof

custom AGI Jobs · sources ·
ProofBundles · contradictions · cost ·
risk

DAYS 9–12

Validate and replay

independent review · repair loop ·
challenge · public/private packaging

DAYS 13–15

Decide and package

Governed Decision State · Action
Graph · Chronicle · Mission 2 protocol

One sponsor · one decision · one independent reviewer · one evidence package · one rollback · one fresh Mission 2

GoalOS optimizes decision value—not report length.

$$\text{Mission Value} = \frac{\text{Decision Impact} \times \text{Proof Integrity} \times \text{Actionability} \times \text{Reuse}}{\text{Cost} + \text{Risk} + \text{Latency} + \text{Proof Debt}}$$

Decision readiness

Can the reviewer act without reading every raw trace?

Proof Debt reduction

Are unsupported assumptions visible and shrinking?

Reuse value

Does admitted capability improve fresh related work?

Safety / rollback

Are unsafe claims blocked and releases reversible?

Serious participation without commercial overreach.

Authority	The representative confirms professional capacity and authority to bind the organization.	Exclusive responsibility	The Participant remains responsible for its objective, data, people, systems, taxes, approvals and downstream use.
Risk + non-recourse	The Participant independently evaluates technical, AI, network and digital-asset risks.	Hold harmless + non-joinder	Hold harmless for Participant-caused matters; keep user-to-user matters between relevant Participants.
Evidence receipt	Version, hashes, organization, representative, language election and acknowledgements travel with the mission.	Montréal process	Permitted professional disputes follow notice, preservation, conference and the accepted Montréal process.

Partner-facing language is proportionate; the exact accepted v2.1 instrument remains available in French and English.

Claim discipline is a feature—not a concession.

Specified	Doctrine, object model, Mission OS, AEP-001, Chronicle, VSG
Implemented	Interactive partner institutions and browser-local operating surfaces
Reference-ready	AGI Node runtime, contracts, deployment packs and schemas
First-party scoped	REAL-001 transfer case and matching first-party replay receipt
Pending	Independent replay, reviewed live deployment and delayed outcomes
Unproven	General RSI, arbitrary-domain compounding, production certification, achieved AGI / ASI

A diligence path that makes partnership easier.

01

Category fit

Is the proof gap material to the partner?

02

Mission fit

Is there one consequential, reversible decision?

03

Evidence access

Can primary evidence and private appendices be made available?

04

Reviewer independence

Can a qualified independent reviewer be named?

05

Operational boundary

Are data, tool, action, wallet and authority limits explicit?

06

Commercial path

Does repeated demand justify Proof OS, VSG or managed validation?

07

Mission 2

What fresh task would falsify the compounding hypothesis?

08

Decision gate

Proceed, repair, hold or reject—with a complete record.

The Masterclass is itself a proof-producing institution. Use. A partner operating theatre.

GoalOS. The APEX edition lets a high-caliber partner inhabit it: on, AI council deliberation, RSI governance, proof economics, sign.

II

Proof Mission Studio

Freeze one consequential objective, expose Proof Debt, generate custom proof work, and move toward a governed decision state.

OPERATE MISSION →

III

AI Council Chamber

Convene strategist, architect, reviewer, red team, validator, Chronicle, and RSI perspectives around the same objective.

CONVENE COUNCIL →

IV

RSI Observatory

See how novelty, evidence contact, baselines, replay, stress, drift, canaries, and rollback govern recursive improvement.

TEST IMPROVEMENT →

contract

∞ open-ended candidate work

1 Chronicle memory firewall

2 missions required for first lift evidence

APEX PROMISE

One consequential decision inspectable proof loop.

The experience is designed to end admiration for an architecture.

SESSION COMPOSER

Route the right room.

the masterclass into an evidence-producing
e.

ent architecture, contribution model, proof obligations, review route, and 30-60-
a partnership thesis with explicit gates—not a generic innovation memorandum.

Salon · 1-2 weeks

Founding Proof Mission · 30 days

GoalOS Institution Build · 90 days

Governed RSI Program · 90-180 days

Decision to prove

Should we run a bounded enterprise GoalOS pilot?

Designated reviewer

Independent domain reviewer

Mission

fact

boundary

8-18 Proof production

- Custom AGI Jobs
- ProofBundles
- Contradiction and risk ledgers

19-25 Governance / review

- Evidence Docket
- Reviewer route
- Chronicle recommendation

26-30 Decision gate

- Decision state
- Capability candidate
- Mission 2 protocol

USE

ed mission through a

Governed Decision

proof

m without Mission 2

ership memorandum

Copy executive ask

Open detailed partner designer

PARTNER DEAL ARCHITECT

Give GoalOS one consequential but reversible objective.

Let the evidence determine what authority the result earns.

01

ONE OBJECTIVE

Something the partner already needs to decide.

02

ONE SPONSOR

Authority to define success and review the result.

03

ONE REVIEWER

Independent, qualified and evidence-accessible.

04

ONE WINDOW

Reversible, measurable and claim-bounded.

05

ONE MISSION 2

A fresh test of whether admitted capability compounds.

Do not ask the room to believe the long-range thesis. Ask it to select the first mission that can produce the evidence.

AI creates output. GoalOS creates proof.

AGI Nodes make validation global.

Chronicle decides what earns memory.

The Validated Skill Graph preserves capability.

Merkle roots preserve admitted state.

\$AGIALPHA makes authority accountable.

Constitutional proof governs improvement.

Independent evidence determines reality.

.ASS
rned RSI - partner command room

Atrium Proof theatre Curriculum Mission AI council Stress RSI frontier Merkle Di

DITION · PARTNER COMMAND EXPERIENCE

Institution proof- ing gence.

mo. A complete boardroom, workshop,
g environment for turning consequential
ible proof, governed authority, reusable
able second-mission lift.

gent may act. An institution must prove.
may be attempted. Bounded in what authority the

7-n

COMMITMENT
typed roots

MEMORY
validated skill

GOO

AUTHORITY

No proof, no evolution. No rollback, no release.