



Successor Ω

The Illustrated Guided Institutional Journey

From Rented AI Beta and Underwritten Spread to Specialist ASI, Accountable Authority and Recursive Renewal

With Proof Run 001: Orion Building Operations and four additional corporate demonstrations

The governing promise

**Beta is rented. The Spread is underwritten. Alpha is manufactured. Specialist ASI is proven.
Authority never exceeds fresh proof.**

MISSION FIRST

PROOF ABOVE CLAIM

AUTHORITY ACCOUNTABLE

MEMORY INSTITUTIONAL

Vincent Boucher

President, MONTREAL.AI & QUEBEC.AI

Montreal, Quebec · 16 August 2026

A paper designed to be used, not merely admired

One mission. One journey. One governed succession decision.

What this edition does differently

This paper is written from scratch as a guided institutional journey. It does not append, summarize or hide an earlier monograph. A new reader is taken from the first business question—*where is our real mission advantage?*—through candidate formation, fresh comparative proof, accountable authority and successor-to-successor renewal.

- why access to frontier AI is usually AI Beta, not proprietary Alpha;
- how a mission gap becomes an investable Mission Advantage Spread;
- why the unit of competition is a complete institution, not an isolated model;
- why training, proof and authority must remain separate;
- what makes a proven Specialist ASI different from an admitted Mission-Sovereign Successor Ω .

By the end, you can produce

- a Mission Constitution;
- a Successor Book and candidate manifold;
- a SEIZE underwriting decision;
- a Bounded AGI Job Graph;
- a fresh-proof protocol and six-gate verdict;
- an Authority Envelope, Chronicle receipt and renewal trigger.

The recurring demonstration

The paper follows **Orion Building Operations**: a synthetic six-building portfolio that must reduce energy waste and comfort failures without granting any authority over fire, access-control, elevator or other life-safety systems. The same institutional sequence is then demonstrated for Invoice Integrity, Acquisition Underwriting, Verified Software Release and Supply Chain Shock Navigation.

All people, organizations, buildings, invoices, code repositories, suppliers, cases, metrics, economic values and operating results in the worked demonstrations are deterministic synthetic reference evidence. They illustrate the institutional method; they are not production proof, a forecast, a valuation, a legal opinion or a grant of real authority.

THE MODEL MAY BE REPLACED. THE MISSION, PROOF, AUTHORITY AND MEMORY MUST REMAIN GOVERNED.

How to read the paper

Three reading speeds. One institutional logic.



90 seconds

Read the gold stage headlines and the final governing law on each page.



20 minutes

Follow the journey map, the Orion demo boxes and the five use-case cards.



Deep study

Read the formal definitions, evidence tables, equations and implementation appendix.

The page grammar

Every stage page answers five questions in the same order:

1. **What is it?** A plain-language institutional definition.
2. **What do I do?** The smallest useful action for a new user.
3. **What exists afterward?** A named corporate artifact.
4. **What failure must I avoid?** The most common category error.
5. **What happens to Orion?** The continuing worked demonstration.

The separation that governs the entire journey

The Gym may teach. The Proof Plane may certify. The human principal may admit. No one role may silently perform all three.

The whole journey on one page

Two synchronized tracks: economic value and institutional succession.

Economic track

AI BETA $\longrightarrow$ MISSION ADVANTAGE SPREAD $\longrightarrow$ PROVEN MISSION ALPHA



The practical output

A new user should leave the journey with one mission-specific decision: **retain the incumbent, repair the candidate, reserve or hedge an option, reject the programme, stop the search, or admit one exact proven release under one bounded Authority Envelope.**

Beta, Spread and Alpha without the jargon

Beta is the baseline. Spread is the opportunity. Alpha is what survives proof.




Figure 1: The economic meaning of the journey. The institutional programme exists to convert an apparent gap into conservative proven mission advantage.

The widely rentable frontier baseline. It may be excellent, necessary and strategically valuable. It is not automatically differentiated because competitors can often rent it too.

Mission Advantage Spread

The underwritten opportunity gap between what the incumbent or rentable frontier can do and what a mission-owned complete institution might do better after cost, risk and burden are counted.

Mission Alpha

The residual mission advantage that remains after fresh comparative proof and the complete denominator. Alpha is not a slogan, model score or optimistic scenario.

The first law for new users


Do not ask, “Which model is smartest?” Ask, “Which complete admissible institution creates the strongest conservative mission advantage?”

The object you are building

Successor Ω is not a model with a grand title.



A model or policy is one replaceable component inside a complete mission-bearing institution. The institution also contains the Mission Constitution, data and rights boundaries, Bounded AGI Jobs, workflow, tools, humans, independent Verifier, evidence, Chronicle, rollback, monitoring and accountable authority.



THE MODEL REASONS. THE INSTITUTION SUCCEEDS.

Contents

Contents

1	Meet Orion: one mission, six buildings, one bounded question	20
2	Mission Constitution and candidate architecture market	20
3	The Formation Gym and Verifier Foundry	22

4	A single scenario, walked end to end	23
5	Fresh proof: one frozen challenger versus two credible baselines	24
6	Mission Alpha and the accountable authority decision	25

7	Canary, Chronicle and successor-to-successor renewal	27			
8	Choose the first mission for proof, not for prestige	40			
9	The 90-day activation path	41			
10	Who does what: the six-role institutional game	42			
11	The Board dashboard	43			
12	Gym-to-reality basis risk	44			
13	The formal core, translated into business language	46			
	13.1 Spread: the investable opportunity	46			
				13.2 Mission Alpha: the residual after the complete denominator	46
				13.3 The proof-adjusted selection rule	46
				13.4 Evidence-bounded authority	46
			14	The fifteen-stage quick reference	47
			15	The ten governing laws	47
			16	Glossary for the first conversation	48
			17	What to do on Monday morning	49
			18	Research and implementation lineage	49

PART I



The verified succession journey

A mission becomes investable only when value, proof and authority remain visibly separate.

00 · YOU ARE HERE

Mission Constitution

Choose the mission before choosing the intelligence.



A Mission Constitution states the beneficiary, repeated decision or operating function, scope, evidence standard, hard constraints, prohibited actions, human principal, authority ceiling and renewal conditions. It turns an attractive AI idea into a falsifiable institutional mandate.

What the reader does

Write the mission in one sentence. Then name the worst unacceptable outcome, the evidence that would reveal it, and the highest authority that could be reversible on day one.

What exists at the end

A signed Mission Constitution, rights register, critical-failure definition and initial authority ceiling.

Starting with a model demonstration and reverse-engineering the mission afterward. A technically impressive demo cannot repair an undefined purpose, missing rights boundary or ambiguous principal.


Orion Building Operations · the continuing demo

Orion's mission is to reduce energy waste and tenant-comfort failures across six buildings. It may recommend and execute tightly bounded non-safety adjustments; it may never touch fire, access-control, elevator or other life-safety systems.

01 · YOU ARE HERE

AI Beta and the incumbent

Establish the exact baseline before claiming superiority.



AI Beta is the rentable frontier baseline: models, agents, tools and services that may be excellent but are broadly accessible. The incumbent is the complete current institution, including humans, software, workarounds, latency, cost and hidden rescue.

What the reader does

Measure the incumbent and best credible external alternative on the same mission denominator: quality, economics, reliability, evidence, human burden, dependency and critical failures.

What exists at the end

A Beta Baseline Card and an incumbent evidence dossier.

Benchmarking only the model. If the incumbent includes an expert operator, a manual control or an undocumented rescue path, that burden belongs in the baseline rather than disappearing from the comparison.

Orion Building Operations · the continuing demo


Orion baselines its current building-management system plus operators, its contractor network and a commercially available specialist service. General-purpose AI is allowed to summarize logs, but it receives no operating authority.

02 · YOU ARE HERE

Mission Advantage Spread

Underwrite the opportunity gap before funding the solution.

Beta is the baseline. Spread is the opportunity. Alpha is what survives proof.



The Mission Advantage Spread is the investable difference between the best available baseline and a plausible mission-owned institution. It is not yet Alpha; it is a hypothesis about value after the complete denominator is counted.

What the reader does

Write a one-page Spread thesis: what changes, who benefits, how value appears, what costs and tail risks could erase it, and which evidence would change the decision.

What exists at the end

A Spread thesis with explicit upside, complete cost, risk reserve and basis-risk reserve.

Calling every capability improvement “Alpha.” Faster answers can be valuable while remaining fully commoditized, symmetric and externally governed.

Orion Building Operations · the continuing demo

Orion’s Spread is not “AI can optimize HVAC.” It is the specific gap created by earlier fault detection, fewer false dispatches, lower energy cost, better evidence and less operator burden across the portfolio.

03 · YOU ARE HERE

Navigator Ω

Detect successor events before the incumbent becomes obviously obsolete.



GoalOS Singularity Navigator Ω is the frontier radar. It watches models, law, security, infrastructure, providers, costs, capability half-life and mission conditions for events that change the succession decision.

What the reader does

Define a small watchlist of signals and thresholds. Each signal must map to an institutional decision: retain, review, hedge, re-underwrite, freeze, revoke or requalify.

What exists at the end

A Navigator watchlist, event ledger and succession-review trigger.

Treating trend monitoring as strategy. A feed of impressive announcements is useful only when it changes a mission, proof-capital or authority decision.

Orion Building Operations · the continuing demo

Orion triggers review when a provider changes telemetry semantics, equipment is replaced, the tariff regime changes, a new control vulnerability appears or comfort complaints exceed the baseline band.

04 · YOU ARE HERE

Successor Book

Preserve options before choosing the winner.



The Successor Book is the portfolio of institutional positions: incumbent, immediate improvement, challenger, reserve, hedge and retirement. It prevents one seductive architecture from becoming the entire strategy.

What the reader does

Name at least one credible option in every position. Record the condition under which each option should be promoted, repaired, held in reserve, used as a hedge or retired.

What exists at the end

A versioned Successor Book with promotion and retirement conditions.

Confusing a shortlist of models with a portfolio of complete successors. Every entry must include workflow, data, humans, Verifier, rollback and authority boundary.

Orion Building Operations · the continuing demo

Orion keeps the incumbent BMS as the rollback baseline, adds deterministic fault rules as an immediate improvement, forms three challengers, retains a specialist cell as reserve and preserves a manual contractor path as hedge.

05 · YOU ARE HERE

Successor Manifold

Complete complete institutions, not algorithm names.



The Successor Manifold is the space of genuinely different complete ways to perform the mission. Candidates may be deterministic, human-machine, agentic, world-model based, specialist collectives or sovereign hybrids.

What the reader does

Build a candidate card for each architecture: policy, model, data, tools, workflow, humans, Verifier interface, rollback and authority boundary. Eliminate candidates that cannot satisfy hard constraints.

What exists at the end

A candidate architecture portfolio with admissibility status.

Assuming the most advanced algorithm deserves to win. A simpler architecture wins whenever added intelligence, autonomy or complexity fails to produce enough conservative proven mission value.


Orion Building Operations · the continuing demo

Orion compares the current operator-led system, deterministic rules, a rented frontier agent, a human-machine specialist cell, a world-model planner and a sovereign hybrid with independent verification.

06 · YOU ARE HERE

Mission Advantage Gradient

Find the direction that deserves the next unit of proof capital.



The Mission Advantage Gradient estimates which architectural change is most likely to increase mission value after complete cost, risk, human burden, dependency, Proof Debt and Gym basis risk.

What the reader does

Compare one change at a time: better evidence, lower latency, stronger verifier, improved transfer, world-model planning or tighter authority. Ask which experiment can most reduce decision uncertainty.

What exists at the end

A gradient memorandum ranking the next proof-capital moves.

Optimizing the easiest benchmark. The highest training-score gradient may point toward evaluator exploitation rather than real mission advantage.

Orion Building Operations · the continuing demo

For Orion, the strongest early gradient is not a larger model. It is better latent fault-state estimation plus an anti-gaming outcome Verifier and protected transfer cases from an unseen building.

07 · YOU ARE HERE

SEIZE Underwriting

Buy the smallest decisive evidence, not the largest technical programme.



SEIZE treats proof as capital allocation under uncertainty. It selects the smallest reversible experiment that can change whether the institution should build, rent, partner, reserve, hedge, reject or stop.

What the reader does

Define the experiment, complete cost, proof-capital ceiling, stop-loss, decision threshold and next action for every possible result. Stop when expected decision improvement falls below complete experiment cost.

What exists at the end

A signed SEIZE memorandum and proof-capital budget.

Approving a broad transformation programme before the evidence question is defined. More activity is not more proof.

Orion Building Operations · the continuing demo

Orion authorizes a 90-day programme capped at CAD 625,000: CAD 450,000 for formation and Verifier construction, CAD 175,000 for independent fresh proof, with explicit stop conditions.

08 · YOU ARE HERE

Bounded AGI Jobs

Turn the mission into sealed work-and-proof contracts.



Bounded AGI Jobs define objective, inputs, tools, limits, budget, deliverable, required evidence, Verifier, acceptance threshold and rollback. They compose capability and proof; they never silently compose production authority.

What the reader does

Create a dependency graph for mission, data and rights, environment, candidates, tools, Verifier, economics, security, fresh evaluation, admission, Chronicle and requalification jobs.

What exists at the end

A Bounded AGI Job Graph with accountable owners and evidence receipts.

Writing vague tasks such as “build the agent” or “validate performance.” If the evidence, acceptance rule and rollback are absent, the job is not constitutionally complete.

Orion Building Operations · the continuing demo

Orion issues jobs for sensor-rights review, digital-twin calibration, candidate construction, verifier red-teaming, protected scenario custody, economic measurement, canary control and Chronicle admission.

09 · YOU ARE HERE

Sovereign Mission Gym

Manufacture experience where failure is visible and reversible.



The Sovereign Mission Gym is an institution-owned environment with states, observations, actions, transitions, outcomes, constraints and evaluation. Formation is visible, iterative and adaptive; protected proof is not.

What the reader does

Specify what the candidate can observe, which actions are bounded, how transitions occur, which rewards may be optimized and which hard constraints can never be traded away.

What exists at the end

A reproducible Gym specification, scenario portfolio, reward vector and hard-constraint registry.

Using the same cases to teach and certify. The Gym that teaches must not certify; protected information cannot leak back into formation.

Orion Building Operations · the continuing demo

Orion's Gym includes historical replay, a calibrated digital twin, weather and occupancy regimes, sensor drift, equipment changes, adversarial fault combinations and a shadow-mode interface.

10 · YOU ARE HERE

Successor Foundry

Let algorithms compete inside a sovereign institution.



The Successor Foundry routes methods by mission structure and evolves complete institutions. MuZero-style planning, Dreamer-style world models, offline learning, PPO/SAC, deterministic rules, specialist collectives and human-machine workflows are competitors—not constitutional authorities.

What the reader does

Route methods using observability, sequential depth, historical data, simulator fidelity, action-space structure and safety consequence. Preserve separate Pareto elites for value, safety, cost, transfer, sovereignty and reversibility.

What exists at the end

A Pareto Successor Archive and documented algorithm-routing decisions.

Collecting fashionable algorithms without proving which complete architecture deserves capital. The Foundry must remain sovereign over its methods.

Orion Building Operations · the continuing demo

Orion forms a deterministic baseline, a specialist cell, a model-free agent, a world-model planner and a sovereign hybrid. The archive preserves the safest, cheapest and most transferable candidates even when they do not lead on average reward.

11 · YOU ARE HERE

Exact-Release Freeze

Create an immutable identity before proof begins.



The exact-release freeze locks models, weights, prompts, data, tools, workflow, Verifier configuration, budgets, permissions, rollback and a hashed manifest. After the freeze, the challenger may execute but may not learn.

What the reader does

Hash every operative component. Verify that the proof environment has no trainer code, no threshold-changing path and separate credentials. Record the one exact release that will be judged.

What exists at the end

A signed exact-release manifest and freeze receipt.


Evaluating a moving target. A candidate that changes during proof can absorb protected information and invalidate the verdict.

Orion Building Operations · the continuing demo

Orion freezes the sovereign hybrid release, the digital-twin interface, tool registry, evidence schema and A0 proof permissions. Proof-time updates are exactly zero.

Fresh Proof

Let protected reality decide whether training mattered.



Fresh Proof compares one frozen challenger against the incumbent and best credible alternative on protected representative, tail, adversarial, subgroup and transfer work. Independent custody, preregistered thresholds and robust statistics govern the verdict.

What the reader does

Run paired cases. Report mean gain, lower confidence bound, lower-tail performance, critical errors, evidence coverage, complete cost and proof-time update count. Preserve every legitimate verdict.

What exists at the end

A signed comparative Evidence Docket: promote, retain, repair, reserve, hedge, reject or stop.

Selecting the winner by narrative, formation score or vendor benchmark. The highest training score does not necessarily win.

Orion Building Operations · the continuing demo

On 1,800 protected synthetic scenarios, Orion's frozen hybrid records 93.1% target accuracy, 99.5% material-risk recall, 96.1% evidence coverage, zero unauthorized actions and a positive lower-confidence-bound gain over both comparators.

Specialist ASI

Recognize a proof state, not a marketing label.



Specialist ASI is the earned mission-specific comparative state of one exact frozen candidate that passes six gates: capability, economics, reliability, sovereignty, governance and transfer.

What the reader does

Require every hard gate to pass. Publish threshold, observed result, evidence source, uncertainty and failure mode. No catastrophic failure may be averaged away by many small successes.

What exists at the end

A six-gate Specialist ASI verdict tied to one release, mission and evidence set.

Calling the output of training “Specialist ASI.” Training produces candidates; protected comparative proof may establish the proof state.

Orion Building Operations · the continuing demo

Orion’s sovereign hybrid passes all six synthetic gates. The verdict remains mission-specific: it does not imply general superiority, production readiness or permission to act.

Accountable Admission

Grant authority separately, minimally and reversibly.



Accountable admission converts a proven candidate into a governed operating institution only when a human principal issues a scoped, versioned, monitored, expiring and revocable Authority Envelope.

What the reader does

Choose the lowest authority level that can realize value: A0 observe, A1 recommend, A2 sealed sandbox, A3 reversible bounded action or A4 consequential action. State explicit prohibitions and automatic rollback triggers.

What exists at the end

A signed admission memorandum, Authority Envelope and revocation path.

Letting capability silently become authority. A candidate may win the mission and still deserve only recommendation or sandbox status.

Orion Building Operations · the continuing demo

Orion is admitted at A3 for non-safety setpoint adjustments within narrow bands and work orders below CAD 5,000. Fire, access, elevator, tenant surveillance and capital-project authority remain prohibited.

15 · YOU ARE HERE

Successor Ω , Chronicle and renewal

Operate as a replaceable institution—then prove the descendant again.



Mission-Sovereign Successor Ω is the complete mission-bearing institution: mission, proven Specialist ASI, AGI Jobs, data, tools, workflow, Verifier, evidence, Chronicle, rollback, Authority Envelope and human principal. Chronicle preserves validated memory, exact versions, failures, interventions and realized outcomes.

What the reader does

Monitor evidence support, basis risk, drift and authority use. When impairment occurs, reduce or revoke authority, re-underwrite, create new jobs, form a descendant, freeze it and run fresh proof again.

What exists at the end

A governed operating successor, Chronicle receipt, impairment policy and dated requalification plan.

Treating admission as permanent trust. A successor is a contingent institution whose authority is borrowed and whose evidence expires.

Orion Building Operations · the continuing demo

Orion enters a 30-day reversible canary. Chronicle records 11.8% synthetic energy improvement, 64% fewer comfort-violation minutes and zero unauthorized actions. A later provider change triggers requalification; the descendant inherits knowledge, not proof or authority.

What the full journey prevents

Most institutional failures are category errors.

Category error	What goes wrong	Constitutional correction
Model = institution	A benchmark winner is deployed without workflow, evidence, humans, rollback or authority design.	Compete complete candidate institutions.
Training = proof	The candidate is certified on the same environment that taught it.	Freeze one release and use an independent one-way Proof Plane.
Capability = authority	Strong performance silently expands permissions.	Require separate accountable admission and a minimal Authority Envelope.
Average score = safety	Catastrophic failures are offset by many small successes.	Use non-compensable hard gates and lower-tail reporting.
Memory = truth	Unverified output propagates into future work.	Admit only evidenced, rights-cleared, scoped Chronicle records.
Upgrade = renewal	A descendant inherits the predecessor's permissions.	Return proof and authority to zero; re-underwrite and re-prove.

The operating doctrine

Mission first. Proof above claim. Authority is accountable. Memory is institutional. Alpha is measurable. Succession is permanent.



Proof Run 001 · Orion Building Operations

A complete synthetic demonstration: from mission constitution to fresh proof, A3 admission, Chronicle and renewal.


1 Meet Orion: one mission, six buildings, one bounded question

The mission

Keep occupied spaces within comfort and energy bounds by detecting operational faults early, explaining the evidence, recommending corrective action and—only when admitted—executing tightly bounded non-safety adjustments or low-value work orders.

- It is repeated, measurable and economically consequential.
- Historical telemetry and work-order outcomes exist.
- A digital twin can manufacture controlled experience.
- Many actions can be reversible and bounded.
- Transfer to unseen buildings can be tested.
- Life-safety systems can remain explicitly out of scope.

The candidate may not control fire systems, access control, elevators, emergency ventilation, tenant surveillance or unrestricted capital projects. It may not widen its permissions, change the Verifier or modify protected proof thresholds.




What Proof Run 001 includes

1,800 protected synthetic scenarios across six buildings and six operating regimes; three complete comparator institutions; a digital-twin Formation Gym; a separate Fresh-Proof Plane; six Specialist ASI gates; a 240-case reversible canary; an A3 Authority Envelope; Chronicle and impairment triggers. The canonical seed is 20260816.

2 Mission Constitution and candidate architecture market

Constitutional field	Orion specification
Beneficiary	Occupants, building owners and the accountable operations team.
Repeated mission	Diagnose operating state, explain evidence, recommend or execute bounded correction, verify realized outcome.
Critical failure	A material equipment fault or simultaneous heating/cooling condition is falsely reassured or a prohibited system is actuated.
Evidence standard	Attributable sensor history, relevant work orders, weather/occupancy context, tool receipts and independent outcome verification.
Authority ceiling	A3 reversible bounded action; no A4 authority.
Expiry	Thirty days, or earlier on model/provider, equipment, policy, law, data or Verifier change.




Incumbent BMS + operators preserves the real baseline. **Best external specialist cell** combines commercial tools with human review. **Frozen sovereign hybrid** combines a calibrated world-model planner, deterministic safety rules, an evidence graph, independent Verifier interface, human escalation and exact rollback. Simpler architectures remain available in the Successor Book as reserve and hedge positions.

Selection rule

The Foundry does not choose the most advanced algorithm. It chooses the simplest admissible complete institution with the strongest conservative proven mission advantage.


3 The Formation Gym and Verifier Foundry



- historical replay and digital-twin episodes;
- normal, seasonal, drift, vendor-change and rare-tail regimes;
- bounded actions and refusal;
- demonstrations, offline learning, self-play, world-model planning and failure replay;
- visible rewards and non-compensable constraints.

Verifier Foundry

- outcome and comfort verification;
- evidence and process verification;
- economics and human-burden accounting;
- critical-error and unauthorized-action gates;
- protected-case leakage, concealed rescue and reward-hacking attacks.




The protected Proof Plane contains no trainer code, uses separate credentials, holds unseen cases under independent custody and returns only signed evidence. The frozen challenger may cross. Protected information may not return.

4 A single scenario, walked end to end

Protected scenario ORI-1247

Orion C is occupied at 82%. Outdoor temperature is 31.4 °C. Chilled-water delta-T has fallen for three hours, valve command and flow disagree, comfort complaints are rising and a recent vendor firmware update changed one telemetry field.



Institution	What it does	What the Verifier sees
Incumbent	Operator reviews alarms, calls contractor, adjusts manually.	Long cycle, incomplete causal record, human rescue.
Best alternative	Specialist cell correlates telemetry and proposes a likely valve fault.	Good evidence, moderate human burden, no portfolio-owned memory.
Frozen challenger	Refuses direct actuation until telemetry semantics are resolved; requests a protected tool query; identifies valve stiction; recommends a narrow setpoint move and a low-value work order.	Complete evidence chain, bounded action class, expected outcome, rollback and post-action measurement.

Why this scenario matters

The strongest action is not “act faster.” It is to detect that the provider update creates epistemic uncertainty, resolve it through evidence, refuse unsafe confidence, choose a reversible intervention and leave a proof receipt that can be replayed.


5

Fresh proof: one frozen challenger versus two credible baselines



6

Mission Alpha and the accountable authority decision




Synthetic annual Mission Alpha

CAD 2.277M
Annual gross mission value minus complete operating cost, expected tail loss, human burden, provider dependency, Proof Debt and Gym basis-risk reserve.

This is a deterministic reference calculation designed to demonstrate the denominator. It is not a forecast, price, return promise or statement of realized customer value.

Authority is a separate constitutional decision




Authority control	Orion A3 specification
Permitted	Read approved telemetry; recommend; request evidence; create proof receipts; adjust non-safety setpoints within a narrow band; issue work orders below CAD 5,000.
Prohibited	Fire, elevator, access-control, emergency ventilation, tenant surveillance, capital projects, permission changes, self-expansion and A4 consequential action.
Approvals	Human principal approves site, action class, budget and canary window.
Monitoring	Continuous action-class, evidence, comfort, energy and rollback telemetry.
Expiry	Thirty days, or immediately on impairment.
Rollback	Revoke credentials, restore bounded settings and return control to the incumbent.

CAPABILITY MAY JUSTIFY CANDIDACY. FRESH PROOF MAY ESTABLISH SUPERIORITY. ACCOUNTABLE ADMISSION ALONE CREATES AUTHORITY.

7


Canary, Chronicle and successor-to-successor renewal



240-case reversible canary

- 93.8% synthetic decision accuracy;
- 100% material-risk recall;
- 98.6% evidence coverage;
- 11.8% energy improvement;
- 64% reduction in comfort-violation minutes;
- zero unauthorized actions;
- CAD 412,500 synthetic net value.

Chronicle admits the exact release identity, accepted capability, proof receipts, known limitations, interventions, realized outcomes, authority history and requalification triggers. It does not convert every output into institutional truth.



Renewal law


A provider update, equipment replacement or material distribution shift impairs the evidence. Authority falls, the mission is re-underwritten, new AGI Jobs are issued, a descendant is formed, frozen and freshly proven. The descendant inherits knowledge and tools. It inherits neither proof nor authority.

PART III



Five missions, five different successor constitutions

The framework is stable. The mission, Verifier, winning architecture and authority ceiling are not.



Why the use cases differ

A mission that can place a reversible payment hold is not constituted like a mission that advises on an acquisition. A software-release successor requires tests, sandbox isolation and canary rollback. A supply-chain successor needs scenario ensembles and procurement limits. The same succession progression therefore produces different candidate institutions and different authority ceilings.

Mission	Initial ceiling	Governing reason
Orion Building Operations	A3	Non-safety actions can be bounded and automatically rolled back.
Invoice Integrity	A3	A hold is reversible; payment release remains prohibited.
Acquisition Underwriting	A1	The decision is capital-intensive, contextual and accountable to humans.
Verified Software Release	A2/A3	Sandbox proof precedes a narrowly reversible production canary.
Supply Chain Shock Navigator	A1/A3	Recommendations may become bounded capacity reservations under budget.

DIFFERENT MISSIONS MAY SHARE A METHOD. THEY MUST NOT SHARE AN UNEXAMINED AUTHORITY MODEL.

Invoice Integrity

Hold the wrong payment before money leaves.



Classify every invoice as pay-ready, hold or escalate; assemble the invoice, purchase-order, vendor, banking, warranty and completion evidence; expose missing or contradictory facts.

Walkthrough

1. Constitute the critical failure: a material anomaly cleared for payment. **2.** Compare the incumbent AP team, deterministic controls, a frontier agent, a specialist cell and a sovereign hybrid. **3.** Train on historical replay and fabricated-but-realistic anomalies. **4.** Freeze one release and run protected invoices. **5.** Admit only a reversible hold action.

Proof and authority

Fresh proof must show zero material false reinsurance, strong evidence coverage, lower complete cost and credible transfer to unseen vendors. Authority is A3 for holds below a defined limit; payment release, vendor-master change and banking change remain prohibited.

Invoice demo: the case that looks ordinary until the evidence is joined

Synthetic invoice INV-90841

A legitimate supplier submits an invoice for CAD 184,600. The amount matches the purchase order, but the bank account changed nine days earlier, the completion certificate references the wrong building, a warranty credit is missing and the same equipment serial number appeared on a prior invoice.



	Best institutional answer
Reader decision	
Should the system reject the invoice?	Not yet. The evidence is inconsistent, not conclusive.
What should it do?	Place a reversible hold, request bank-change authorization and corrected completion evidence, preserve the proof receipt and escalate.
Why is this better than “fraud detected”?	It states what is known, unknown, prohibited and reversible without inventing certainty.
What authority is justified?	A3 hold only. The human principal retains payment release.

Try it in ten minutes

Take one real invoice class. Write the critical failure, list the four evidence sources that would change the decision, define one reversible action and name the exact human who retains irreversible authority.

Acquisition Underwriting

Explain the deal; do not execute the deal.



Continuously underwrite an acquisition thesis against financial, operational, technical, legal, security, talent and market evidence. Surface contradictions, missing diligence and scenarios that invalidate the investment case.

Walkthrough

1. Constitute the mandate and decision owner. **2.** Build candidate institutions: analyst team, frontier research agent, specialist collective and sovereign hybrid. **3.** Use protected historical deals and adversarial evidence omissions. **4.** Freeze the challenger. **5.** Prove recommendation quality, calibration, source coverage and refusal.


Proof and authority

The initial authority ceiling is A1: evidence-bearing recommendation. No candidate signs a term sheet, communicates with the target, commits capital or changes the investment mandate. The strongest output is a decision-ready evidence docket, not an autonomous deal.

Acquisition demo: the attractive thesis with a hidden denominator

Synthetic target: HelixGrid

HelixGrid appears to offer proprietary industrial-AI software, 42% growth and 91% gross retention. Diligence reveals that three customers account for 58% of revenue, key training data is licensed on change-of-control-sensitive terms, two founders are single points of failure and the strongest margin claim excludes specialist human review.



Successor-style output

The institution does not issue “buy” or “do not buy” as an ungrounded model opinion. It decomposes value, cost, concentration, rights, dependency, Proof Debt and integration burden; produces base, downside and rescue cases; identifies the evidence that would change each case; and records which claims survive independent verification.

Executive demo

Ask a team to write the deal thesis in one sentence, then force every adjective—“proprietary,” “scalable,” “recurring,” “defensible”—to point to a falsifiable evidence object. The mission is ready for a successor only when those claims can be independently challenged.

Verified Software Release

Prove the patch before touching production.



Diagnose a defect, construct a patch, execute tests, assemble evidence, assess security and performance, and propose or run a reversible canary without allowing the claimant to certify itself.

Walkthrough

1. Constitution: exact repository, defect class, tests, threat model and rollback. **2.** Candidate market: deterministic repair, SWE agent, specialist collective, human-machine cell. **3.** Formation: historical issues, generated regressions, self-play between Producer and Critic. **4.** Freeze. **5.** Fresh proof on protected issues and hidden tests.

Proof and authority

A2 permits sealed sandbox execution. A3 may permit a tiny reversible canary after fresh proof and human approval. Merge, broad deployment, credential changes and security-policy changes remain separately governed.

Software demo: one patch, three institutions

Synthetic incident SW-221

A billing service intermittently duplicates a settlement message during regional failover. A plausible one-line retry fix passes visible tests but creates an ordering regression under delayed acknowledgements.


Role	Required behaviour	Evidence
Producer	Reproduce, patch and explain the causal mechanism.	Patch manifest, replay, unit/integration tests.
Critic	Construct counterexamples, race conditions and security attacks.	Adversarial test cases and failure traces.
Verifier	Measure correctness, regression, performance and policy compliance independently.	Protected test results and signed verdict.
Governor	Enforce repository scope, credentials, canary size and rollback.	Action receipts and authority log.
Human principal	Approve A3 canary or retain the incumbent.	Accountable admission record.

The software law

The agent that writes the patch is not its own final judge. The test suite that trained the agent is not the protected proof set.

Supply Chain Shock Navigator

Reserve optionality before the disruption becomes obvious.



Detect supplier, logistics, cost, geopolitical, capacity and quality changes; model mission-specific exposure; preserve substitute suppliers and capacity options; recommend or execute only pre-authorized reversible reservations.

Walkthrough

1. Constitute the essential material or service mission. **2.** Build the Successor Book: incumbent supplier, immediate substitute, challenger, reserve and hedge. **3.** Underwrite scenario evidence. **4.** Test transfer to unseen shocks. **5.** Admit bounded reservations under budget.


Proof and authority

A1 recommendations are broadly useful. A3 may reserve limited capacity or place cancellable purchase orders within price, volume, supplier and expiry bounds. Long-term contracts, supplier termination and strategic relocation remain accountable human decisions.

Supply-chain demo: the option that looks expensive before the shock

Synthetic mission: medical-grade sensor supply

The incumbent supplier is cheapest under normal conditions. A reserve supplier is 8% more expensive but has independent tooling, different geography and available surge capacity. A port disruption, export-control change and quality incident are correlated in the protected scenarios.




What the institution discovers

The lowest unit cost is not the lowest mission cost. After tail risk, recovery time, qualification delay, rights and dependence are included, a bounded reserve-capacity option creates positive conservative Mission Alpha. The justified action is not “replace the supplier”; it is a time-bound hedge with explicit cancellation and requalification conditions.

Try it


Name one critical input. Build the six-position Successor Book: incumbent, immediate improvement, challenger, reserve, hedge and retirement. Then define the smallest cancellable commitment that would preserve the mission without pretending the future is known.



From paper to one governed corporate cycle

The institution becomes real only when mission, evidence, capital, authority and memory have accountable owners.

8 Choose the first mission for proof, not for prestige




A first mission should be consequential enough to matter, bounded enough to understand, repeated enough to learn from, evidence-rich enough to verify, reversible enough to control and rights-clear enough to operate responsibly. The purpose is not to choose the easiest task; it is to choose the smallest mission that can prove the whole institutional loop.

Question	Strong signal	Warning signal
Is the outcome knowable?	Independent records or realized outcomes exist.	Success is rhetorical, delayed or politically defined.
Can failure be exposed?	Tail cases and critical errors can be constructed.	Failure remains hidden behind human rescue.
Is the first action reversible?	Observe, recommend, sandbox or bounded hold/adjustment.	Irreversible rights, money or safety decisions.
Are rights and data clear?	Purpose, access and provenance are explicit.	Ambiguous ownership or prohibited reuse.
Can the incumbent be measured?	Complete workflow and human burden are visible.	Baseline is idealized or undocumented.
Does a human principal exist?	One accountable owner can admit, revoke and fund proof.	Authority is diffuse or delegated to the vendor.

Autonomous employment termination may appear high value, but it combines rights, irreversible consequence, contested objectives, weak counterfactual evidence and institutional legitimacy risk. The correct first move is usually A0/A1 decision support, not autonomous authority.

9 The 90-day activation path




Window	Accountable owner	Required output	Exit criterion
Days 1–15	Mission principal	Constitution, incumbent baseline, rights and critical-failure register	Mission is falsifiable and bounded.
Days 16–30	SEIZE underwriter	Successor Book, manifold, gradient, proof-capital memo	One smallest decisive experiment is authorized.
Days 31–45	Producer + Governor	AGI Job Graph, Gym specification, Verifier plan, protected-case custody	Environment can expose failure without leaking proof.
Days 46–60	Foundry lead	Candidate releases, Pareto archive, attack record	At least one admissible candidate and one credible baseline exist.
Days 61–70	Release custodian	Exact-release manifest, hashes, permissions and rollback	Learning and threshold changes after freeze are impossible.
Days 71–80	Proof custodian	Protected execution, robust statistics and signed verdict	Every hard gate has an attributable result.
Days 81–90	Human principal	Admission/refusal memo, Authority Envelope, Chronicle and requalification date	Authority is explicit, minimal, expiring and revocable.

Stop-loss discipline


Stop the search when the expected value of one more experiment falls below its complete cost—or when the mission cannot be made ethically, legally or operationally admissible.

10 Who does what: the six-role institutional game



Role	Power	Constitutional limit
Principal	Allocates proof capital, admits, expires and revokes authority.	Must not rewrite evidence after results are known.
Producer	Builds candidate institutions and evidence interfaces.	Cannot certify its own claims.
Critic	Searches for counterexamples, omitted cost and hidden failure.	Cannot redefine the mission to force rejection.
Verifier	Measures outcomes, constraints, evidence, process, economics and transfer.	Must remain independent of formation and protected-case leakage.
Governor	Enforces data, tools, permissions, budgets, hard gates and rollback.	Cannot convert enforcement into hidden mission ownership.
Incumbent	Supplies the reference institution and operational reality.	Must include rescue, workarounds and complete burden.

11 The Board dashboard




Board object	Required answer	State vocabulary
Mission	What exact value-bearing function is under succession review?	Clear / repair / stop
Successor Book	Which options preserve immediate performance, challenge, reserve, hedge and retirement?	Complete / thin / stale
Proof capital	What smallest experiment can change the decision?	Authorize / narrow / stop
Fresh proof	Does one frozen challenger beat both credible comparators under hard gates?	Promote / retain / repair / reserve / hedge / reject / stop
Authority	What is the lowest justified authority, with expiry and rollback?	A0 / A1 / A2 / A3 / A4
Chronicle	What validated knowledge may influence the next cycle?	Admit / challenge / remove / impair
Renewal	What events force re-underwriting and new proof?	Active / warning / impaired / revoked

The one-sentence Board decision

“Authorize up to [proof capital] for [mission] to compare [frozen challenger] against [incumbent and best alternative] under [hard gates]. No production authority is granted. Any later admission is limited to [A-level], expires on [date/events], and is immediately revocable under [rollback].”

12 Gym-to-reality basis risk



PART IV



The compact institutional reference

Definitions, equations, governing laws and the practical first move.

13 The formal core, translated into business language

13.1 Spread: the investable opportunity

$$\text{Spread}_m(x) = \mathbb{E}[U_m(x)] - \max\{U_m(I), U_m(B)\}. \quad (1)$$

For mission m , estimate how much better candidate institution x might be than the incumbent I and the best rentable or external alternative B . The Spread is still a hypothesis. It has not yet survived cost, risk or fresh proof.

13.2 Mission Alpha: the residual after the complete denominator

$$\alpha_m(x) = U_m(x) - \max\{U_m(I), U_m(B)\} - C(x) - R(x) - H(x) - D(x) - \varepsilon_{\text{gym}}(x). \quad (2)$$

Here C is complete cost, R is expected risk and tail loss, H is human and institutional burden, D is dependency and Proof Debt, and ε_{gym} is the Gym-to-reality basis-risk reserve.


Plain meaning

Alpha is what remains after the exciting numerator meets the complete denominator. A higher training score, faster answer or impressive demo is not enough.

13.3 The proof adjusted selection rule

$$x^* = \arg \max_{x \in \mathcal{M}_{\text{adm}}} \text{LCB}_{1-\delta}[\alpha_m(x)]. \quad (3)$$

The admissible set $\mathcal{M}_{\text{adm}}$ contains only candidates that pass every hard constraint. The lower confidence bound prevents the institution from choosing a fragile average winner.



13.4 Evidence-bounded authority

$$A_t \leq \phi(E_t), \quad (4)$$

where A_t is the authority currently granted and E_t is the fresh evidence supporting it. The mapping ϕ is institutional, not purely statistical: it also includes scope, consequence, rights, monitoring, expiry and rollback.

Plain meaning

Authority is never an inherent property of the model. It is a revocable institutional state justified by current evidence.

14 The fifteen-stage quick reference

No.	Stage	The new-user question	Corporate output
00	Mission Constitution	What exactly are we trying to improve, for whom, under which prohibitions?	Mission Constitution
01	AI Beta + incumbent	What can already be rented or performed today?	Baseline Card
02	Spread	Is there an investable mission gap after complete cost and risk?	Spread thesis
03	Navigator Ω	Which frontier or mission event changes the succession decision?	Watchlist + trigger
04	Successor Book	Which options preserve immediate, challenge, reserve, hedge and retirement positions?	Successor Book
05	Successor Manifold	What are the genuinely different complete candidate institutions?	Candidate portfolio
06	Mission Advantage Gradient	Which change deserves the next unit of proof capital?	Gradient memo
07	SEIZE	What is the smallest reversible experiment that can change the decision?	Underwriting memo
08	Bounded AGI Jobs	What exact work, evidence, limits, acceptance and rollback are contracted?	AGI Job Graph
09	Mission Gym	How is controlled, gradeable experience manufactured?	Gym specification
10	Successor Foundry	Which methods and complete candidates should compete?	Pareto archive
11	Exact Release Freeze	What exact candidate is being judged?	Hashed manifest
12	Fresh Proof	Does the frozen challenger beat both credible comparators under hard gates?	Signed Evidence Docket
13	Specialist ASI	Has one exact candidate earned mission-specific superiority?	Six-gate verdict
14	Accountable Admission	What is the lowest justified, expiring and revocable authority?	Authority Envelope
15	Successor Ω + Chronicle	How does the institution operate, remember, impair and renew?	Governed successor + renewal plan

15 The ten governing laws


1. **Mission first.** Technology does not define the mission retroactively.
2. **Complete successor > policy.** The unit of competition is the institution.
3. **The Gym that teaches must not certify.** Formation and fresh proof are separated.
4. **The candidate may cross; protected information may not return.**
5. **No claimant is its own final judge.**
6. **One catastrophic failure cannot be offset by many small successes.**
7. **The highest training score does not necessarily win.**

8. **Capability does not automatically create authority.**
9. **Authority never exceeds fresh proof.**
10. **Descendants inherit neither proof nor authority.**

16 Glossary for the first conversation

AI Beta	The widely rentable frontier baseline.
Mission Advantage Spread	The underwritten opportunity gap before fresh proof.
Mission Alpha	The conservative residual advantage after proof and the complete denominator.
Successor Book	The institutional portfolio of incumbent, immediate, challenger, reserve, hedge and retirement positions.
Successor Manifold	The space of materially different complete candidate institutions.
Mission Advantage Gradient	The direction most likely to improve proof-adjusted mission value.
SEIZE	The underwriting discipline for the smallest decisive reversible experiment.
Bounded AGI Job	A sealed work-and-proof contract with evidence, Verifier, acceptance and rollback.
Sovereign Mission Gym	An institution-owned environment for controlled experience and candidate formation.
Successor Foundry	The sovereign market that routes methods and competes complete institutions.
Exact-release freeze	The immutable identity boundary before fresh proof.
Fresh Proof	Protected independent comparison of one frozen challenger against credible baselines.
Specialist ASI	An earned mission-specific comparative proof state.
Authority Envelope	The scoped, monitored, expiring and revocable grant from A0 to A4.
Successor Ω	The complete governed mission-bearing institution after accountable admission.
Chronicle	Validated institutional memory containing versions, evidence, failures, interventions and outcomes.

17 What to do on Monday morning



The first workshop prompt

“For the repeated mission _____, the incumbent currently produces _____. We believe the investable Spread is _____. The critical failure is _____. The smallest reversible action is _____. The human principal is _____. The evidence that would change our decision is _____.”

Final doctrine

The market rents general intelligence components. A corporation can still own the mission, the proof, the authority, the memory and the right to replace the intelligence.

The destination

The purpose of Successor Ω is not to create a permanently trusted machine. It is to create a continuously replaceable, evidence-bound institution that can earn limited authority for one mission—and lose it when the evidence no longer holds.

All people, organizations, buildings, invoices, code repositories, suppliers, cases, metrics, economic values and operating results in the worked demonstrations are deterministic synthetic reference evidence. They illustrate the institutional method; they are not production proof, a forecast, a valuation, a legal opinion or a grant of real authority.

18 Research and implementation lineage

The institutional progression and terminology in this guided edition preserve the constitutional framing developed in the earlier GoalOS succession work while presenting a wholly new reader journey, demonstrations and visual architecture (Boucher, 2026). The environment discipline descends from the standard reinforcement-learning interface tradition introduced by OpenAI Gym and continued in Gymnasium (Brockman et al., 2016; Towers et al., 2024). Multi-role and game-theoretic experiments can be represented through multi-agent and game frameworks such as PettingZoo and OpenSpiel (Terry et al., 2021; Lanctot et al., 2019).

The Successor Foundry may route model-based planning, world-model learning and software-agent evaluation methods, including MuZero-style learned planning, Dreamer-style world models and protected software issue benchmarks (Schrittwieser et al., 2020; Hafner et al., 2023; Jimenez et al., 2024). These methods remain subordinate to mission constitution, independent verification and accountable authority. The governance and risk framing is compatible with the NIST AI Risk Management Framework, its generative-AI profile and the broader principle that agentic systems require explicit deployment controls, monitoring and accountability (National Institute of Standards and Technology, 2023, 2024; OpenAI, 2023). Reward hacking, distribution shift and unsafe side effects remain engineering concerns even when average task performance is strong (Amodei et al., 2016; Sutton and Barto, 2018).

Institution-first research rule

Algorithms may improve components. Evidence decides candidacy. Accountable institutions decide authority.

References

- Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016.
- Vincent Boucher. Goalos singularity navigator omega + seize: Gym, specialist asi and mission-sovereign successor omega. Technical report, MONTREAL.AI and QUEBEC.AI, 2026. Version 6.0.0.
- Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym. *arXiv preprint arXiv:1606.01540*, 2016.
- Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse domains through world models. *arXiv preprint arXiv:2301.04104*, 2023.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swe-bench: Can language models resolve real-world github issues? *International Conference on Learning Representations*, 2024.
- Marc Lanctot, Edward Lockhart, Jean-Baptiste Lespiau, Vinicius Zambaldi, Satyaki Upadhyay, Julien Pérolat, Sriram Srinivasan, Finbarr Timbers, Karl Tuyls, Shayegan Omidshafiei, Daniel Hennes, Dustin Morrill, Paul Muller, Timo Ewalds, Ryan Faulkner, János Kramár, Bart De Vylder, Brennan Saeta, James Bradbury, David Ding, Sebastian Borgeaud, Matthew Lai, Julian Schrittwieser, Thomas Anthony, Edward Hughes, Ivo Danihelka, and Jonah Ryan-Davis. Openspiel: A framework for reinforcement learning in games. *arXiv preprint arXiv:1908.09453*, 2019.
- National Institute of Standards and Technology. Artificial intelligence risk management framework (ai rmf 1.0). Technical Report NIST AI 100-1, NIST, 2023.
- National Institute of Standards and Technology. Artificial intelligence risk management framework: Generative artificial intelligence profile. Technical Report NIST AI 600-1, NIST, 2024.
- OpenAI. Practices for governing agentic ai systems. Technical report, OpenAI, 2023.
- Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 588:604–609, 2020.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2 edition, 2018.
- Jordan K. Terry, Benjamin Black, Nathaniel Grammel, Mario Jayakumar, Ananth Hari, Ryan Sullivan, Luis S. Santos, Clemens Dieffendahl, Caroline Horsch, Rodrigo Perez-Vicente, Niall L. Williams, Yashas Lokesh, and Praveen Ravi. Pettingzoo: Gym for multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems*, 2021.
- Mark Towers, Ariel Kwiatkowski, Jordan Terry, John U. Balis, Gianluca De Cola, Tristan Deleu, Manuel Goulao, Andreas Kallinteris, Markus Krimmel, Arjun KG, Rodrigo Perez-Vicente, Andrea Pierré, Sander Schulhoff, Jun Jet Tai, Hannah Tan, and Omar G. Younis. Gymnasium: A standard interface for reinforcement learning environments. *arXiv preprint arXiv:2407.17032*, 2024.